

# Lower Envelope Algorithms in $\mathbb{R}^3$

C++ Implementation, Parallelization and Benchmarking

# Outline

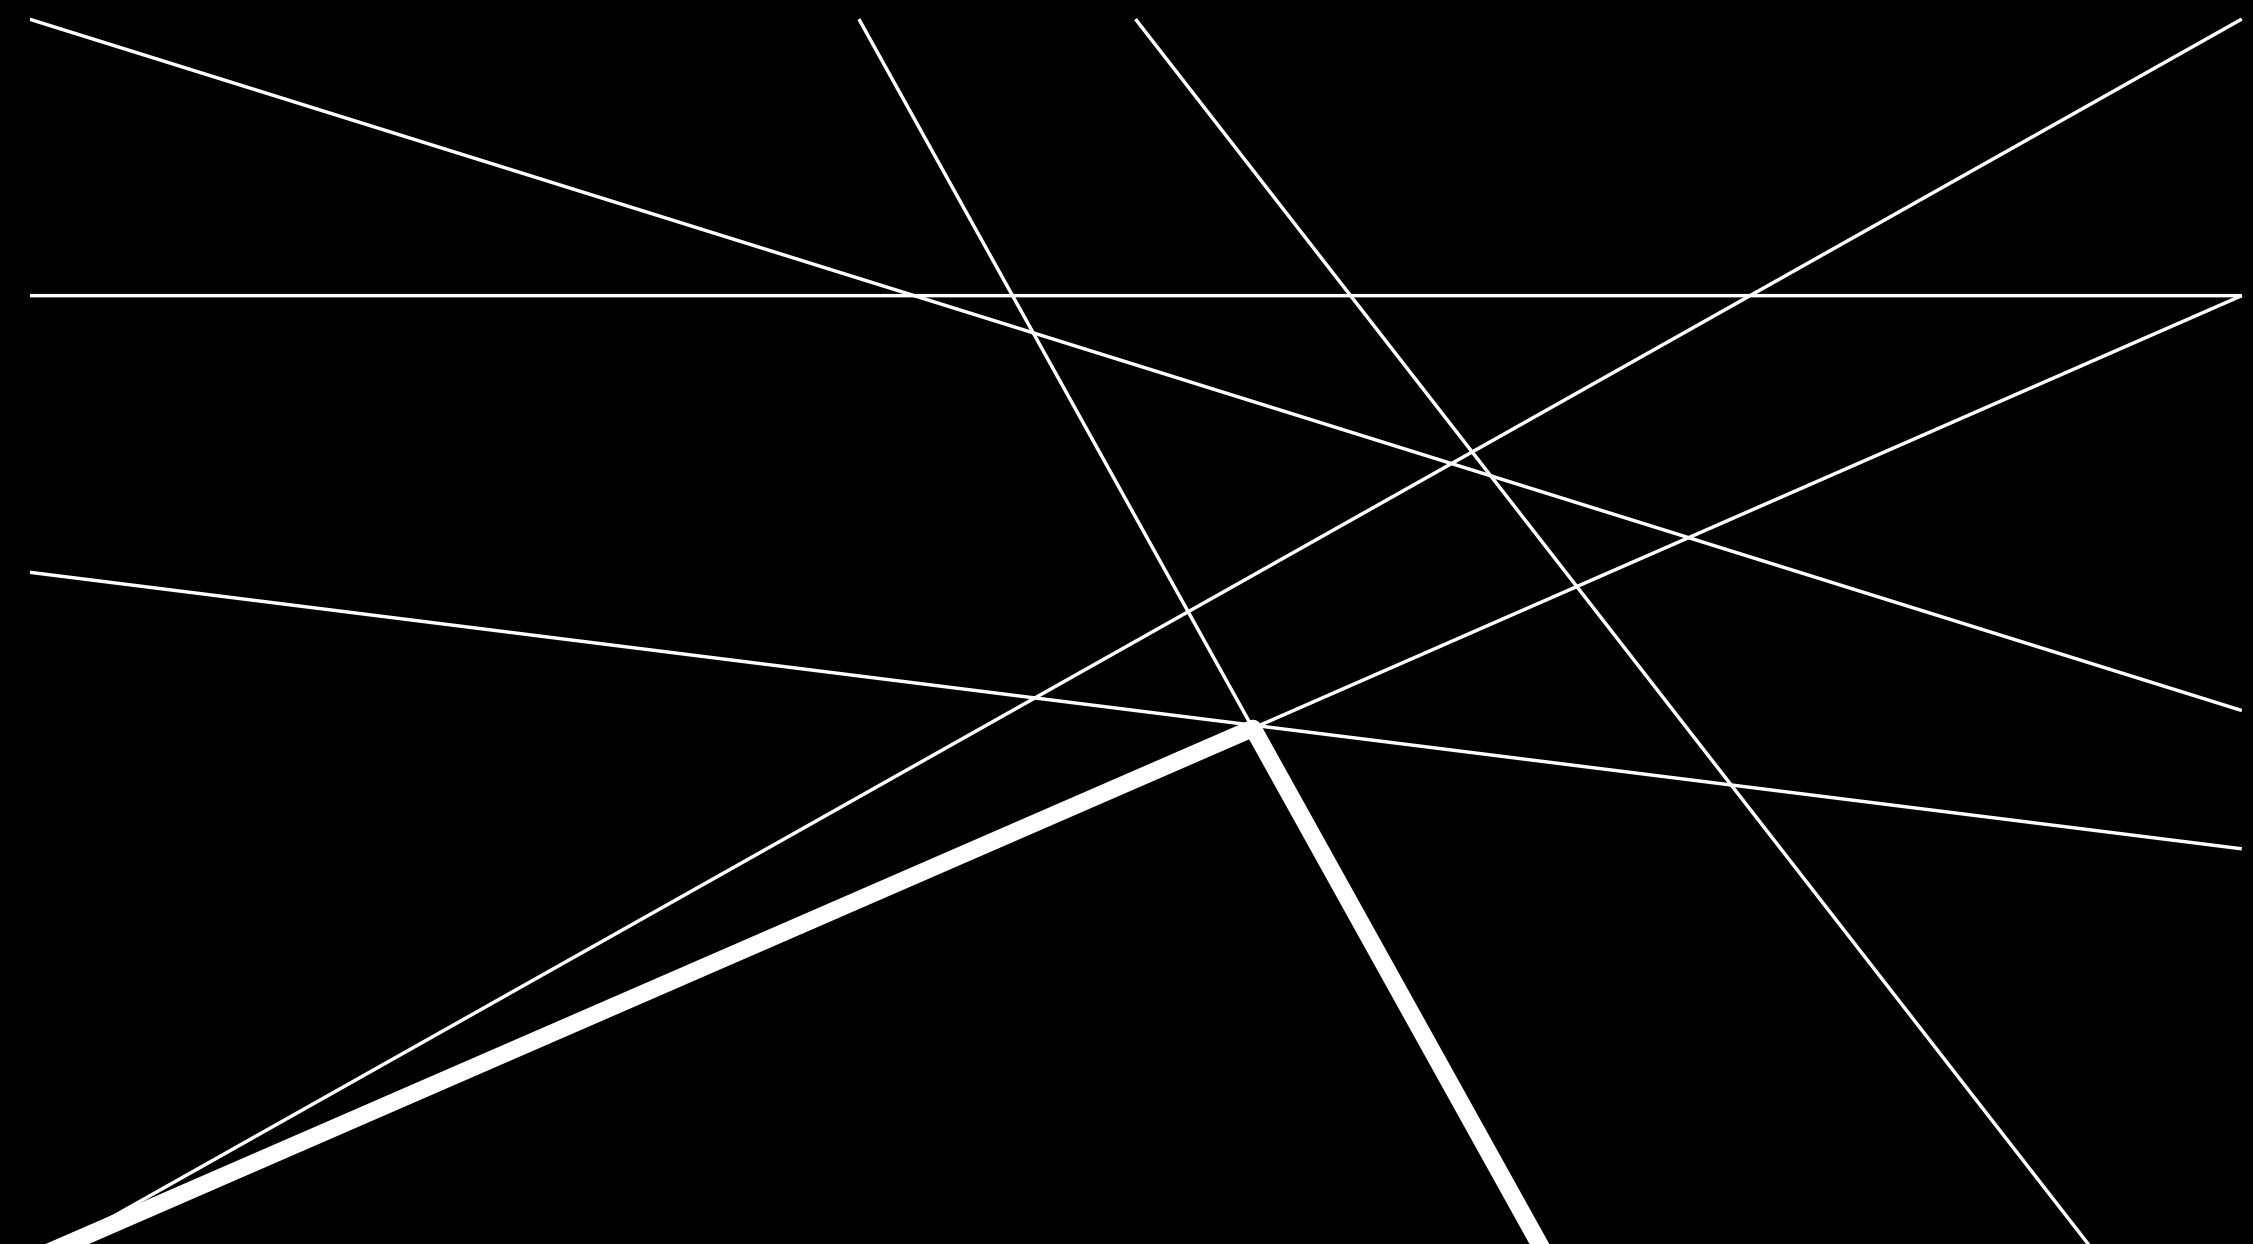
- Introduction
- Motivation and Goal
- Overview
- Algorithms
- Implementation Details
- Visualisation
- Benchmarks and Results

# Introduction

# Introduction (1)

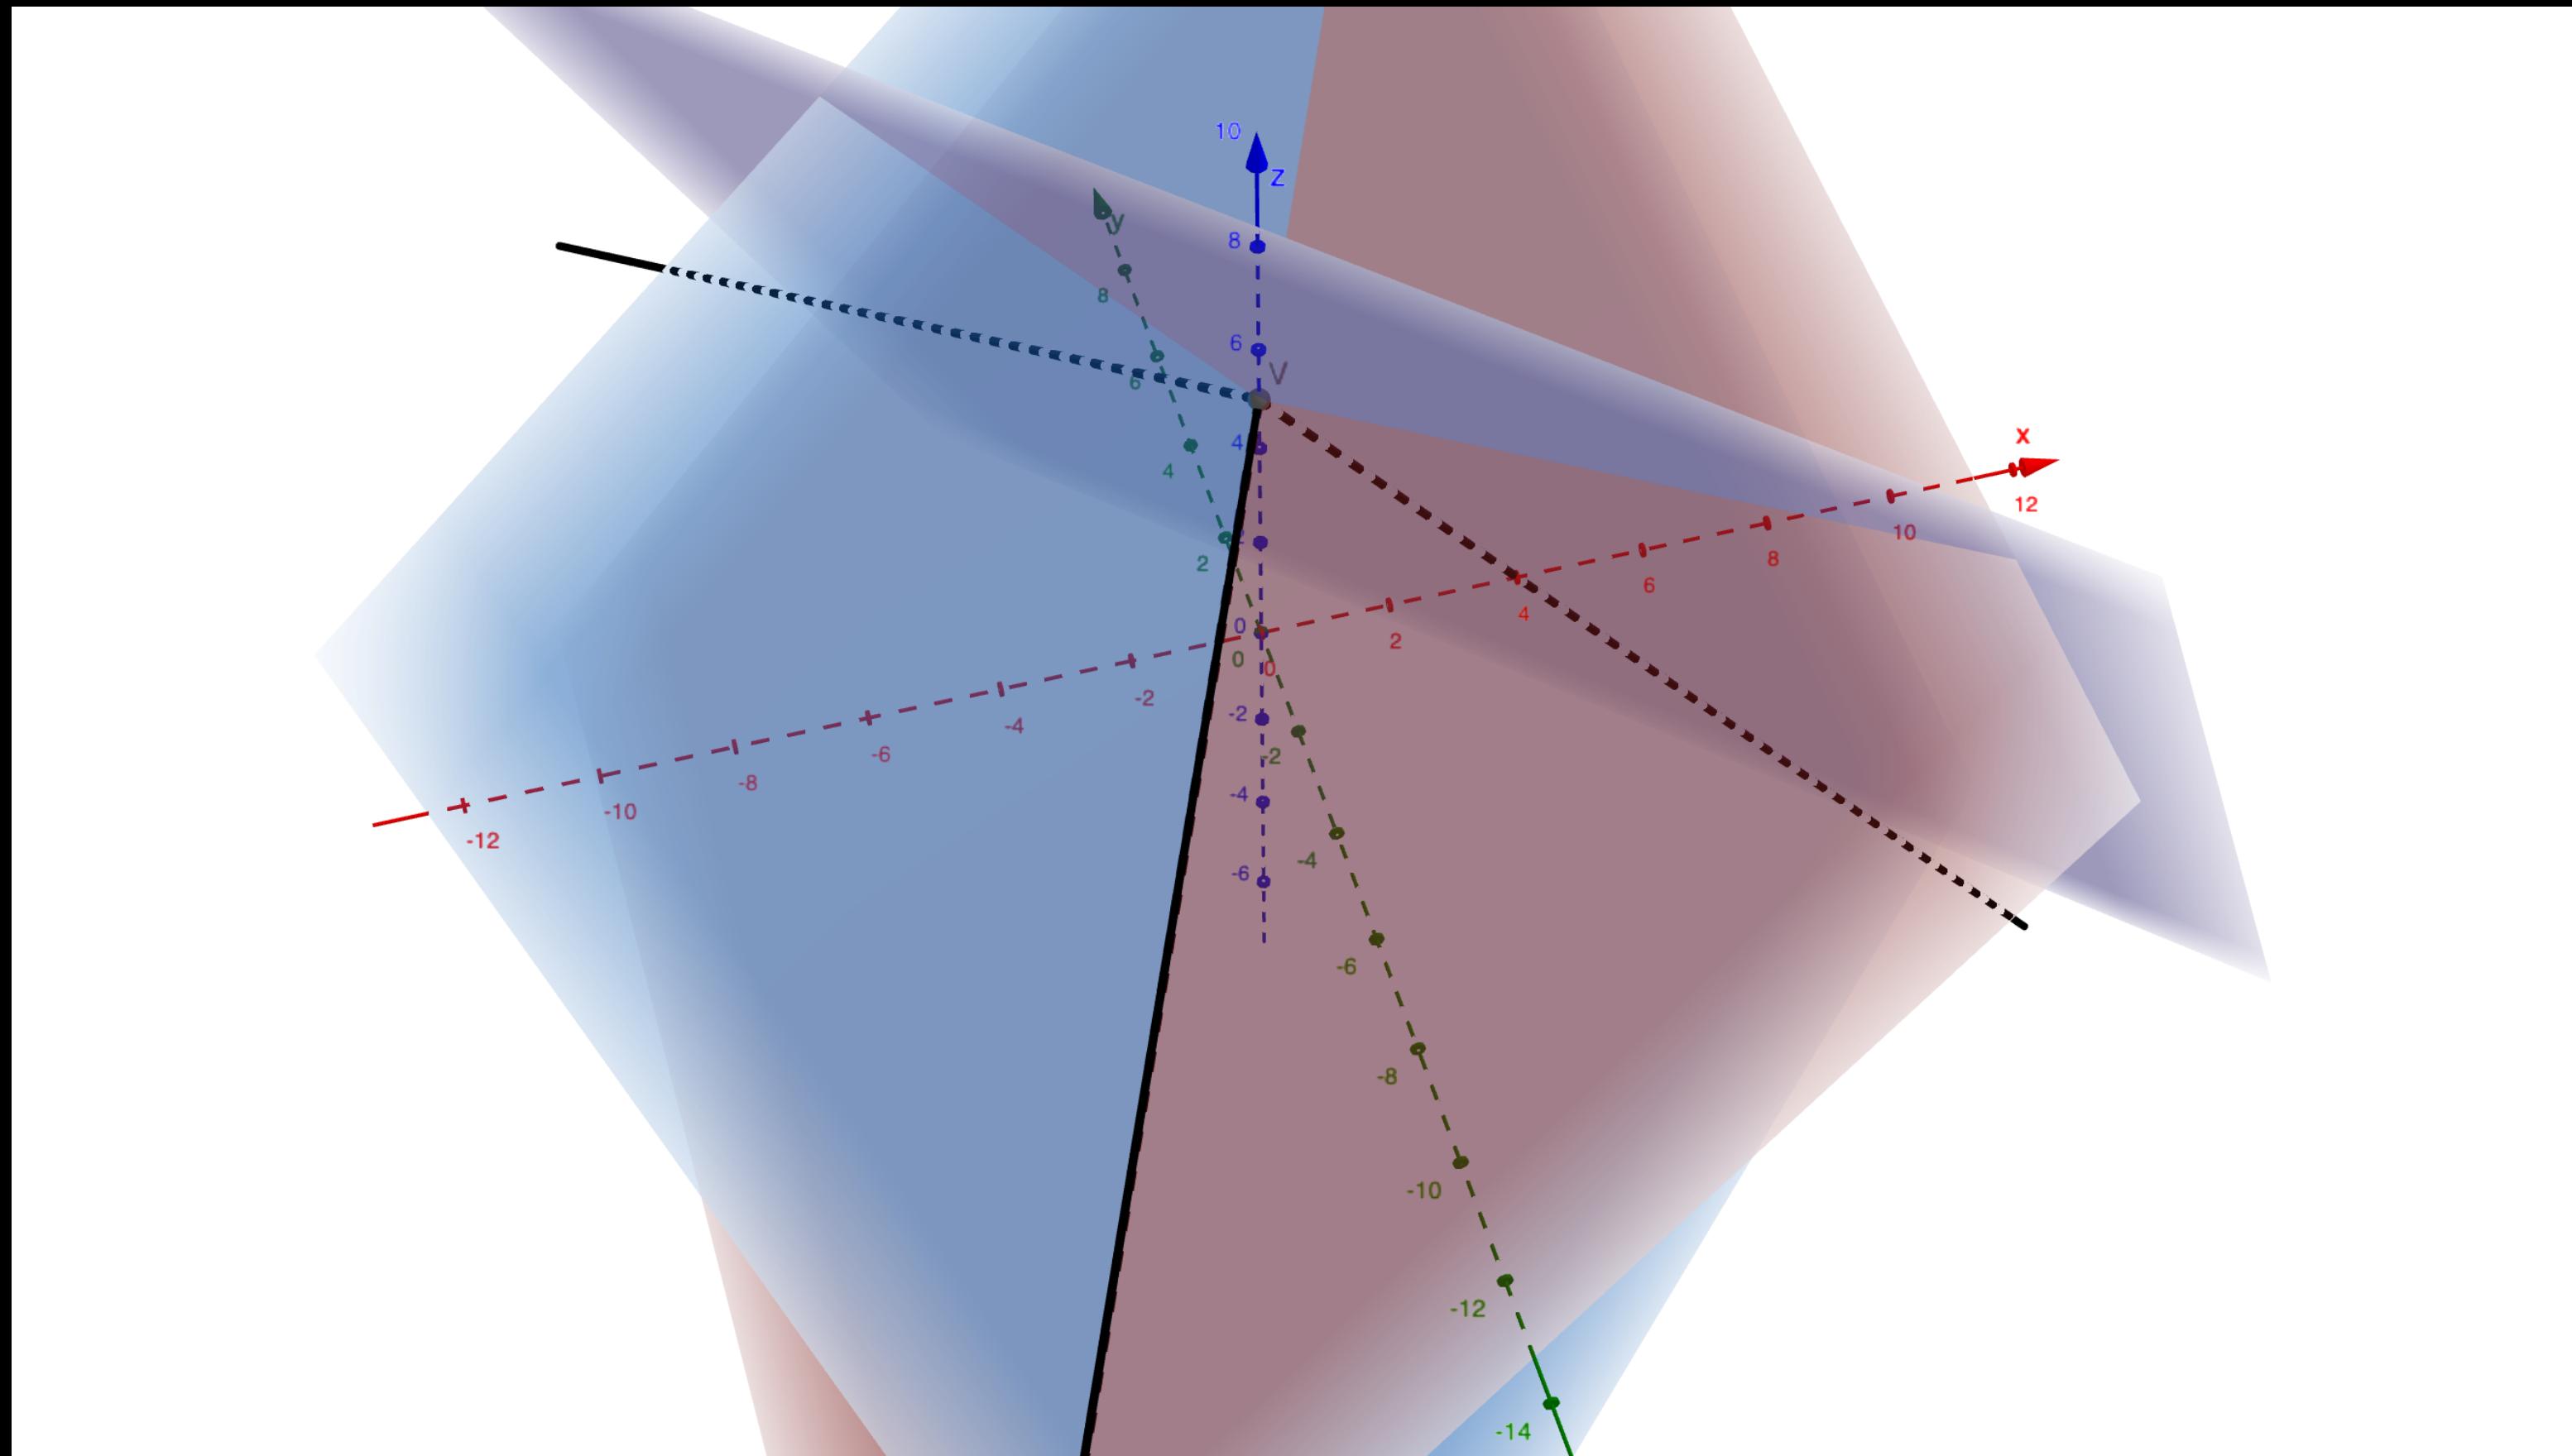
## Lower Envelope $L_0$

- The function whose value at every point is the minimum of the values of the functions in the given set  $H$



# Introduction (2)

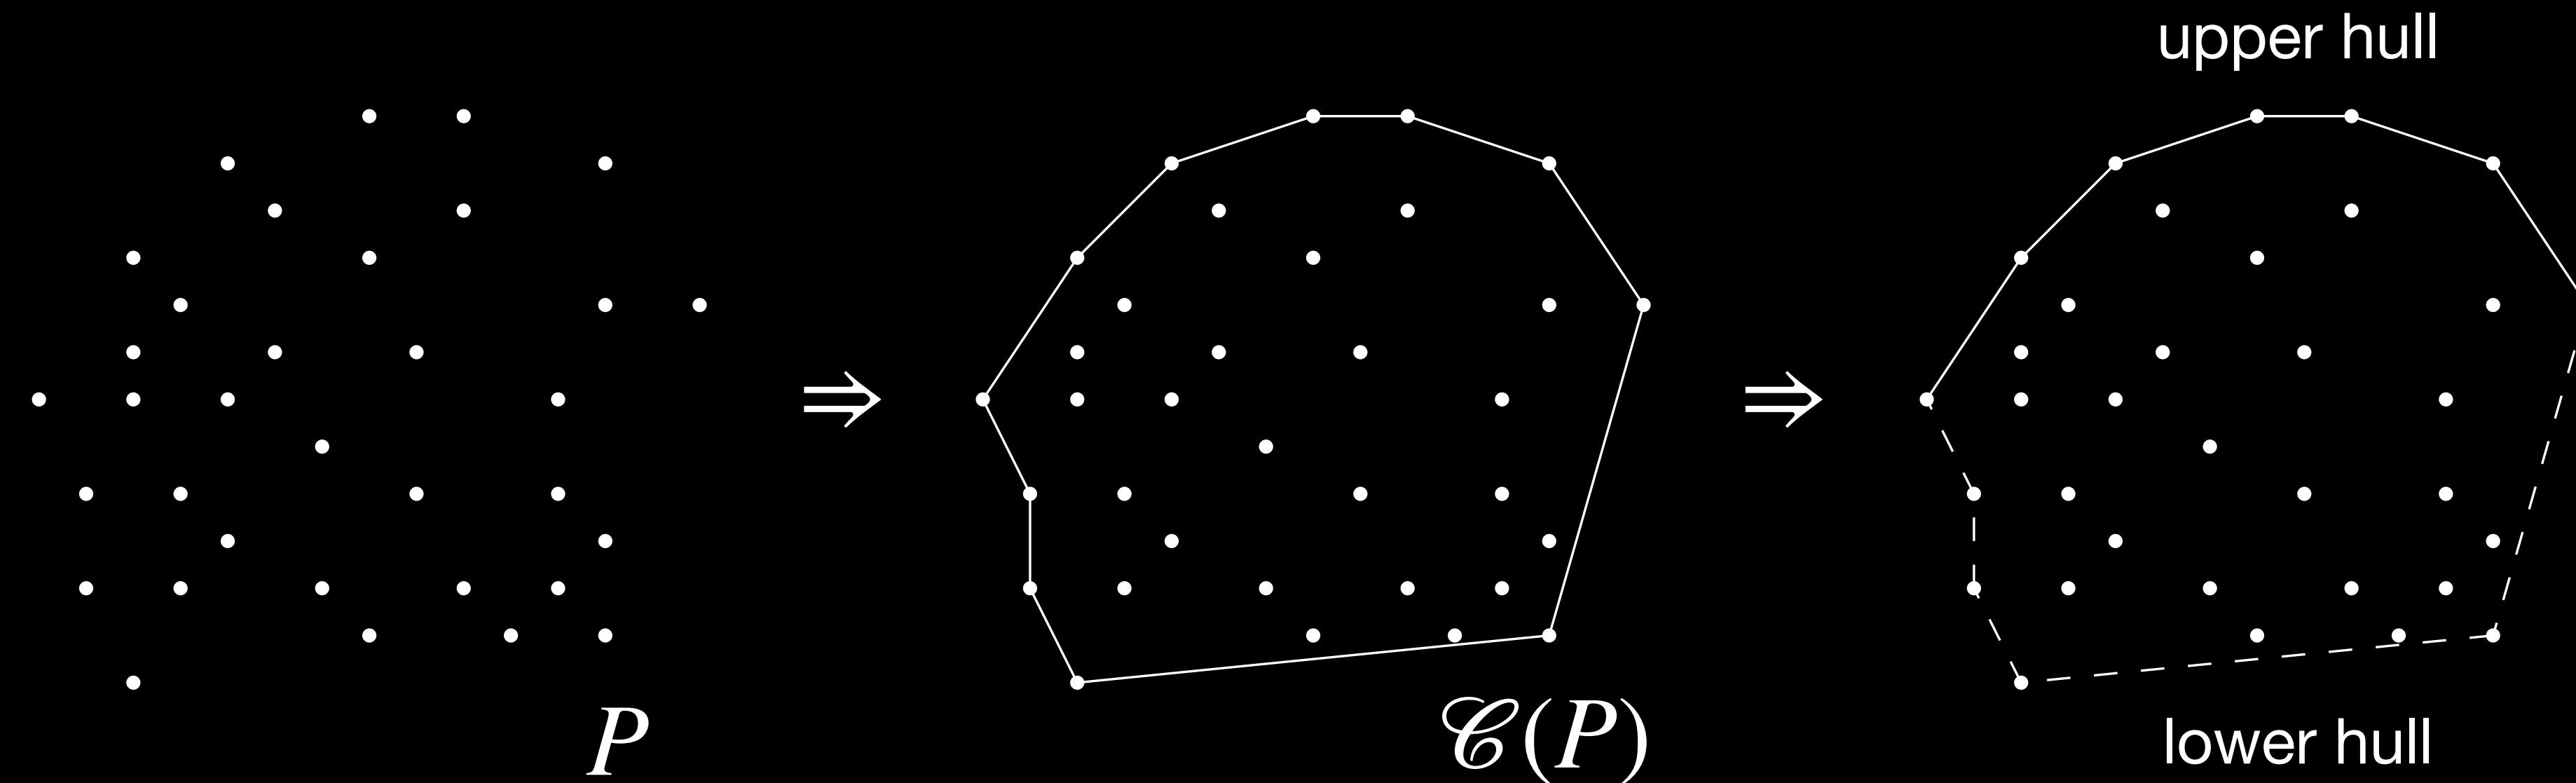
## Lower Envelope $L_0$ (GeoGebra visualisation)



# Introduction (3)

## Convex Hull $\mathcal{C}$

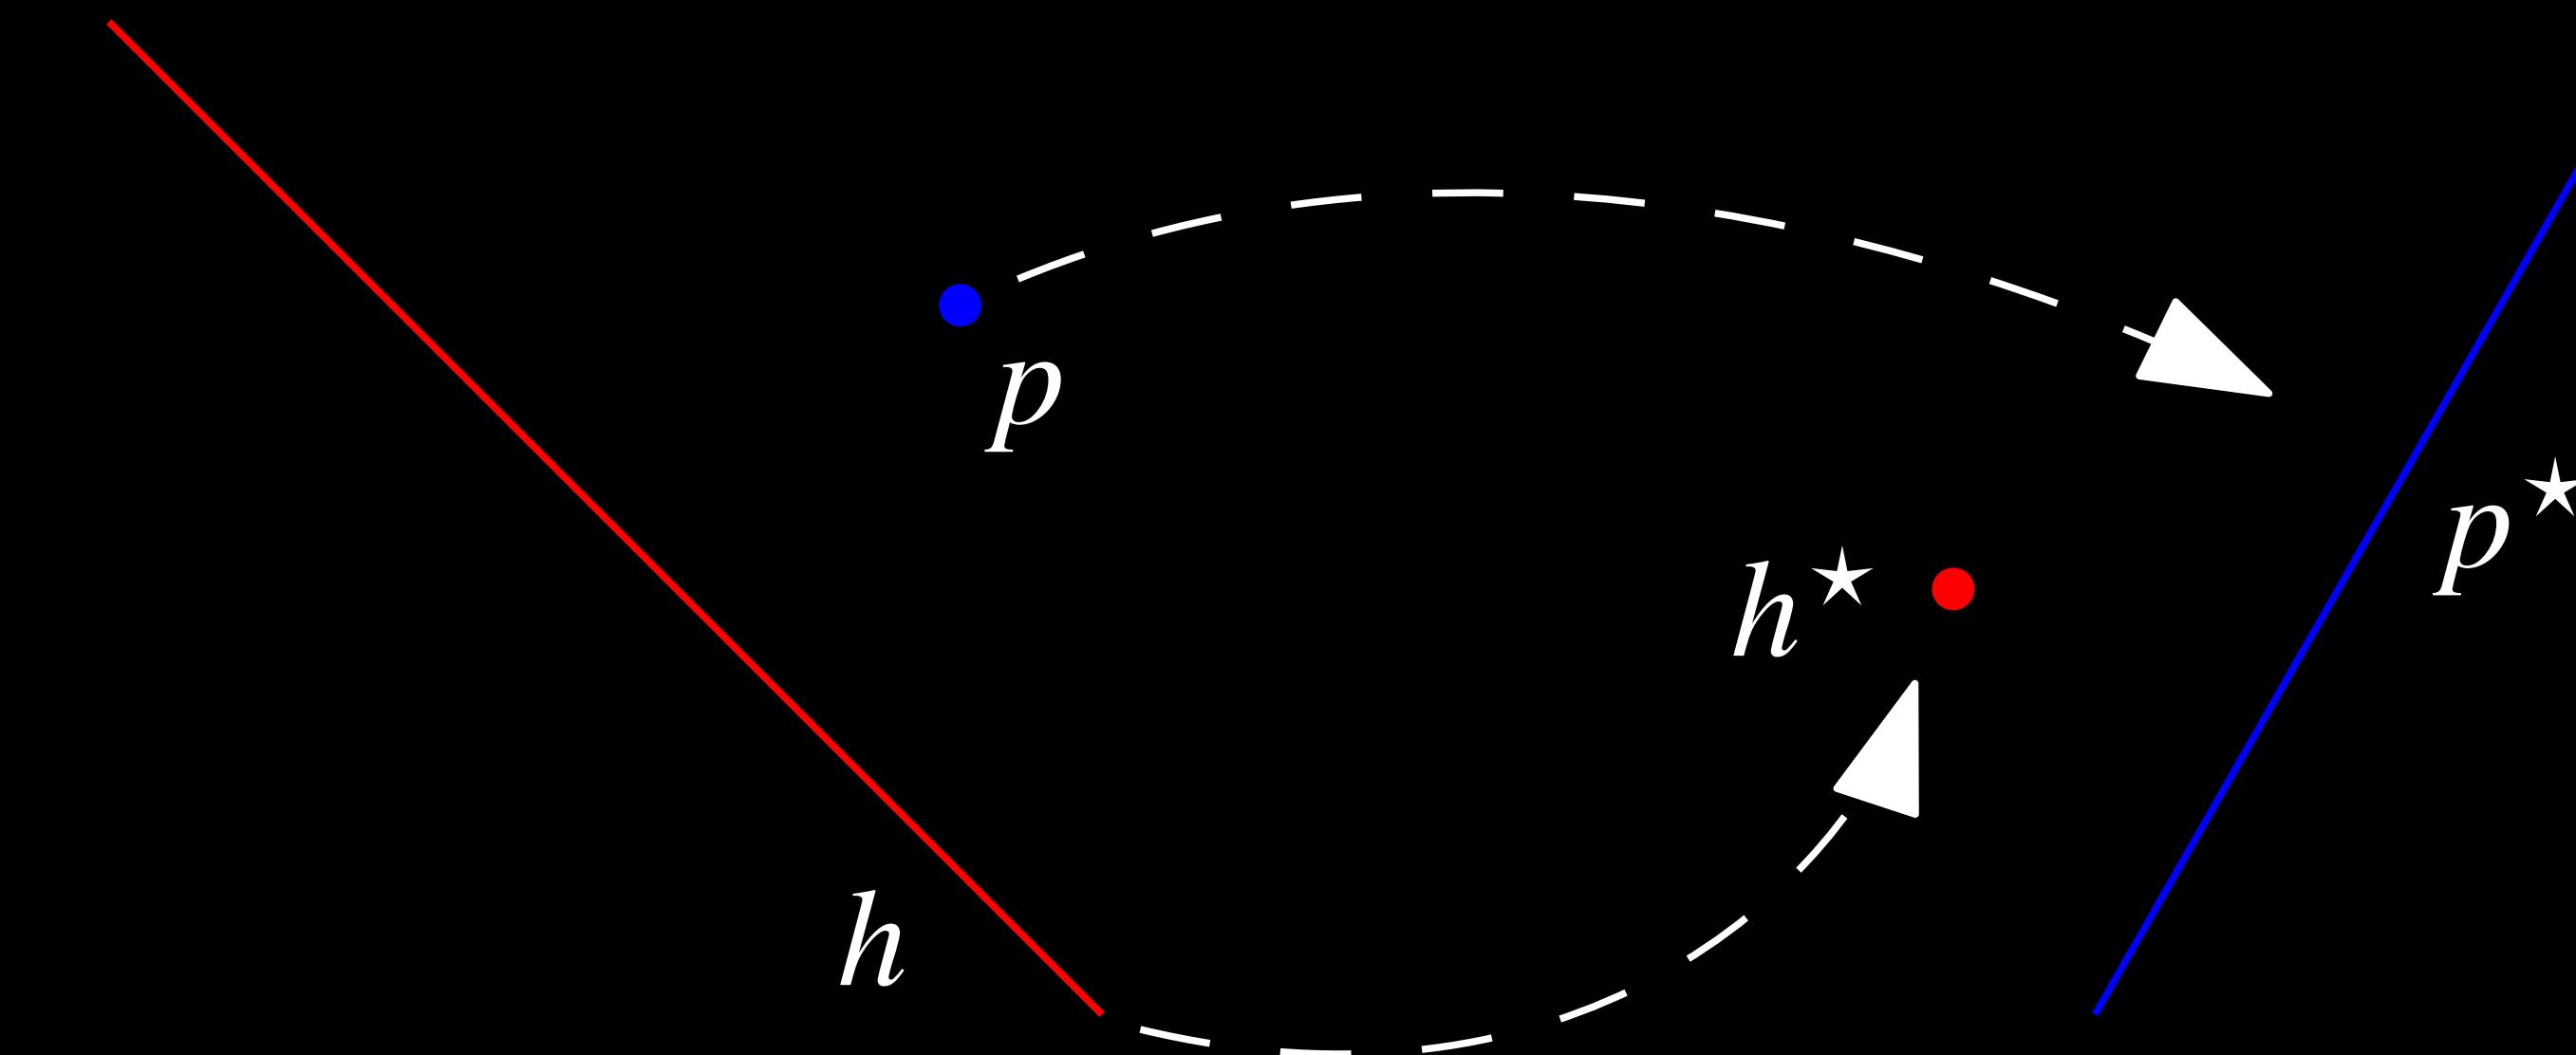
- The smallest convex polygon enclosing all the points in the input set
- Convex hulls can be partitioned in upper and lower hulls



# Introduction (4)

## Duality Transform ( $\cdot$ ) $^*$

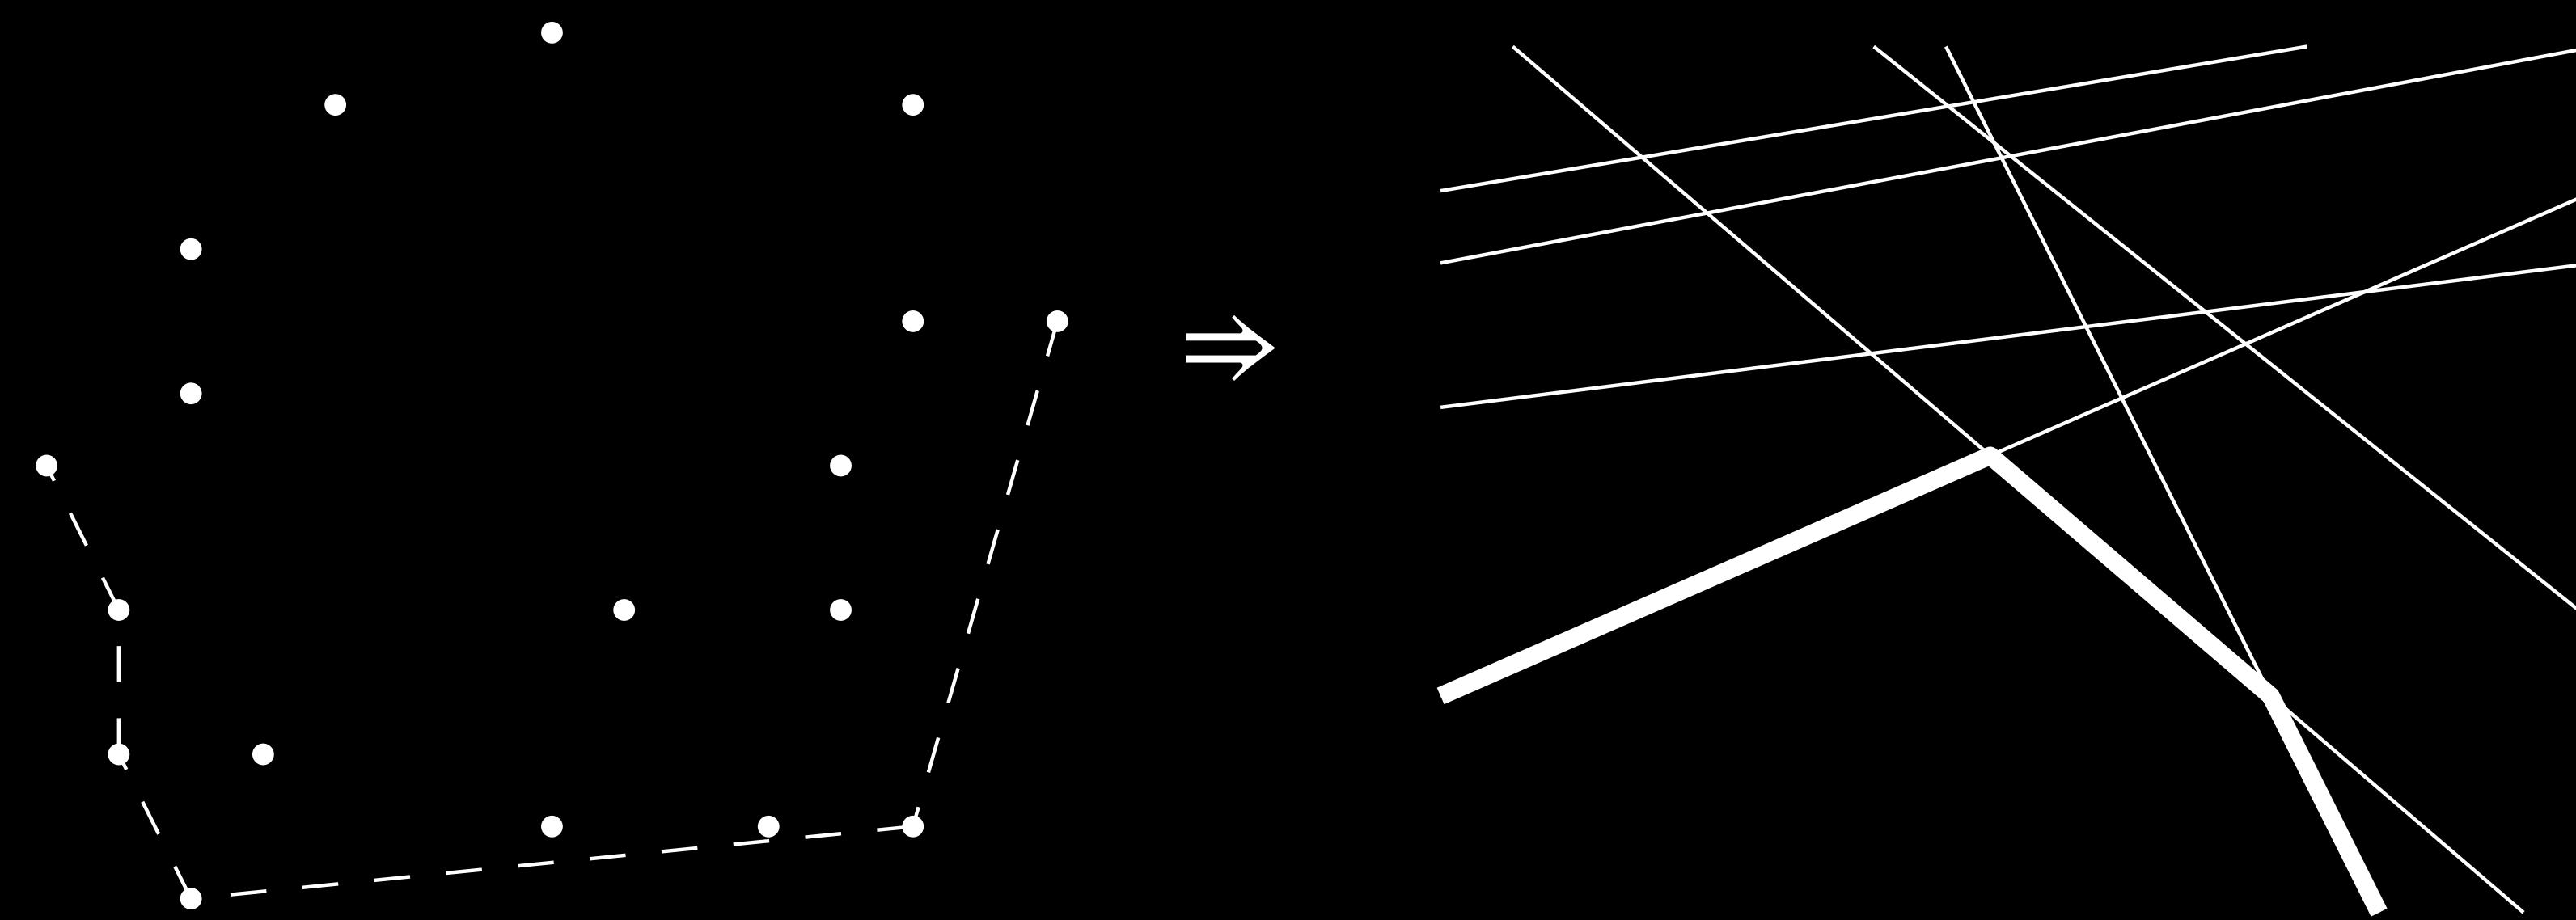
- Transformation from points to hyper-planes and vice-versa



# Introduction (5)

## Dual Relationship

- The lower hull in the primal is the lower envelope in the dual space



# Introduction (6)

## Background

- Frank Staals outlines in [2] an efficient parallel algorithm for computing 3D lower envelopes
- Partially based on T. Chan & E. Chen work [1]

# Motivation and Goal

# Motivation & Goal (1)

## Motivation

- Convex hulls are fundamental to many applications
- Parallel algorithms for computing lower envelopes unlock parallel algorithms for the convex hull problem
- Modern CPUs provide n-fold speedups for embarrassingly-parallel problems
- Unclear if the algorithms and theoretical complexity in [2] are practically achievable

# Motivation & Goal (2)

## Goal

Exploit parallel computing in the context  
of 3D Lower Envelope algorithms

# Motivation & Goal (3)

## Questions

- Are the parallel algorithms in [2] practically feasible?
- What is an efficient way to implement such algorithms?
- How do the algorithms' running times relate to their theoretical complexities?
  - A. As the input size is increased
  - B. As the amount of parallel processors used is increased

# Overview

# Overview (1)

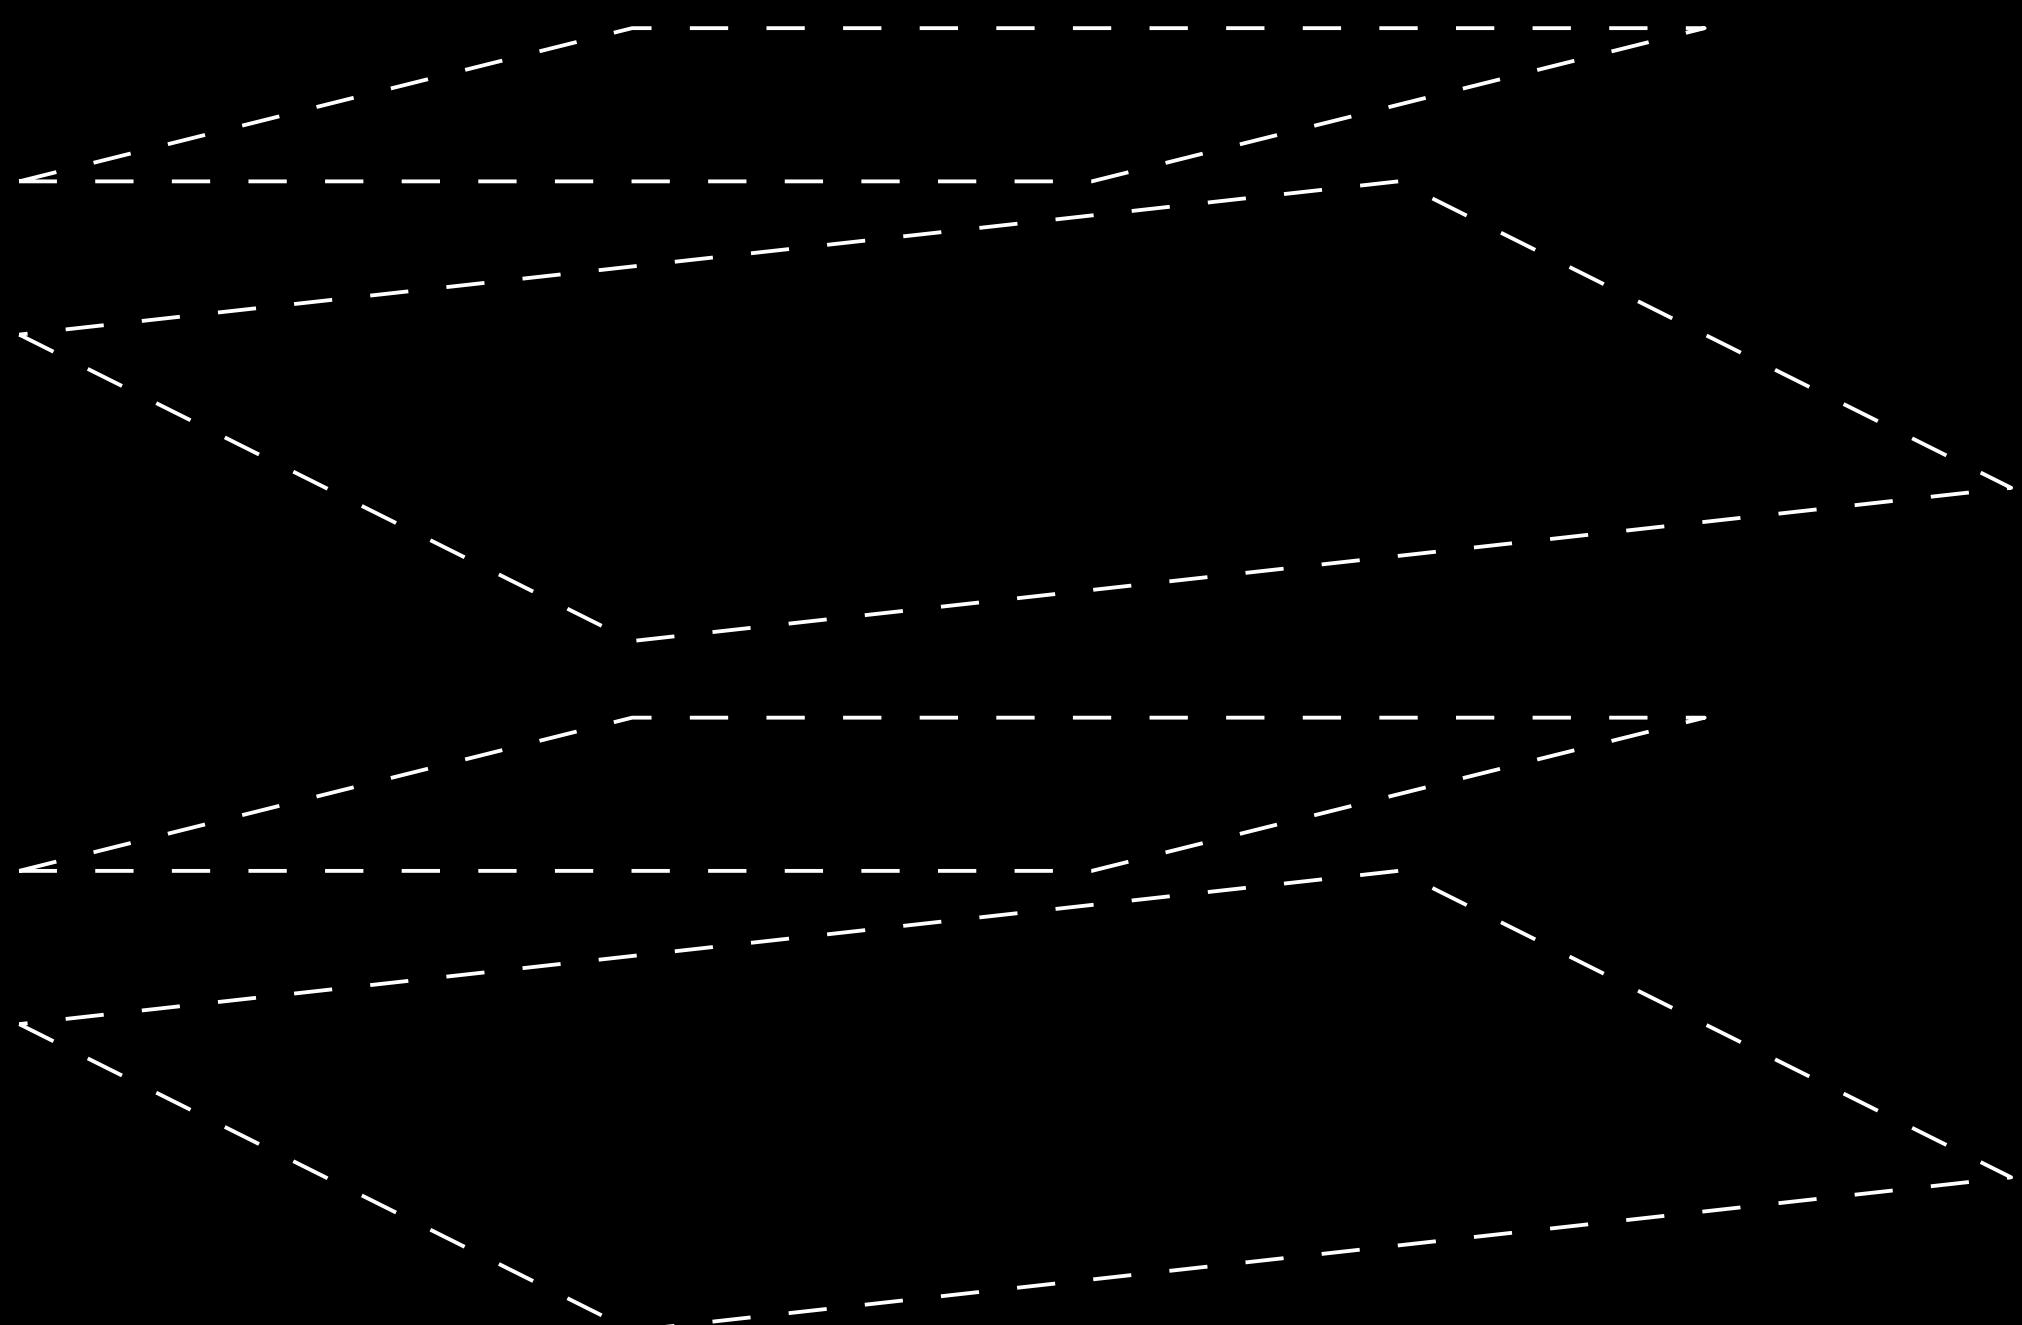
## Efficient $L_0$ computation (main algorithm)

- Efficient divide and conquer, randomised algorithm outlined by F. Staals in [2]
- Given  $n$  input planes and  $p$  processors, the algorithm runs in  $O\left(\frac{n}{p} \log n\right)$

# Overview (2)

## Efficient $L_0$ computation (main algorithm)

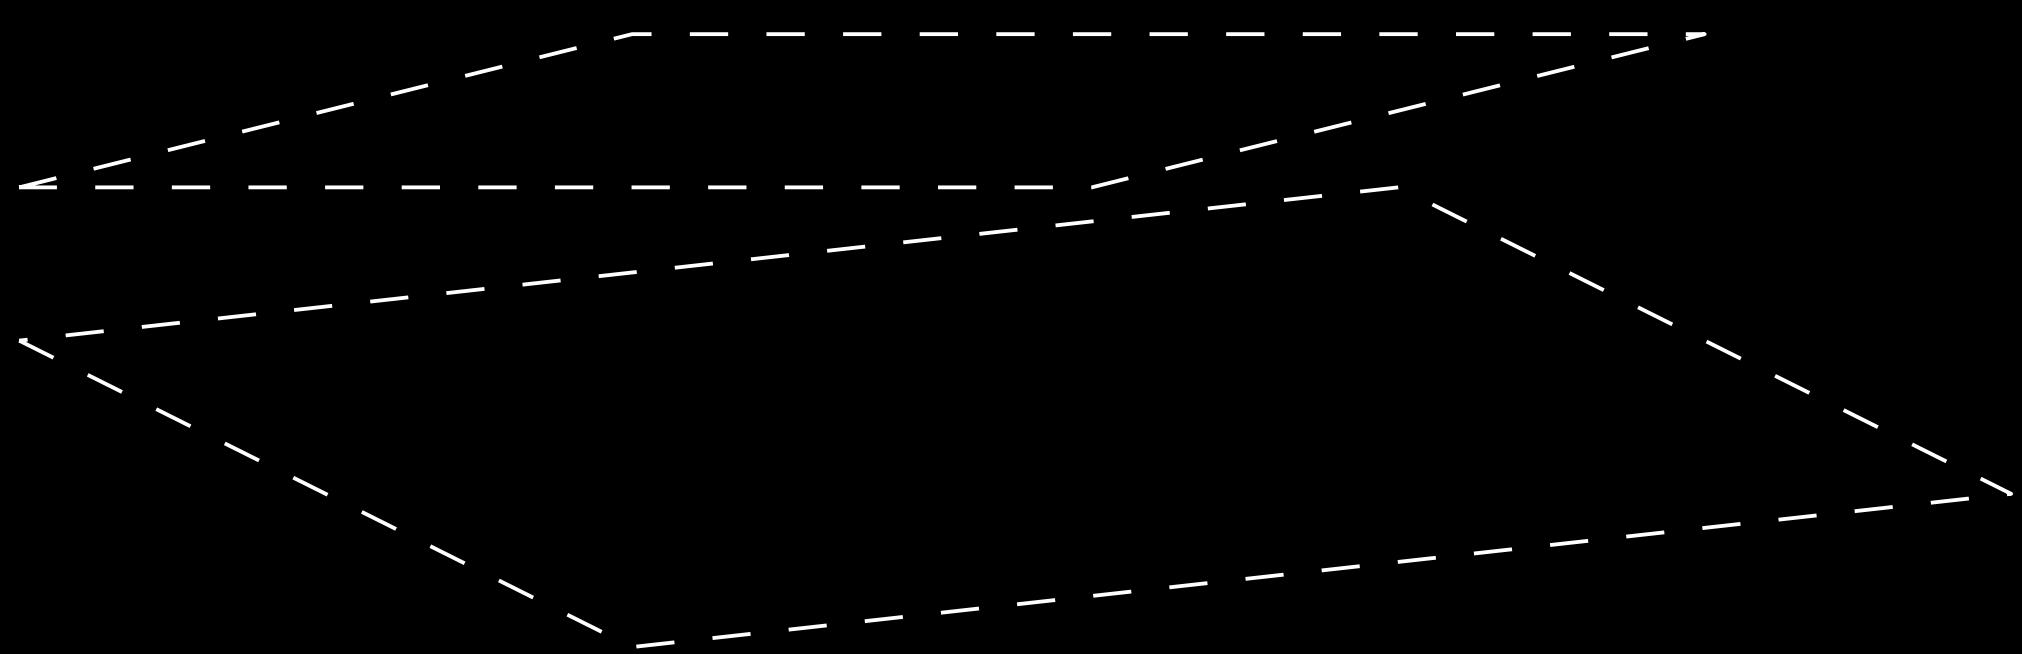
- Given a set of  $n$  planes  $H$ ,  
compute  $L_0(H)$



# Overview (2)

## Efficient $L_0$ computation (main algorithm)

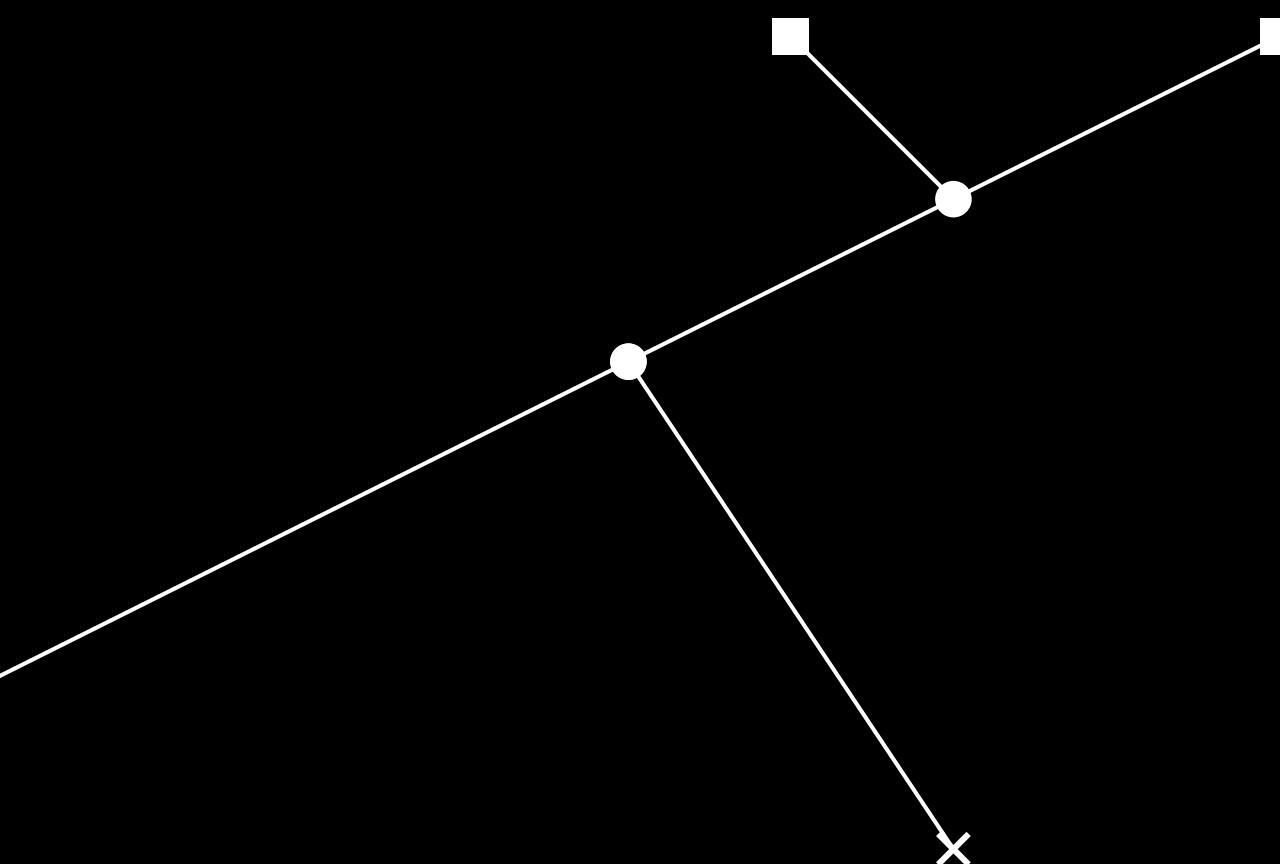
- Given a set of  $n$  planes  $H$ , compute  $L_0(H)$
- Pick sample  $S \subset H$  using double sampling scheme described in [2]



# Overview (2)

## Efficient $L_0$ computation (main algorithm)

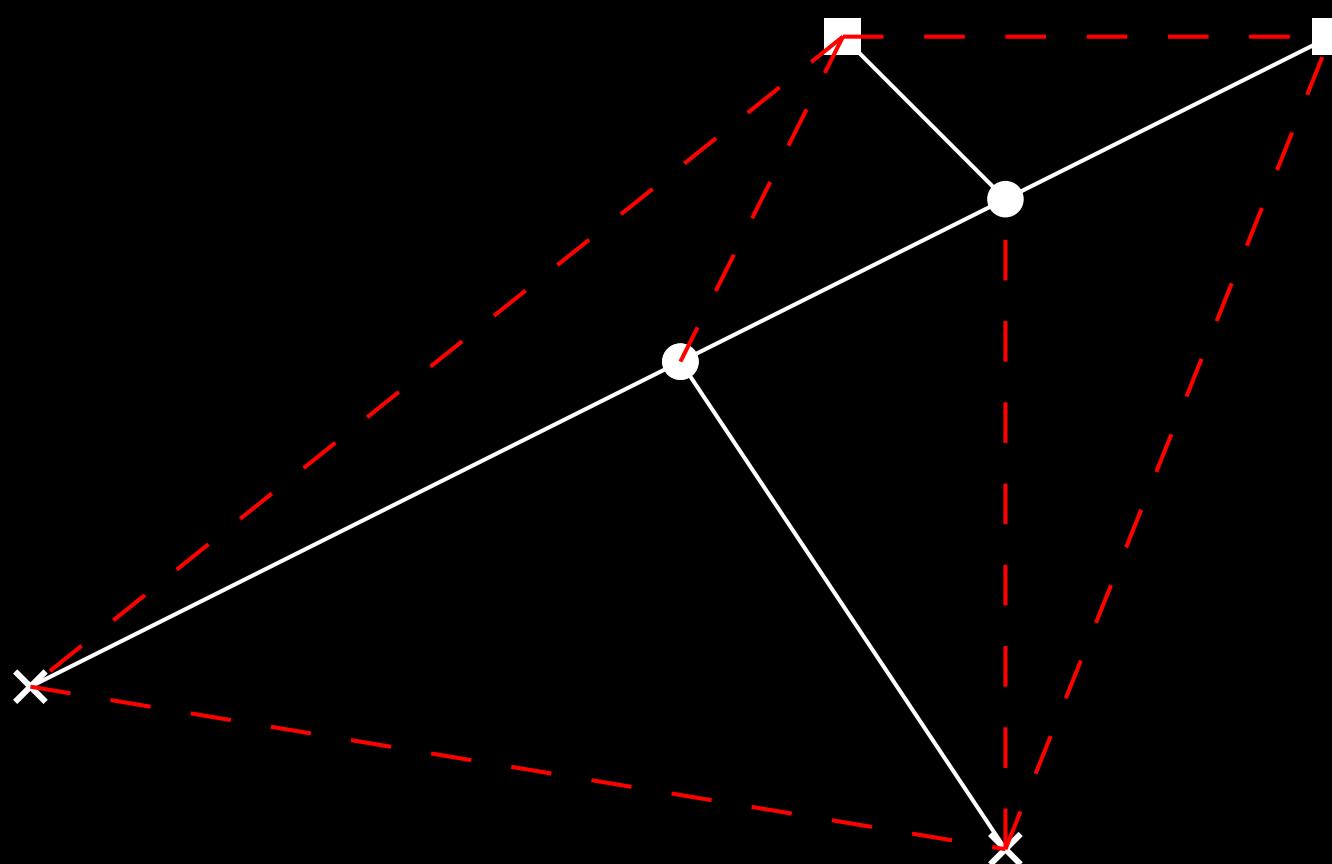
- Given a set of  $n$  planes  $H$ , compute  $L_0(H)$
- Pick sample  $S \subset H$  using double sampling scheme described in [2]
- Compute  $L_0(S)$  and triangulate it



# Overview (2)

## Efficient $L_0$ computation (main algorithm)

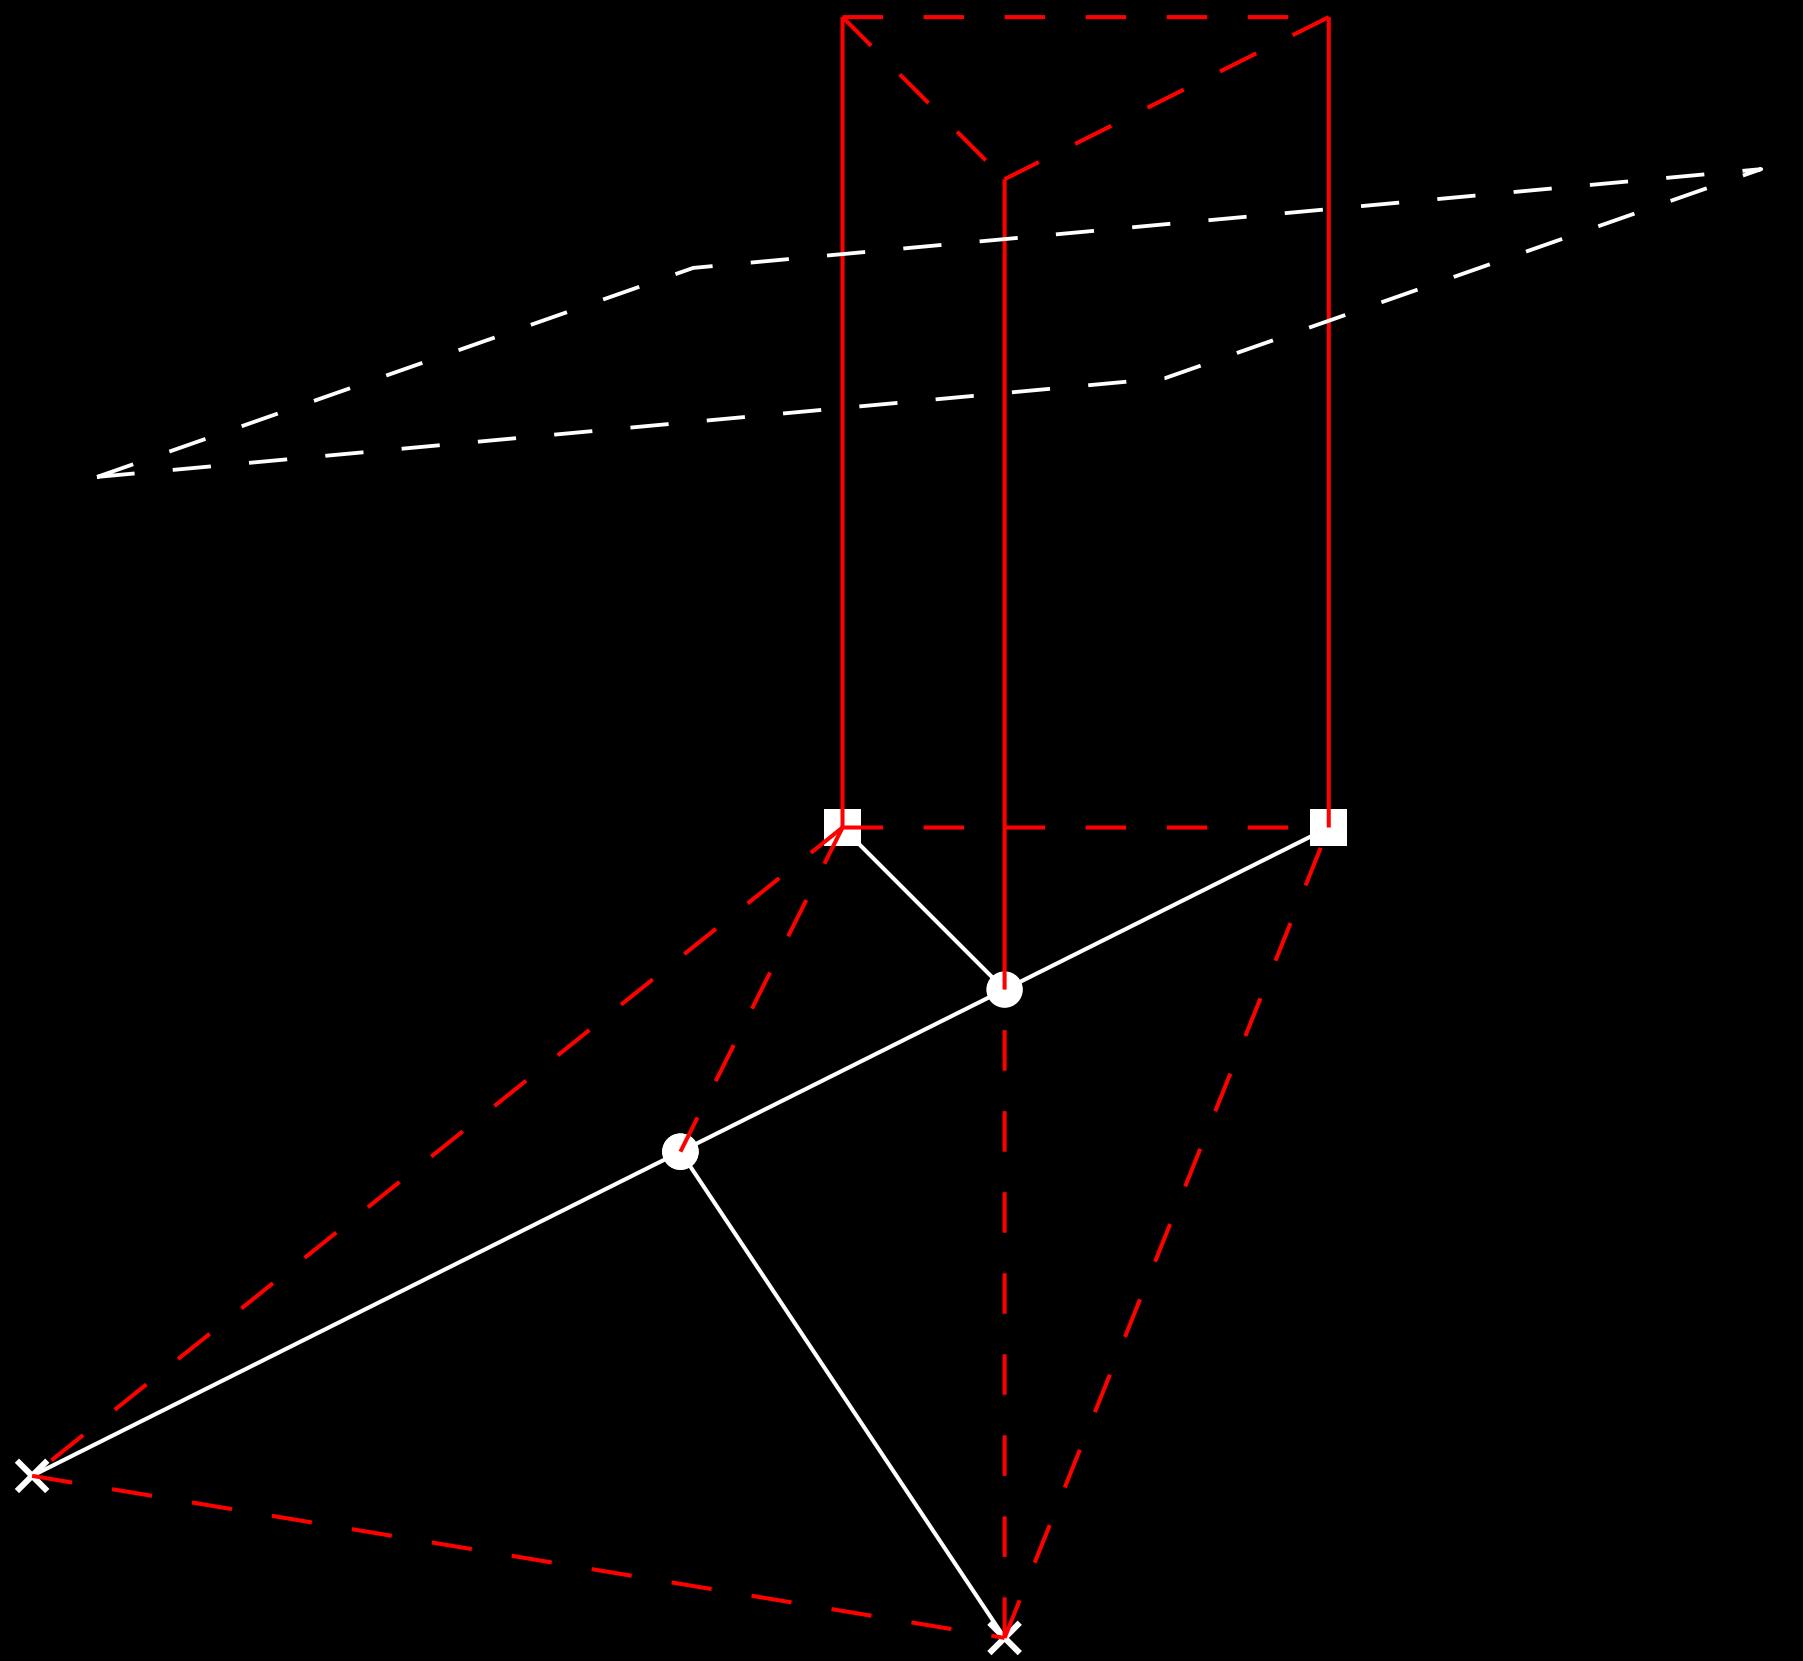
- Given a set of  $n$  planes  $H$ , compute  $L_0(H)$
- Pick sample  $S \subset H$  using double sampling scheme described in [2]
- Compute  $L_0(S)$  and triangulate it



# Overview (2)

## Efficient $L_0$ computation (main algorithm)

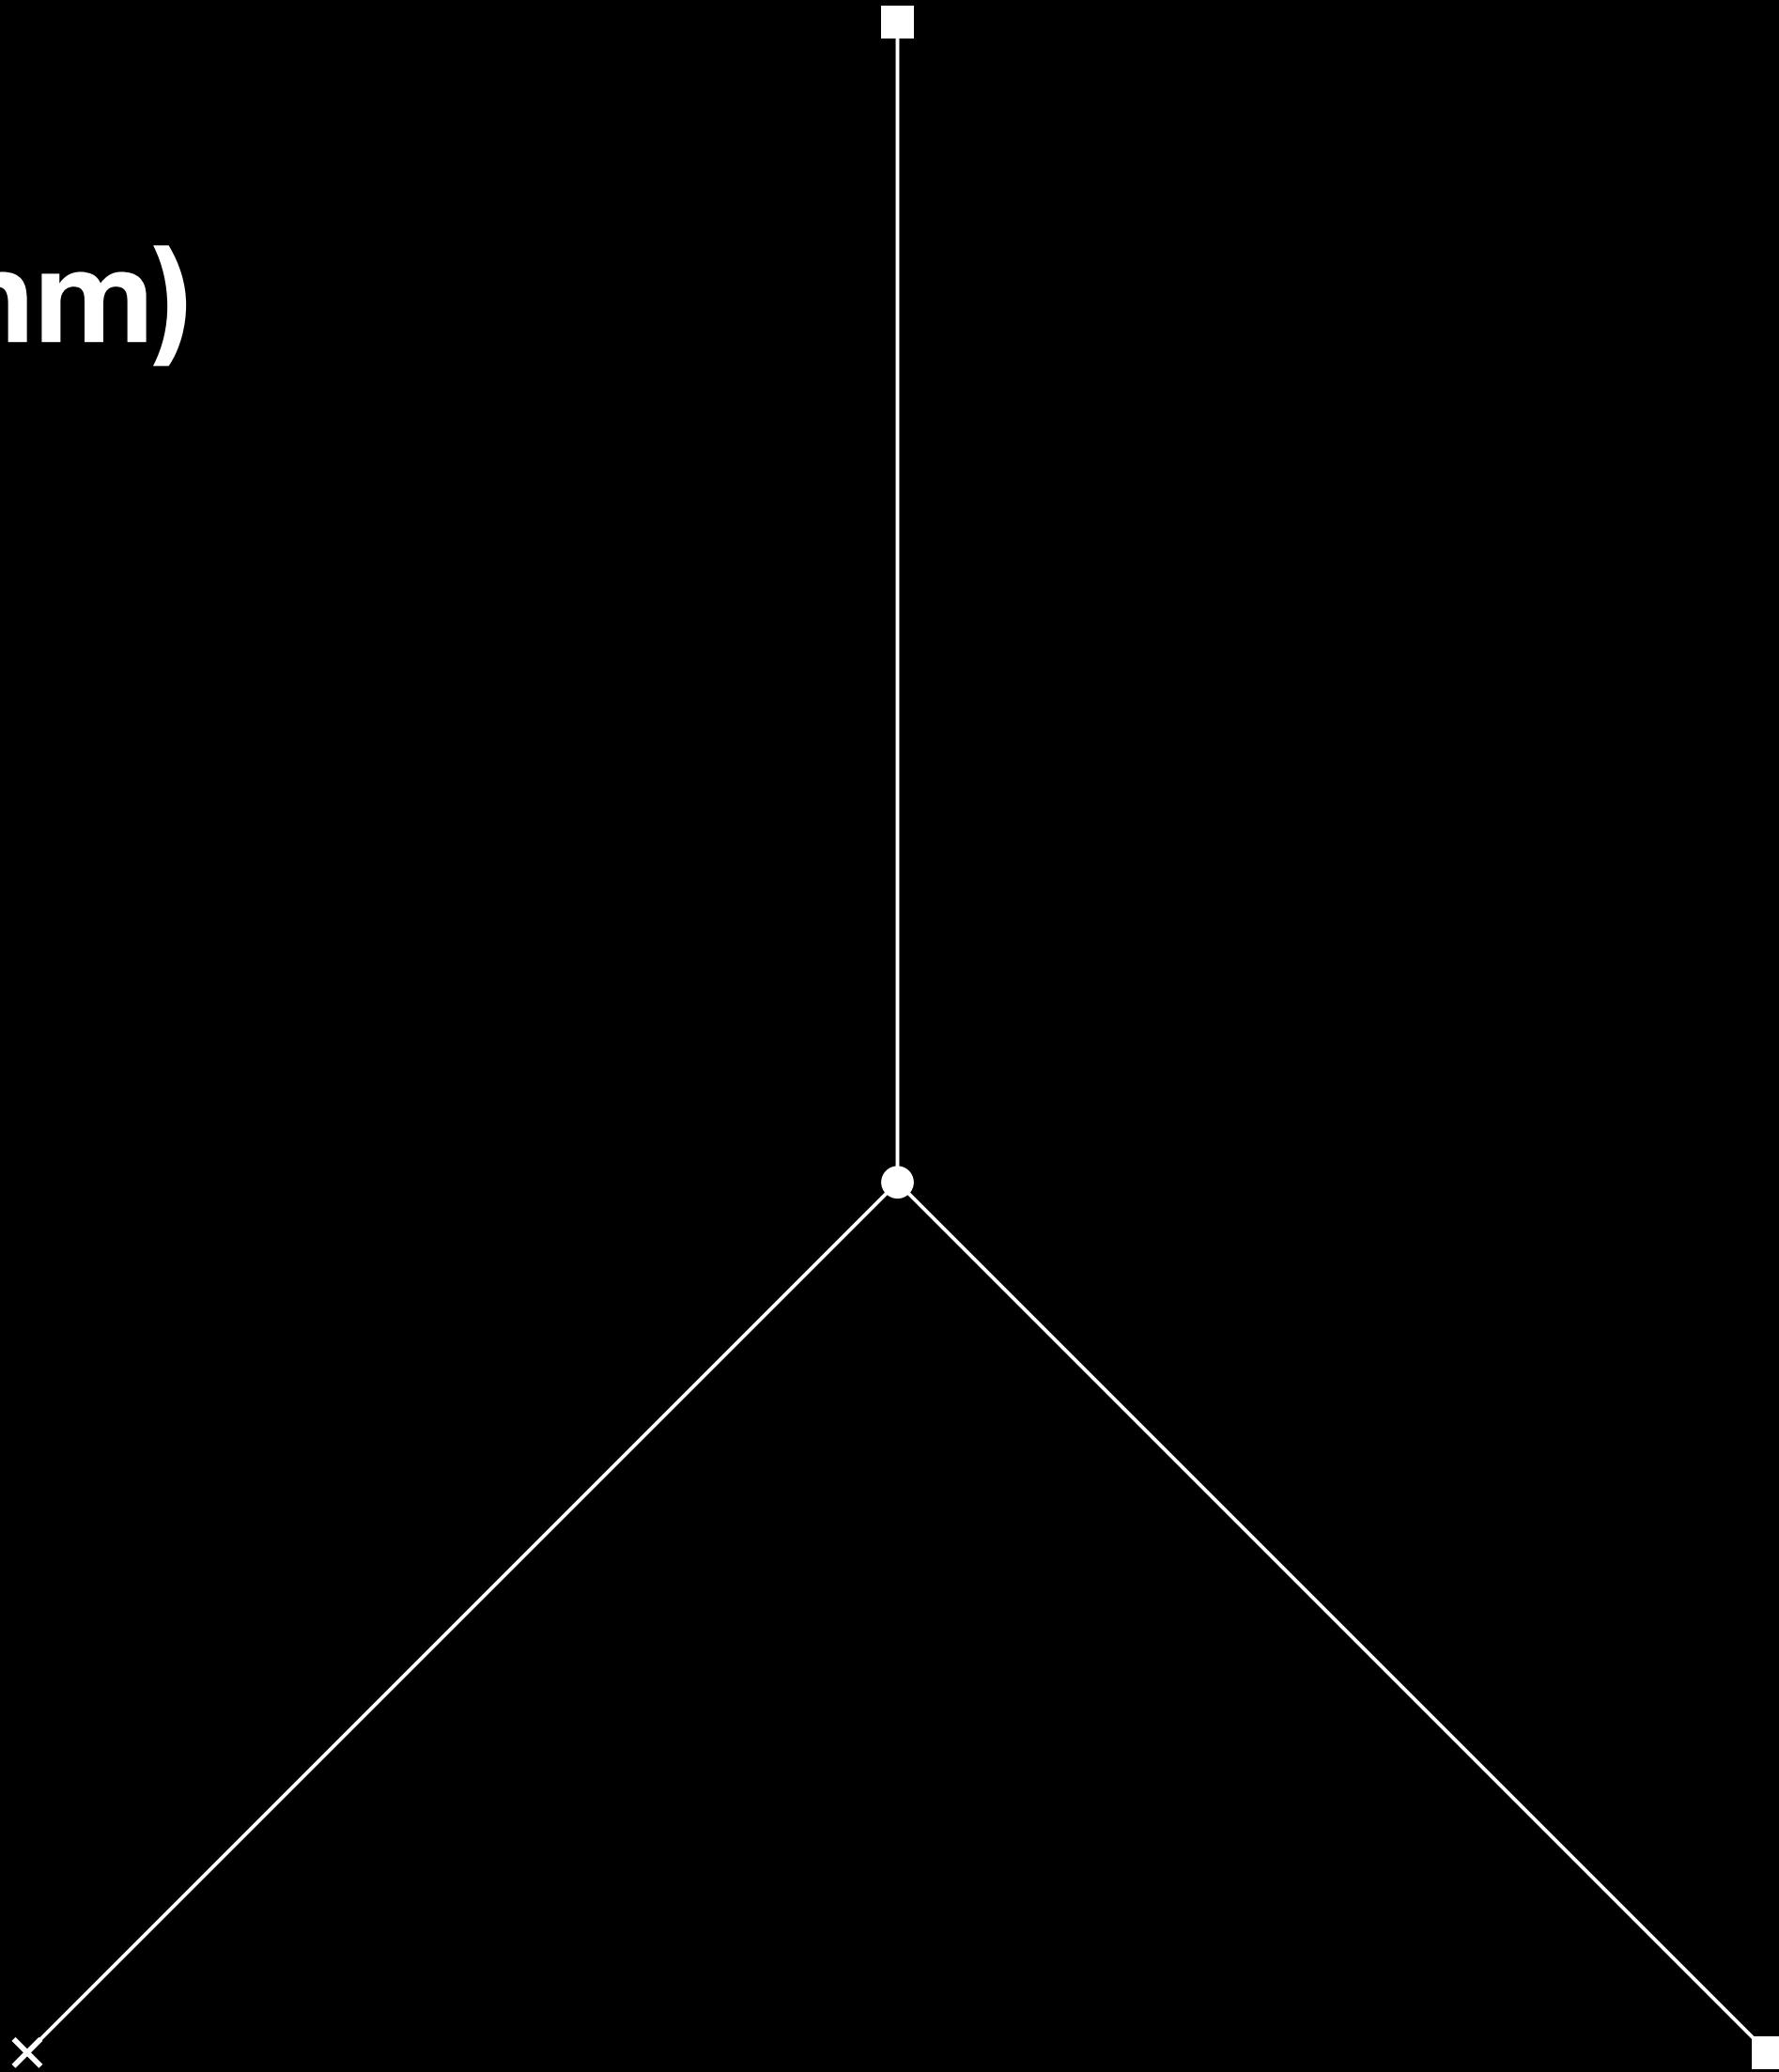
- Given a set of  $n$  planes  $H$ , compute  $L_0(H)$
- Pick sample  $S \subset H$  using double sampling scheme described in [2]
- Compute  $L_0(S)$  and triangulate it
- Compute the conflict lists and solve all the sub-problems for each prism in  $\Lambda_0(S)$



# Overview (3)

## Efficient $L_0$ computation (main algorithm)

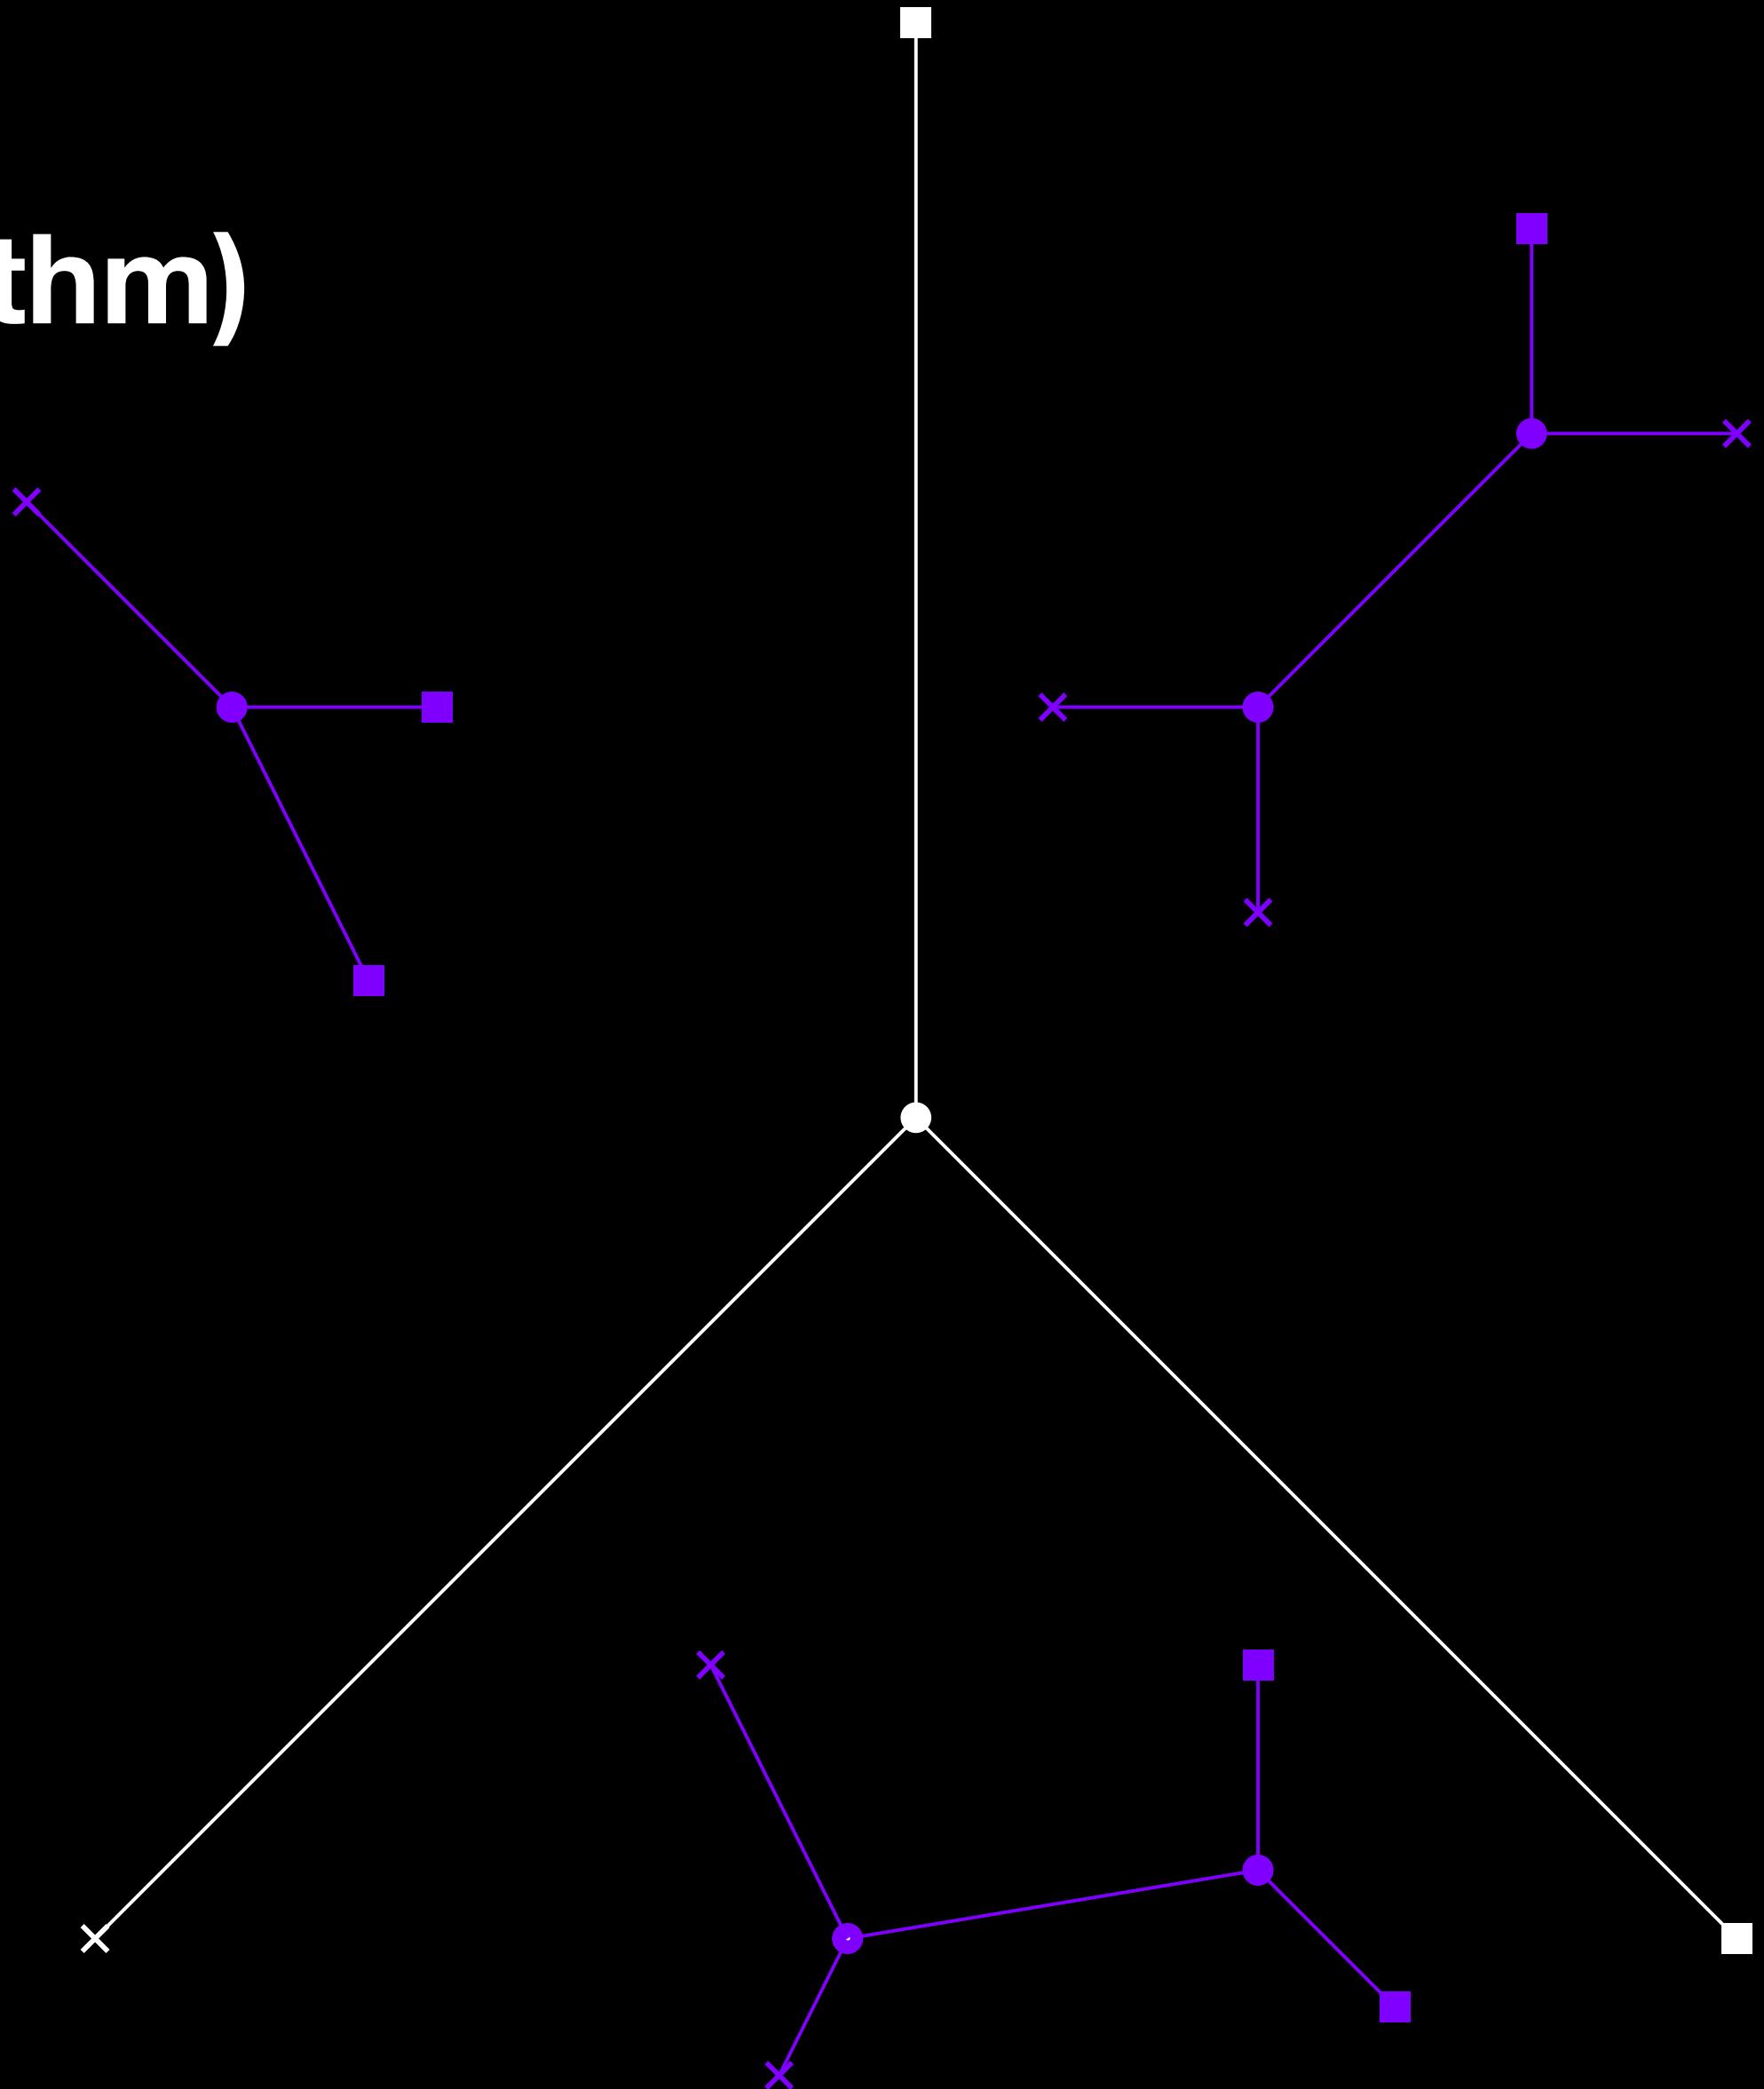
- Compute outer-most lower envelope



# Overview (3)

## Efficient $L_0$ computation (main algorithm)

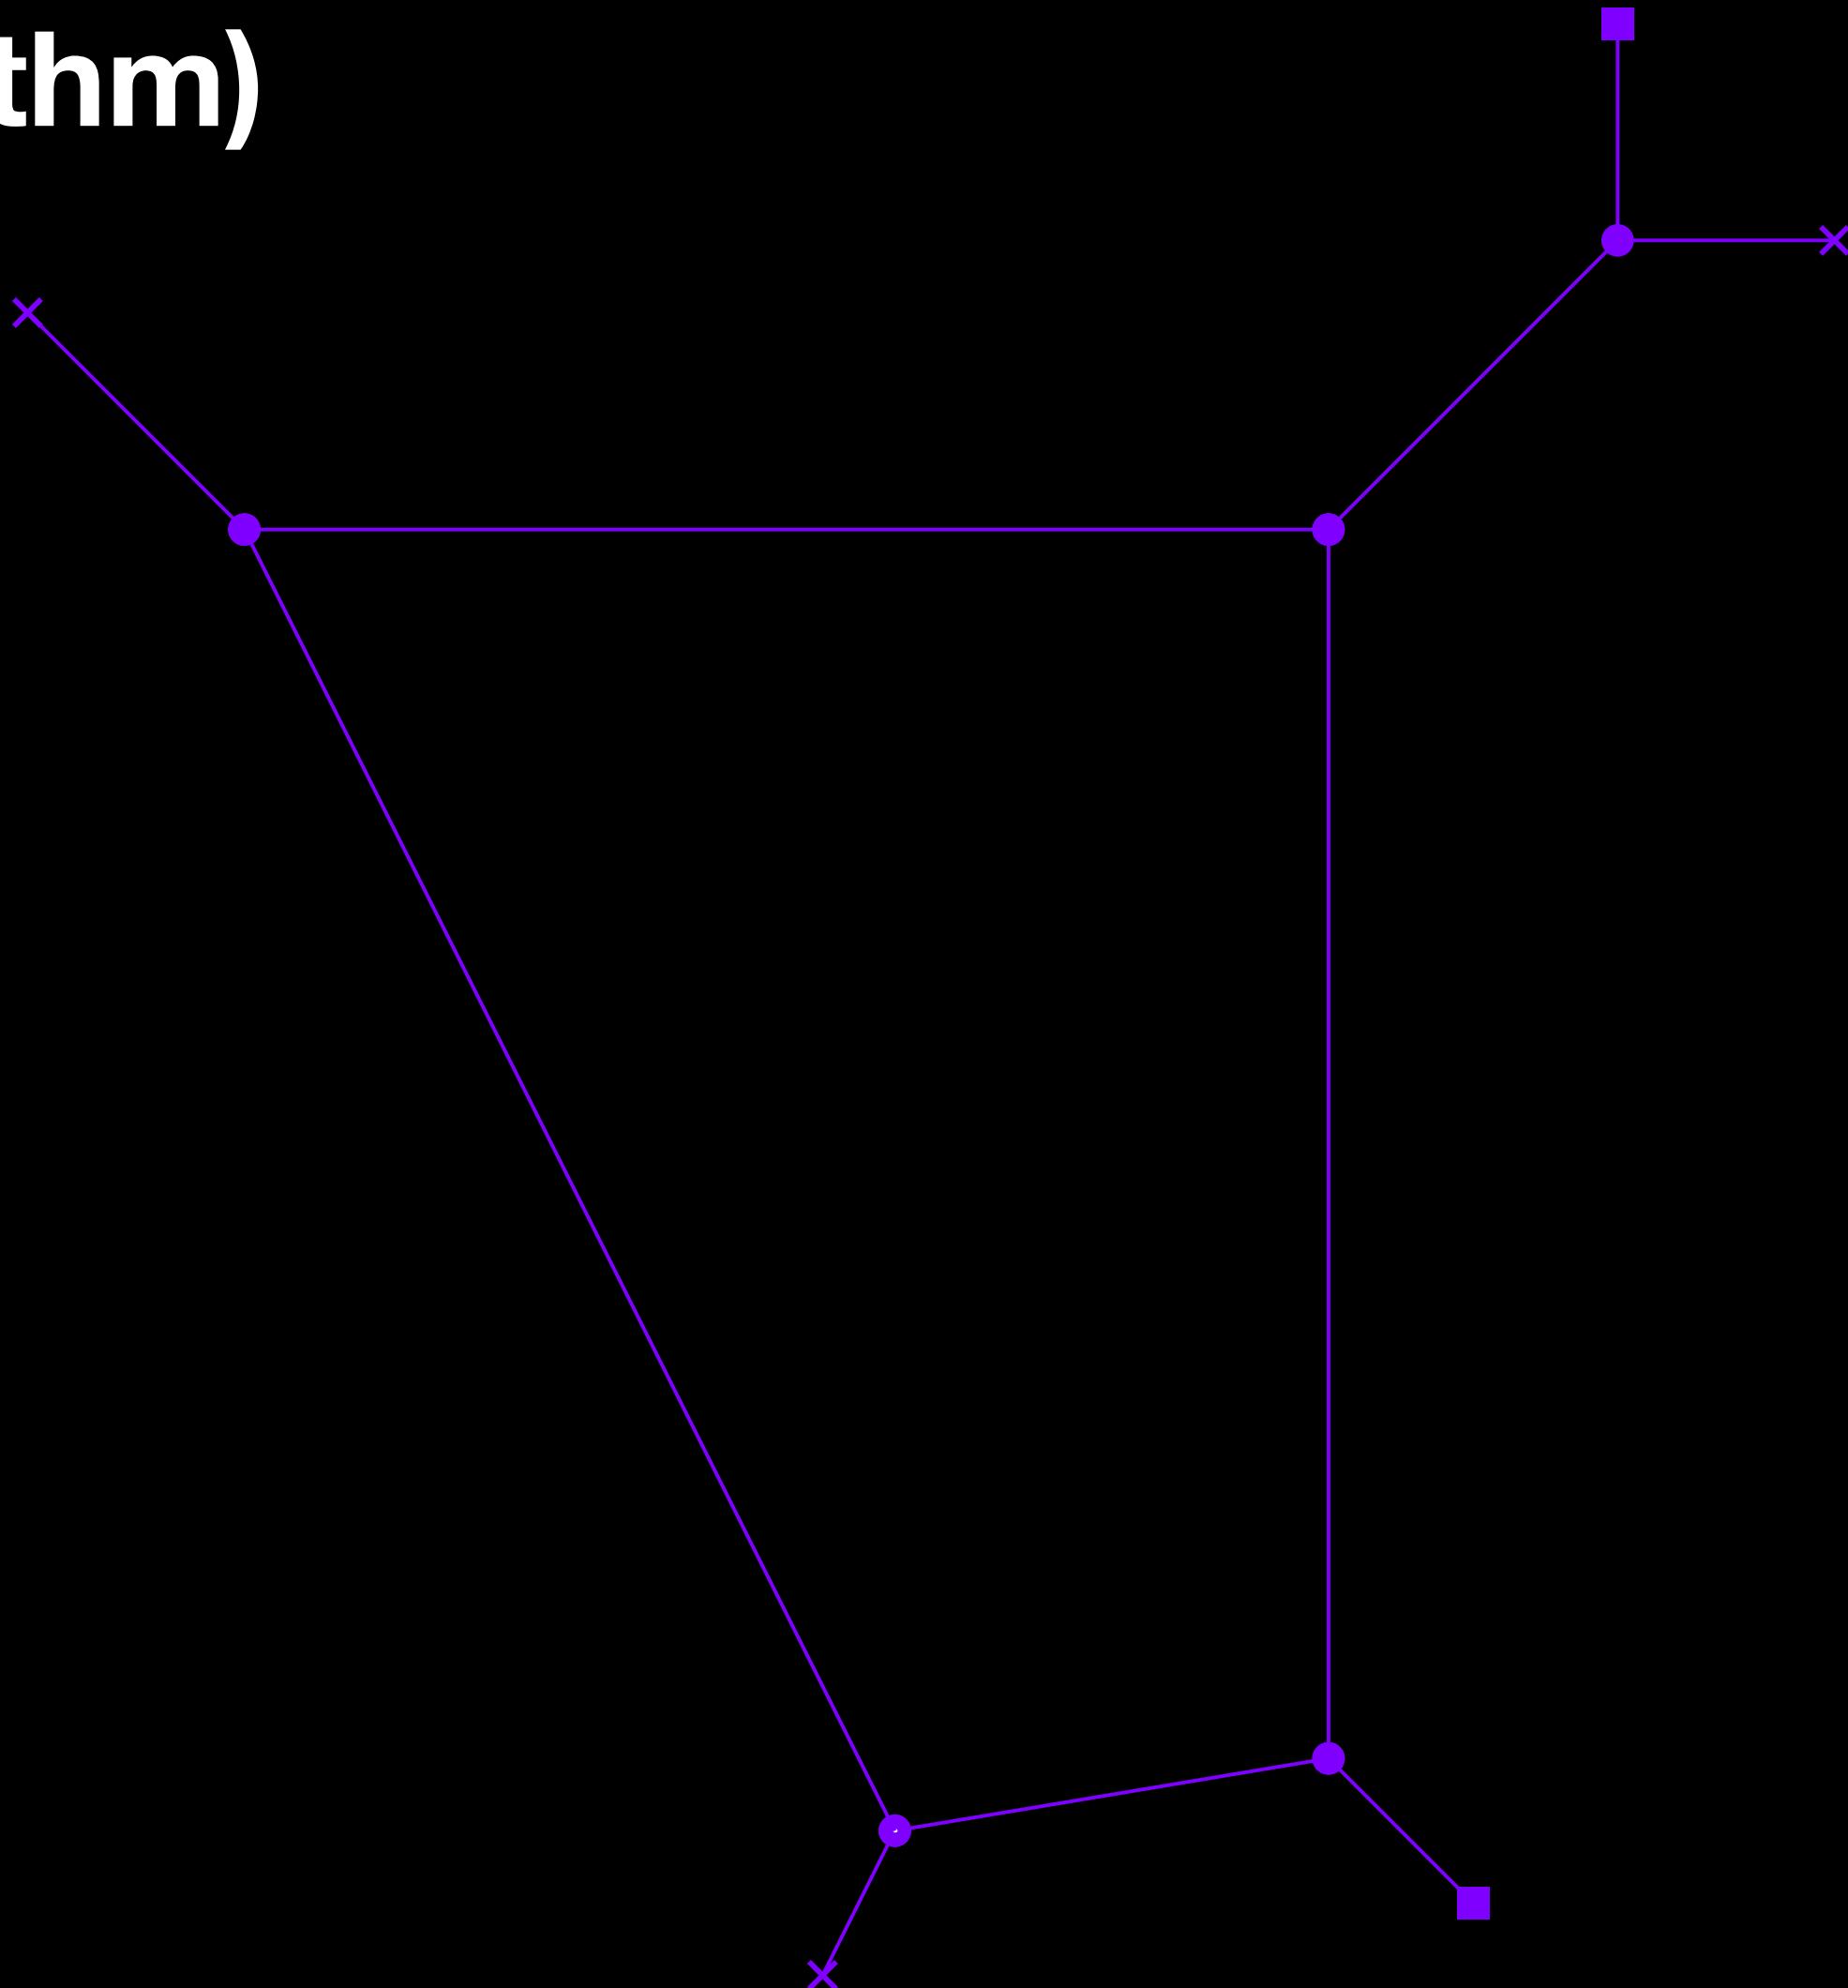
- Compute outer-most lower envelope
- For each prism, compute the conflict lists and find its lower envelope



# Overview (3)

## Efficient $L_0$ computation (main algorithm)

- Compute outer-most lower envelope
- For each prism, compute the conflict lists and find its lower envelope
- Merge adjacent lower envelopes along common edge



# Overview (4)

## Efficient $L_0$ computation (main algorithm)

- Various supporting algorithms are required:
  1. Lower envelope algorithm for small sized  $H$
  2. Lower envelope triangulation algorithm
  3. Conflict lists algorithm
  4. Lower envelopes adjacency and merge algorithm

# Overview (5)

## Efficient $L_0$ computation (main algorithm)

- Focus of this presentation:

**1.Lower envelope algorithm for small sized  $H$  (brute force)**

2.Lower envelope triangulation algorithm

**3.Conflict lists algorithm (batch point location)**

4.Lower envelopes adjacency and merge algorithm

# Overview (6)

## Problem definitions

### 1. The Brute Force algorithm

- Given a set of planes  $H$ , compute  $L_0(H)$  using algorithms from [2]

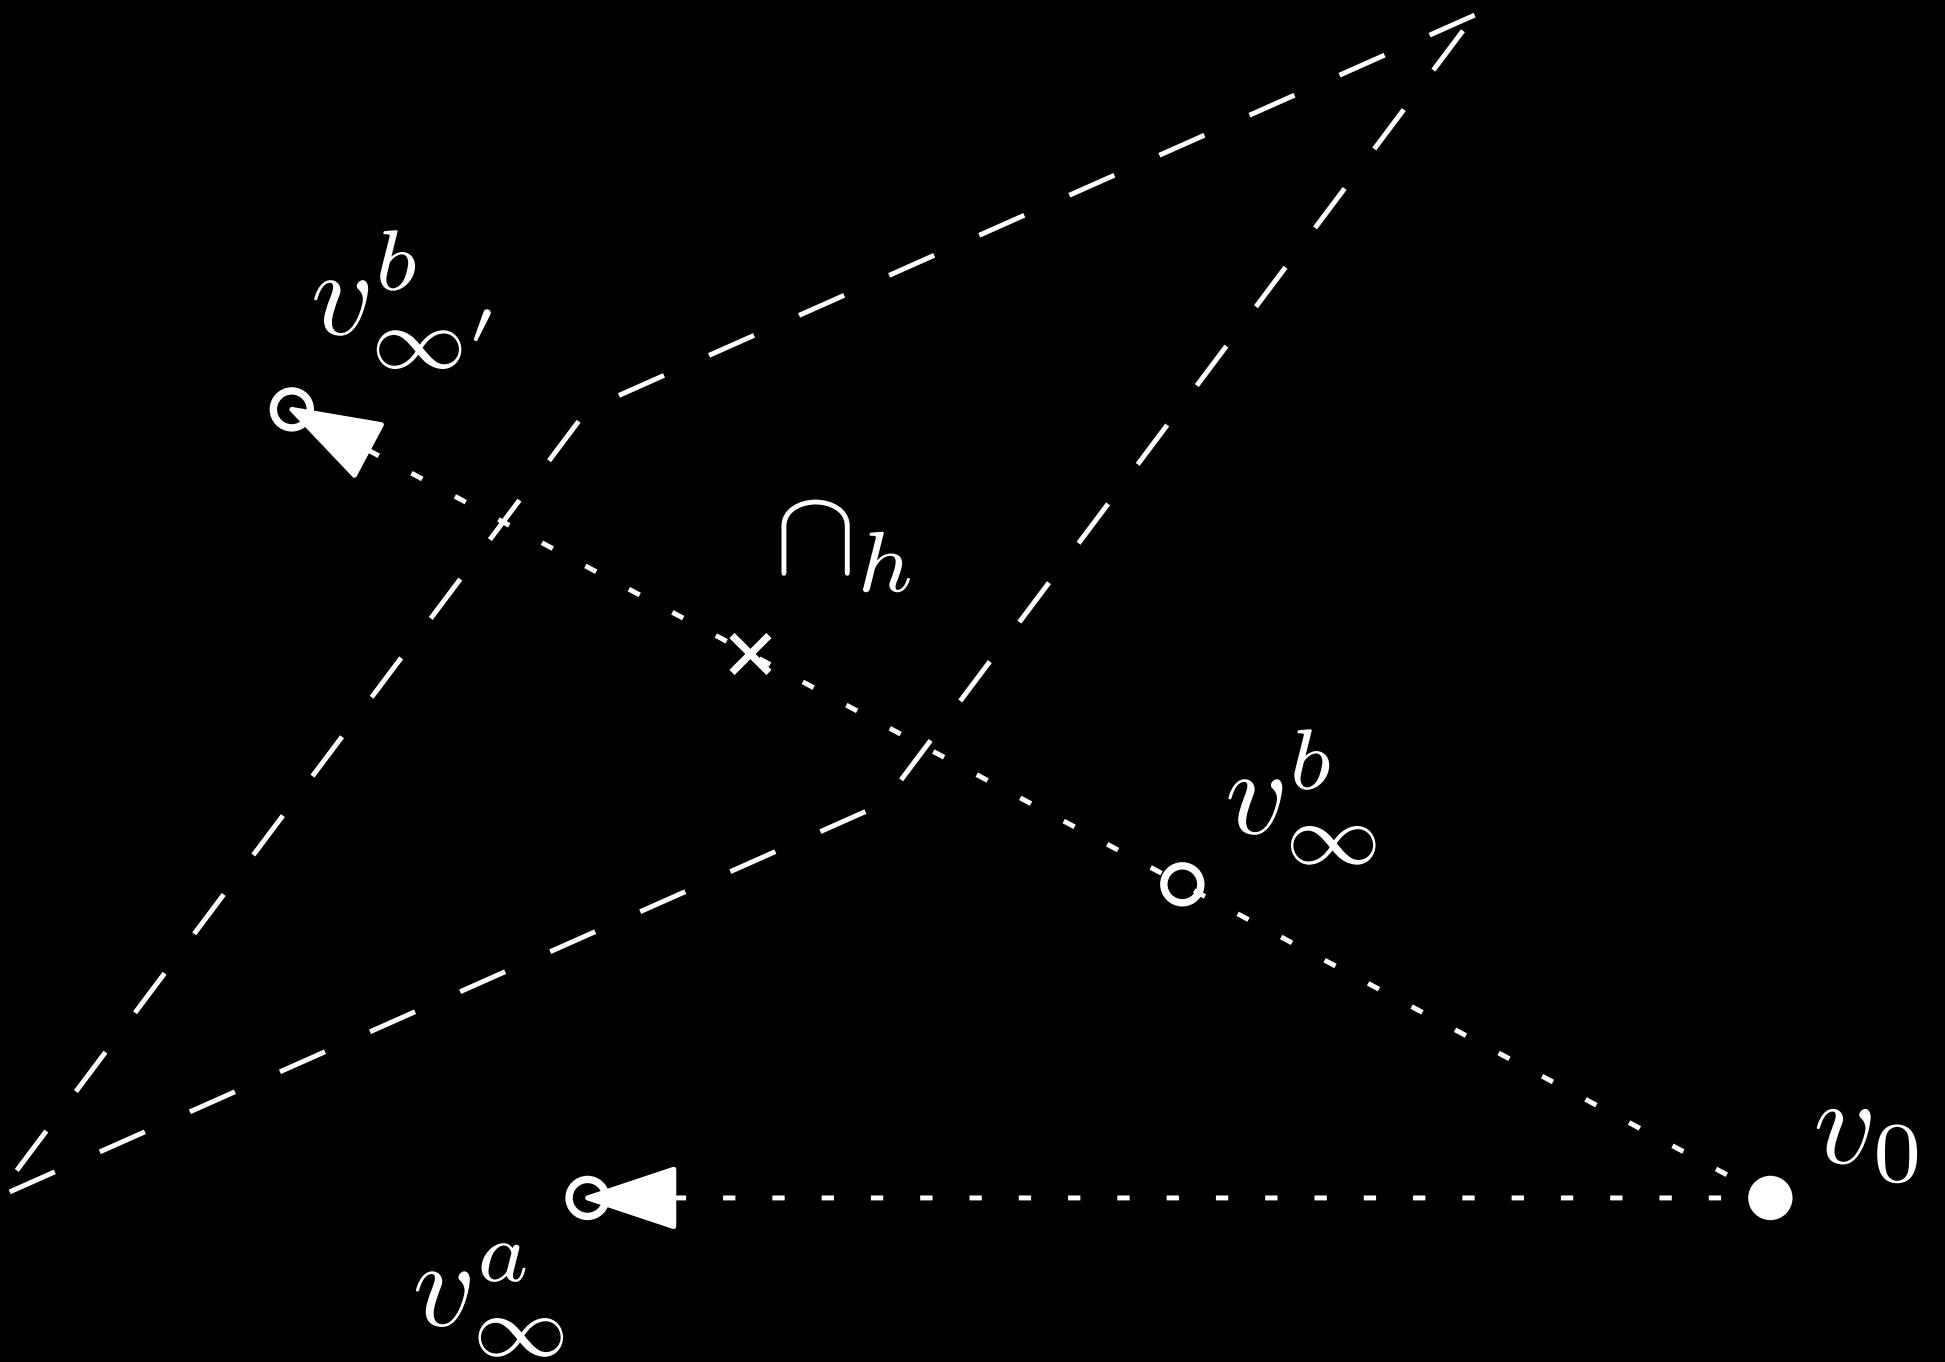
### 3. The Batch Point Location algorithm

- Given a set of planes  $H$  and a set of points  $Q$ , compute the z-sorted sequence of planes above each  $q \in Q$

# Overview (7)

## Limitations

- Number of planes required for the double sampling scheme to work practically unfeasible
- Vertices at infinity may require special treatment
- Incomplete implementation

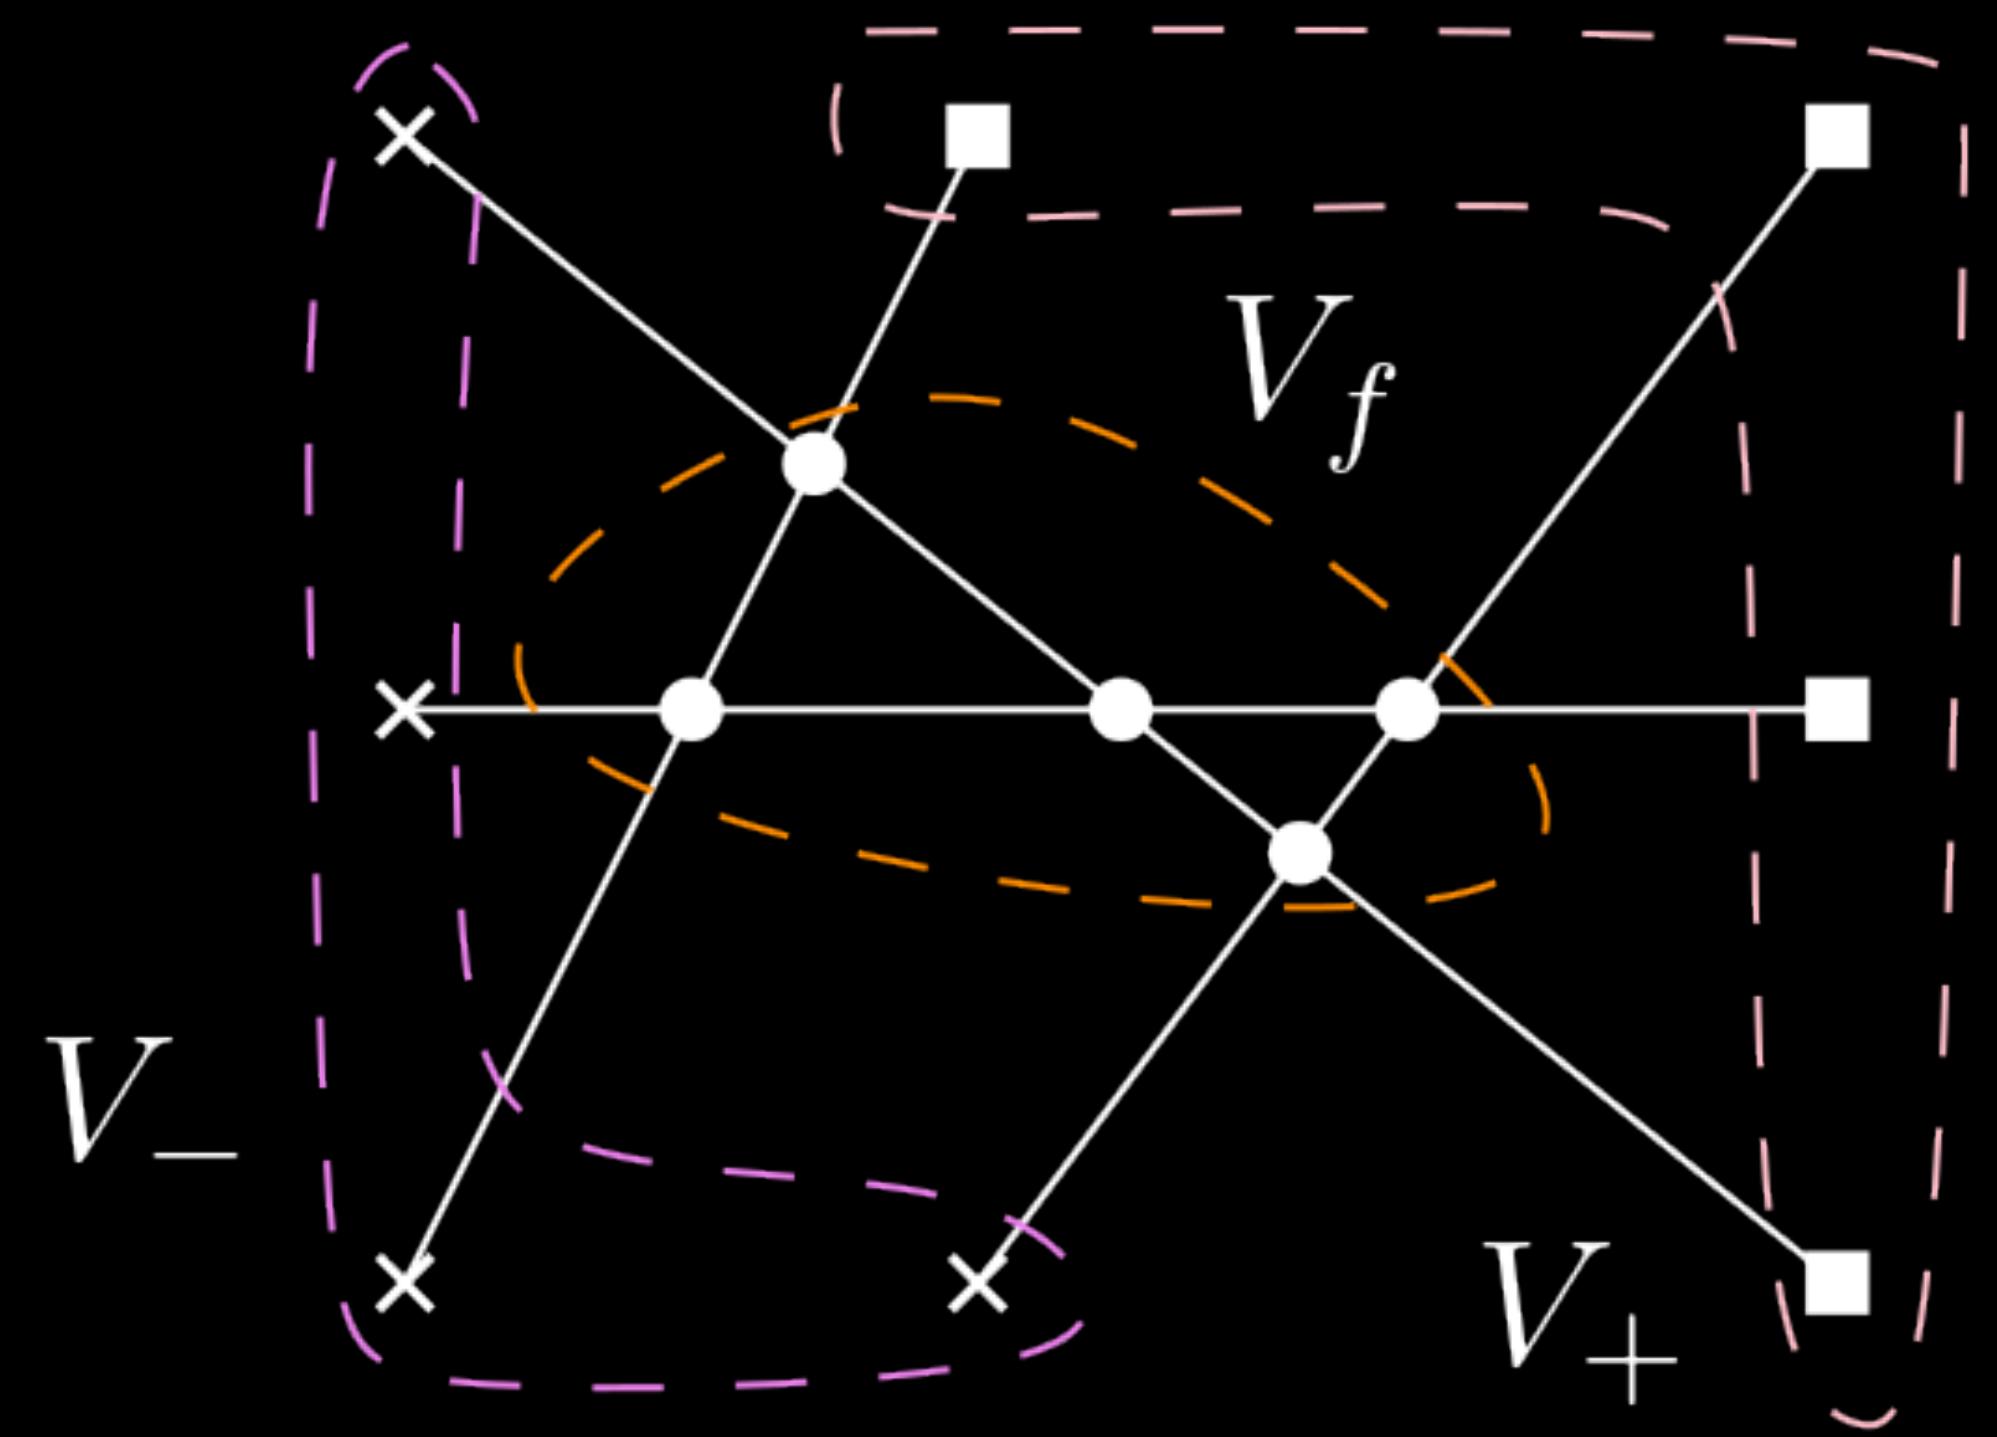


# Algorithms

# Algorithms (1)

## Brute Force $L_0$

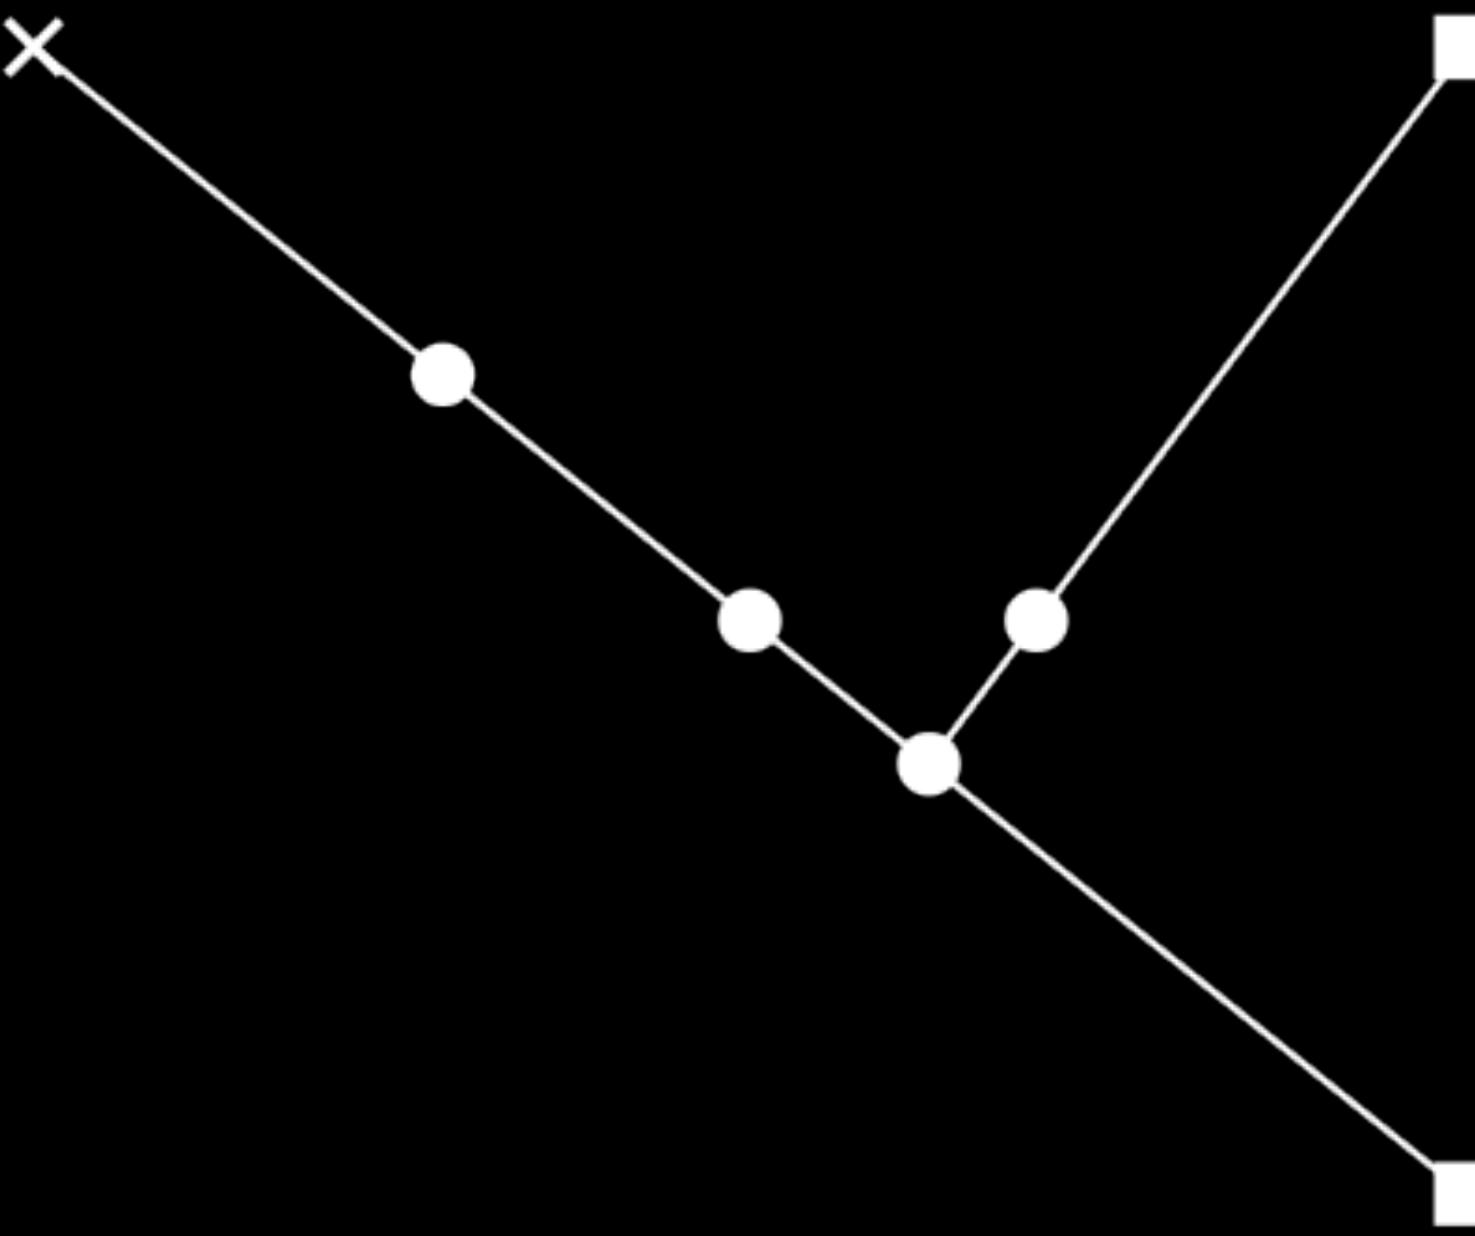
- Given a set of  $n$  planes  $H$ , compute  $L_0(H)$
- Find the set of intersection lines generated by  $H$
- Compute intersection and infinity vertices



# Algorithms (1)

## Brute Force $L_0$

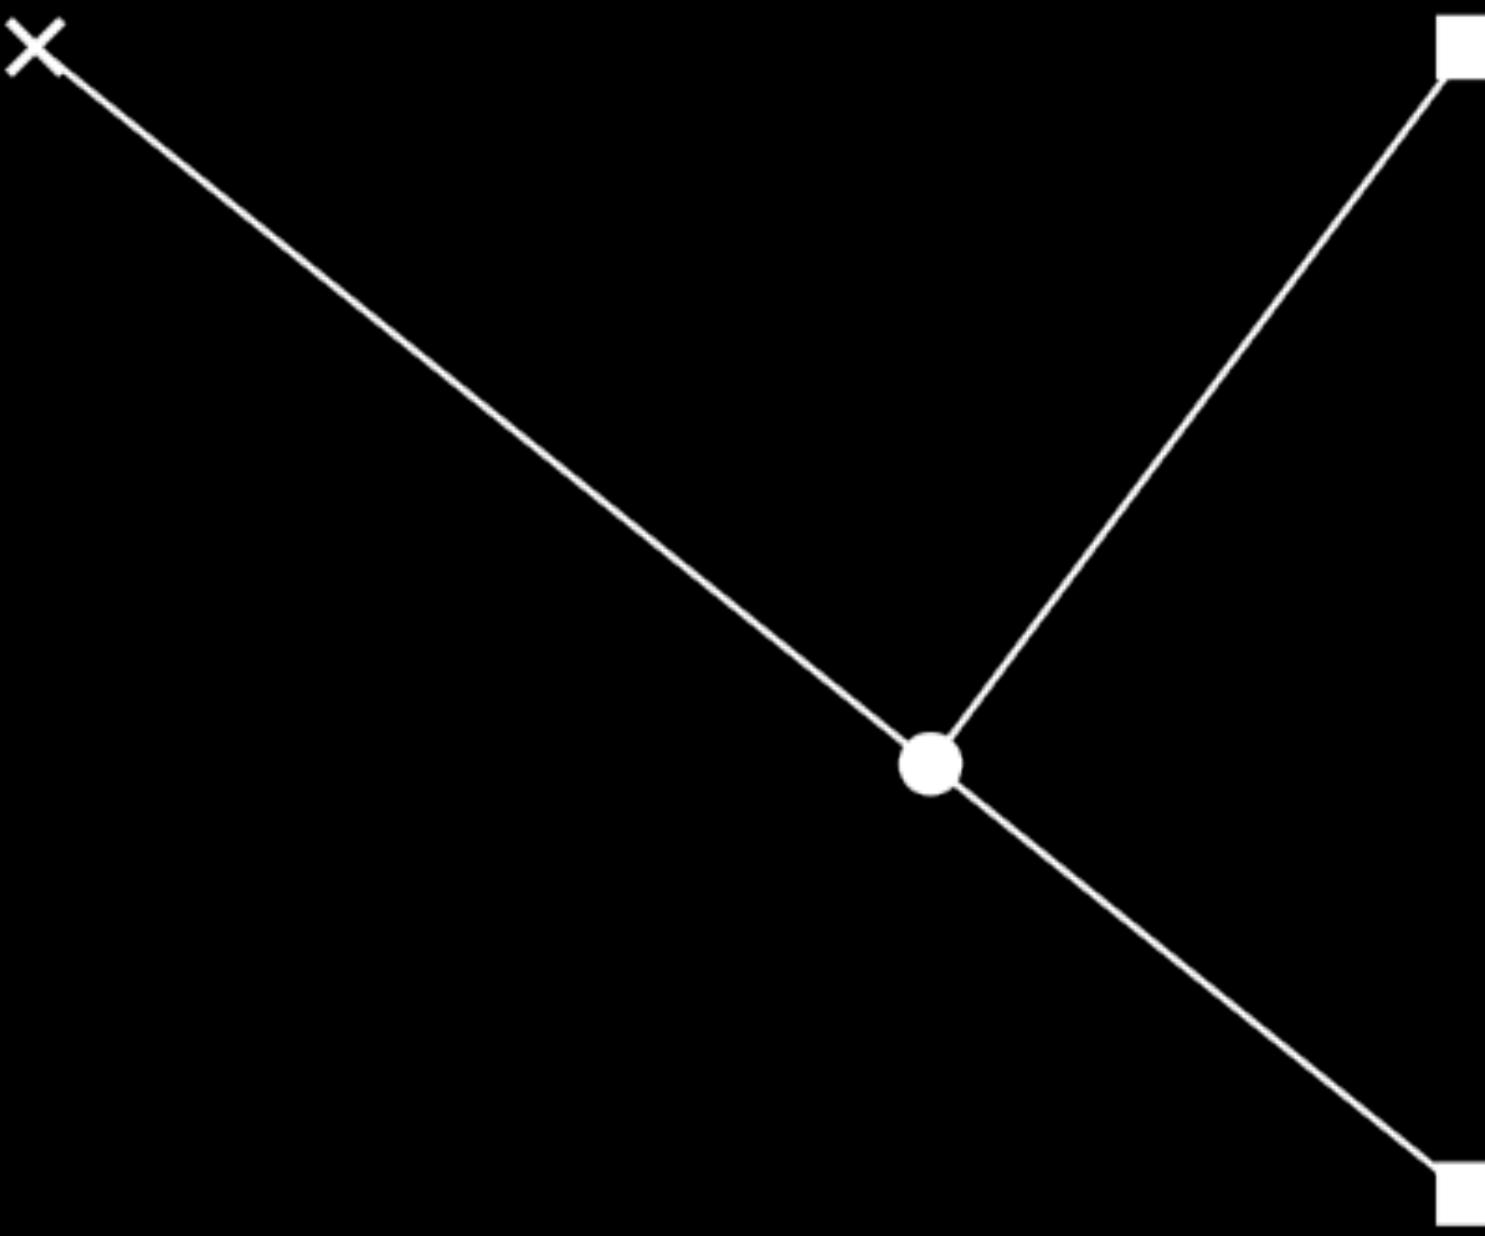
- Given a set of  $n$  planes  $H$ , compute  $L_0(H)$
- Find the set of intersection lines generated by  $H$
- Compute intersection and infinity vertices
- Filter out invalid segments



# Algorithms (1)

## Brute Force $L_0$

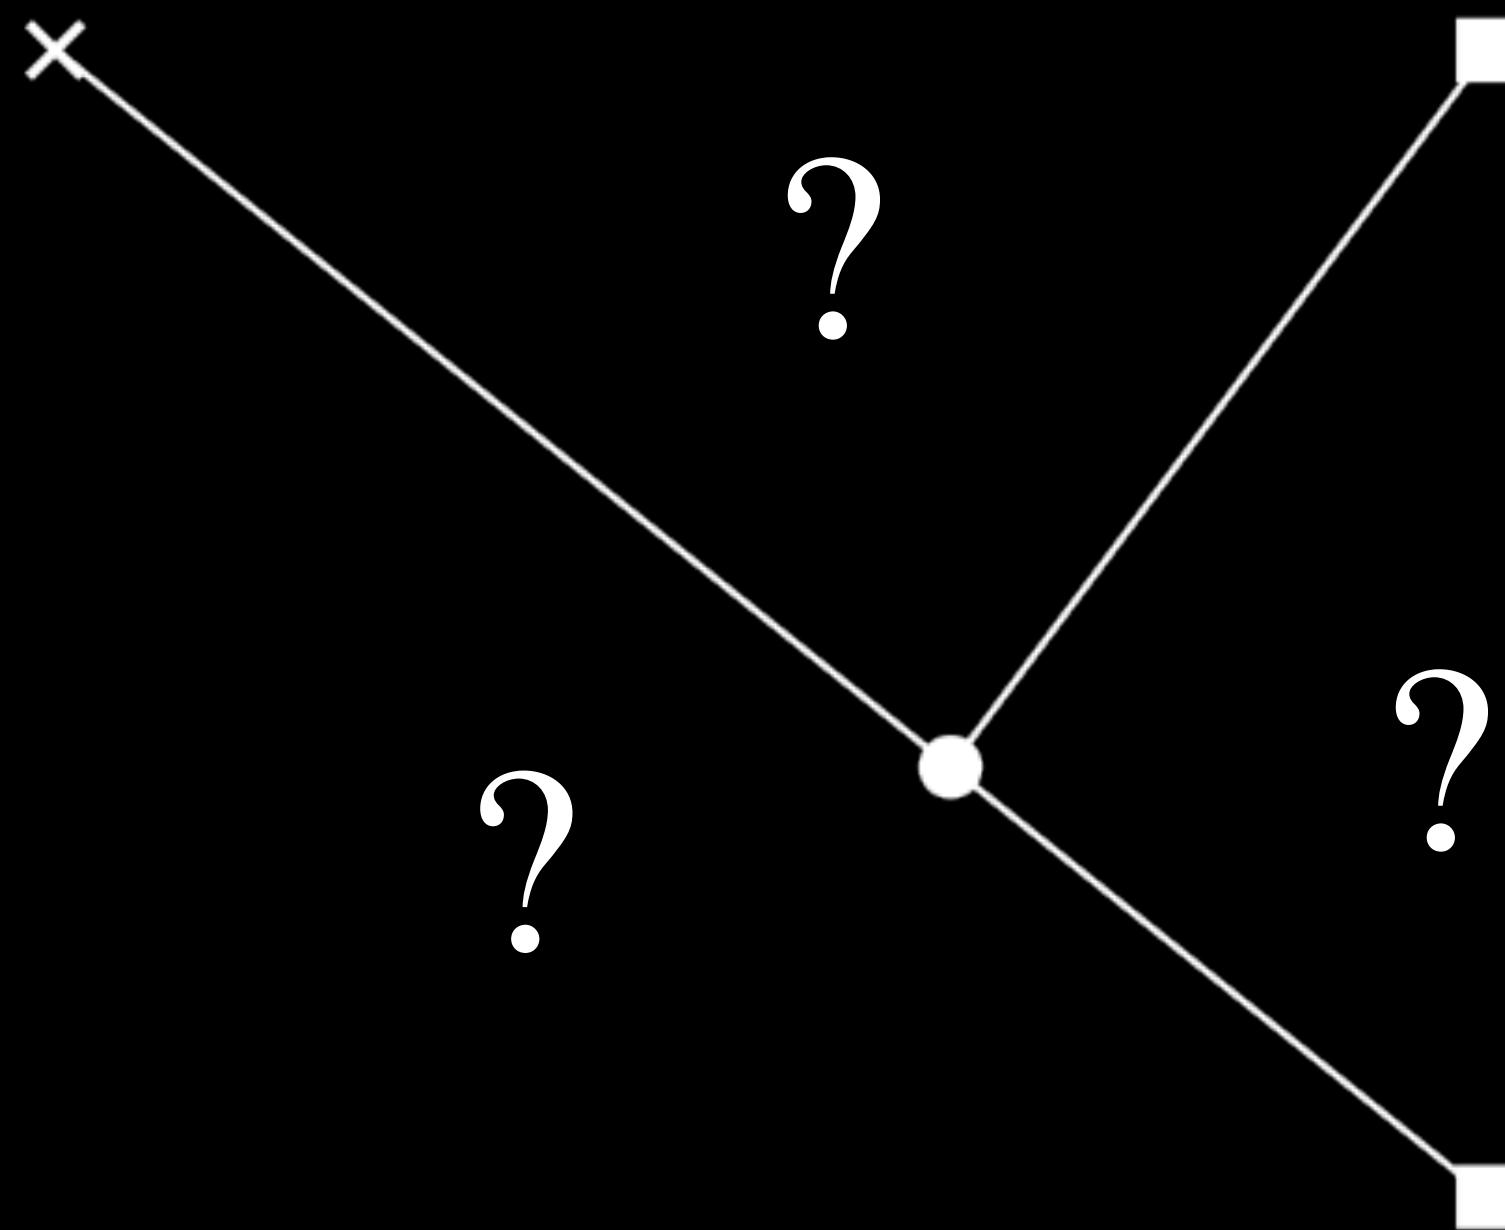
- Given a set of  $n$  planes  $H$ , compute  $L_0(H)$
- Find the set of intersection lines generated by  $H$
- Compute intersection and infinity vertices
- Filter out invalid segments
- Filter out invalid vertices



# Algorithms (2)

## Brute Force $L_0$

- For each face in  $L_0(H)$ , find the plane in  $H$  defining its surface



# Algorithms (3)

## Brute Force $L_0$

---

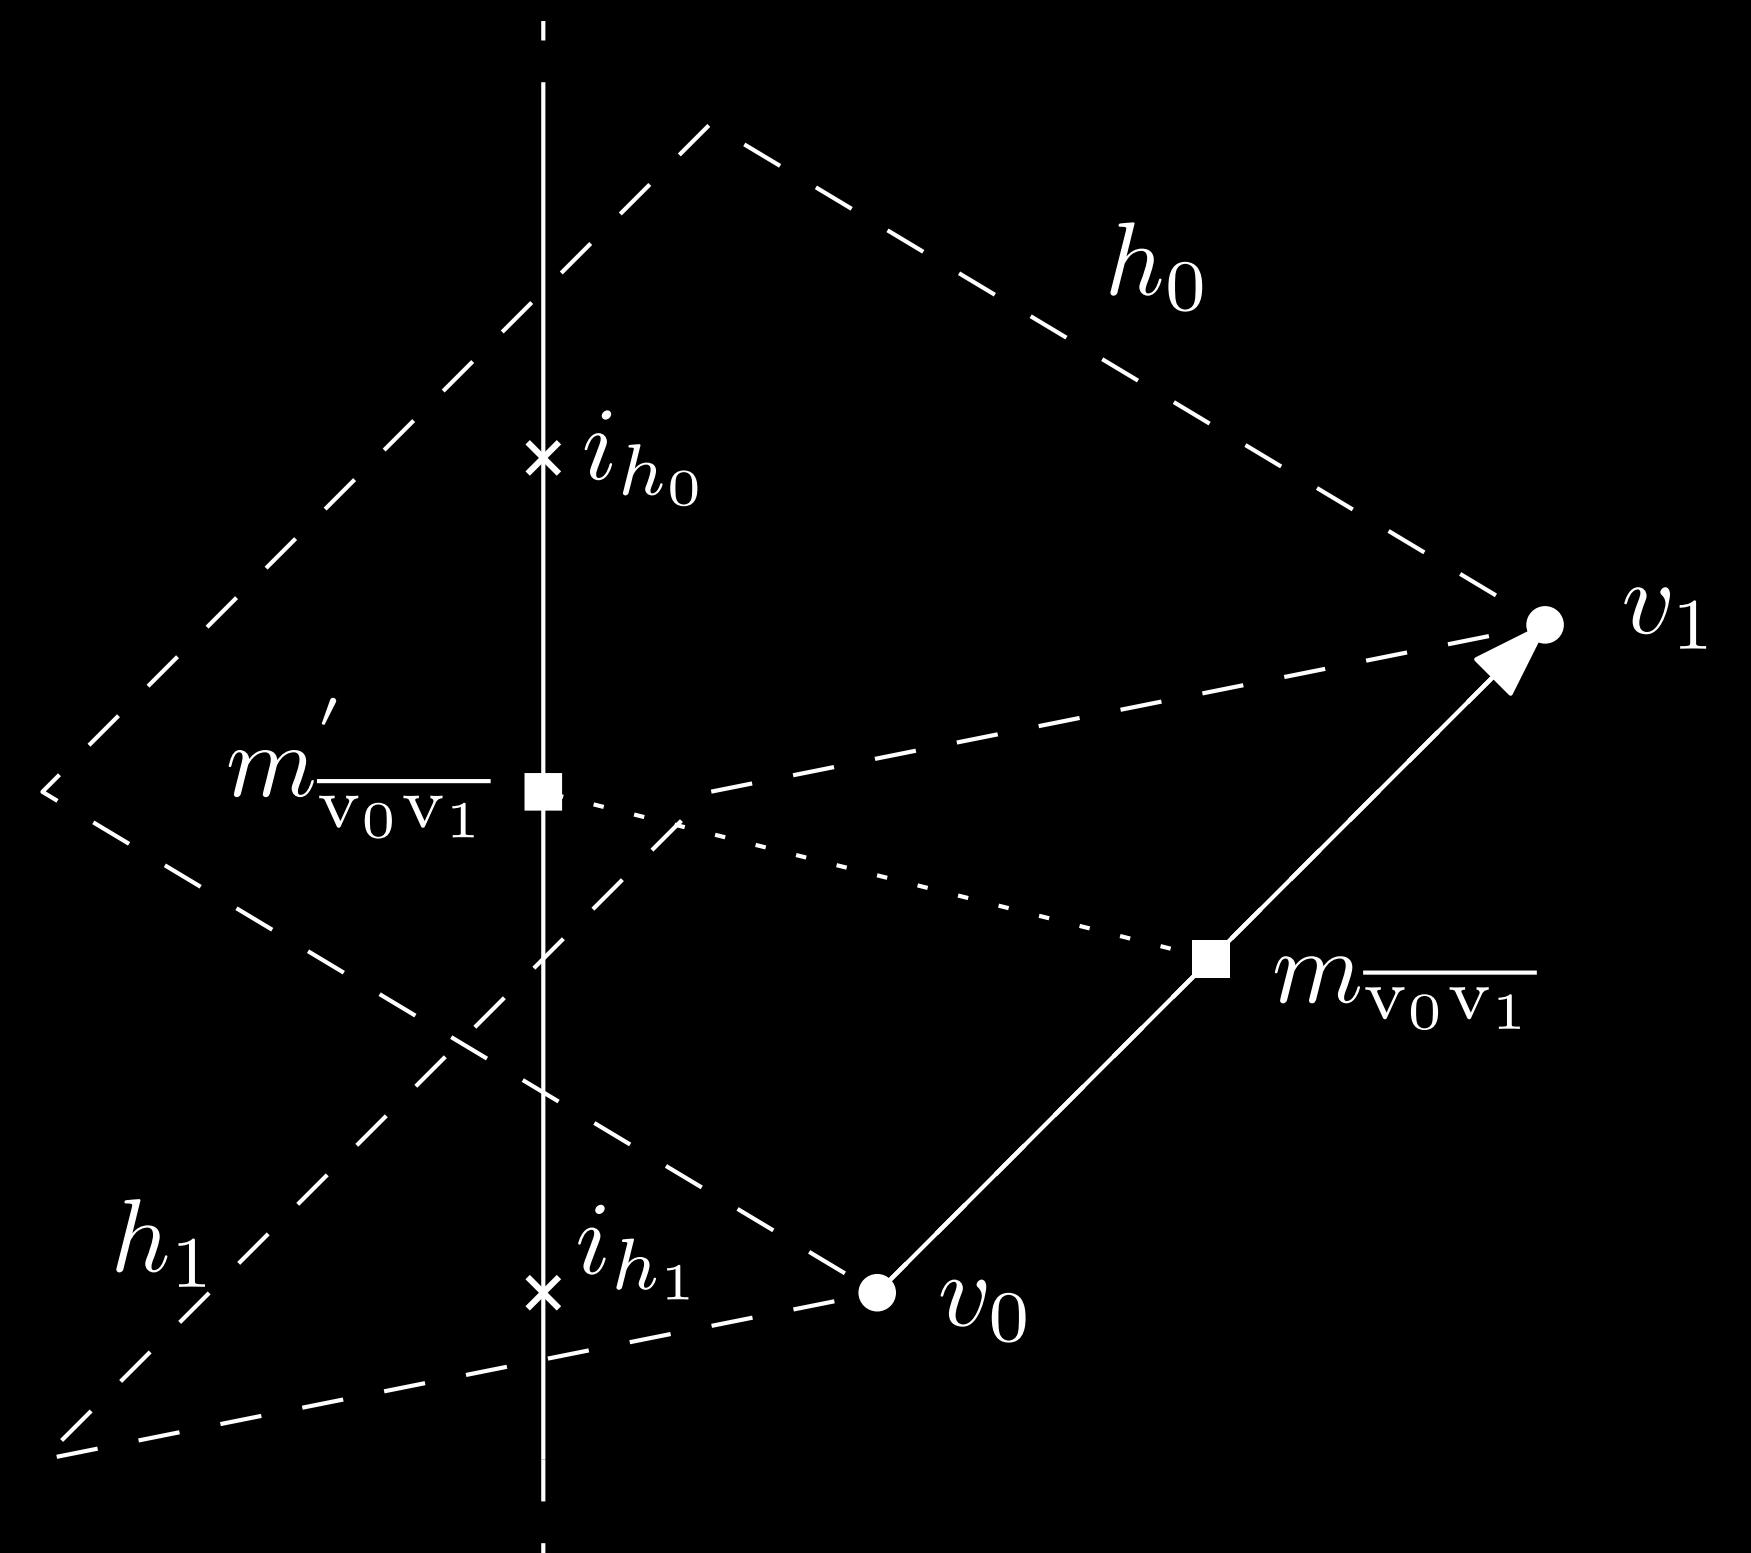
**Algorithm 4:** Compute  $H_{\text{left}}(v_0, v_1)$ 

---

**Data:**  $e = \overrightarrow{v_0 v_1} \in E$  and set  $H' \subset H$  generating  $e$

**Result:**  $H_{\text{left}}(v_0, v_1)$

- 1  $Z \leftarrow (0, 0, 1);$
  - 2  $T \leftarrow 1000;$
  - 3  $m_{\overrightarrow{v_0 v_1}} \leftarrow \frac{v_0 + v_1}{2};$
  - 4  $m'_{\overrightarrow{v_0 v_1}} \leftarrow -(\hat{e} \times Z) \cdot T;$
  - 5  $l \leftarrow \text{Vertical line passing through } m'_{\overrightarrow{v_0 v_1}};$
  - 6 **return**  $\arg \min_{h \in H'} (l \cap h)_z$
- 



# Algorithms (4)

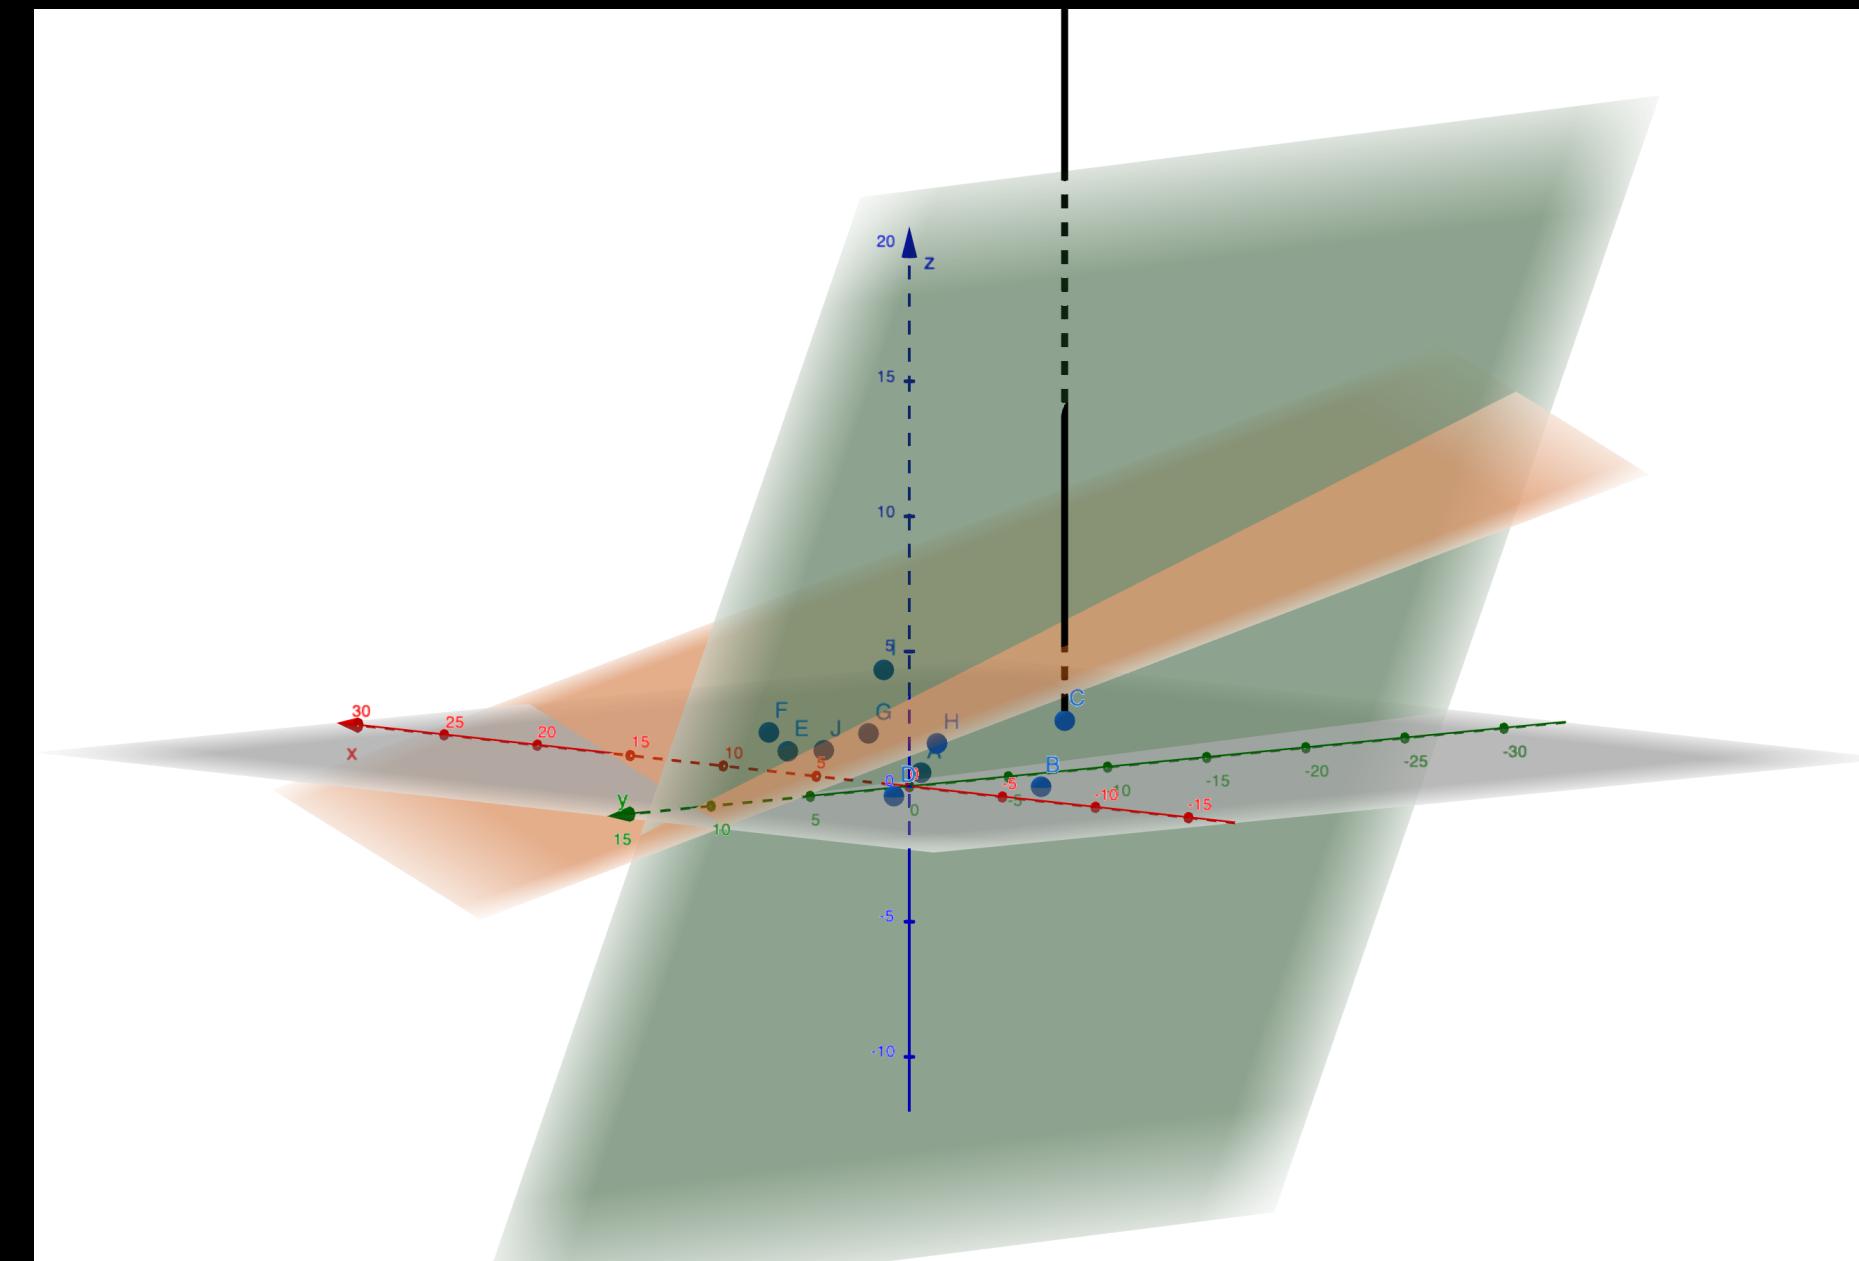
## Brute Force $L_0$

- Given  $n$  input planes and  $p$  processors, the algorithm runs in  $O\left(\frac{n^4}{p}\right)$

# Algorithms (5)

## Batch Point Location (GeoGebra visualisation)

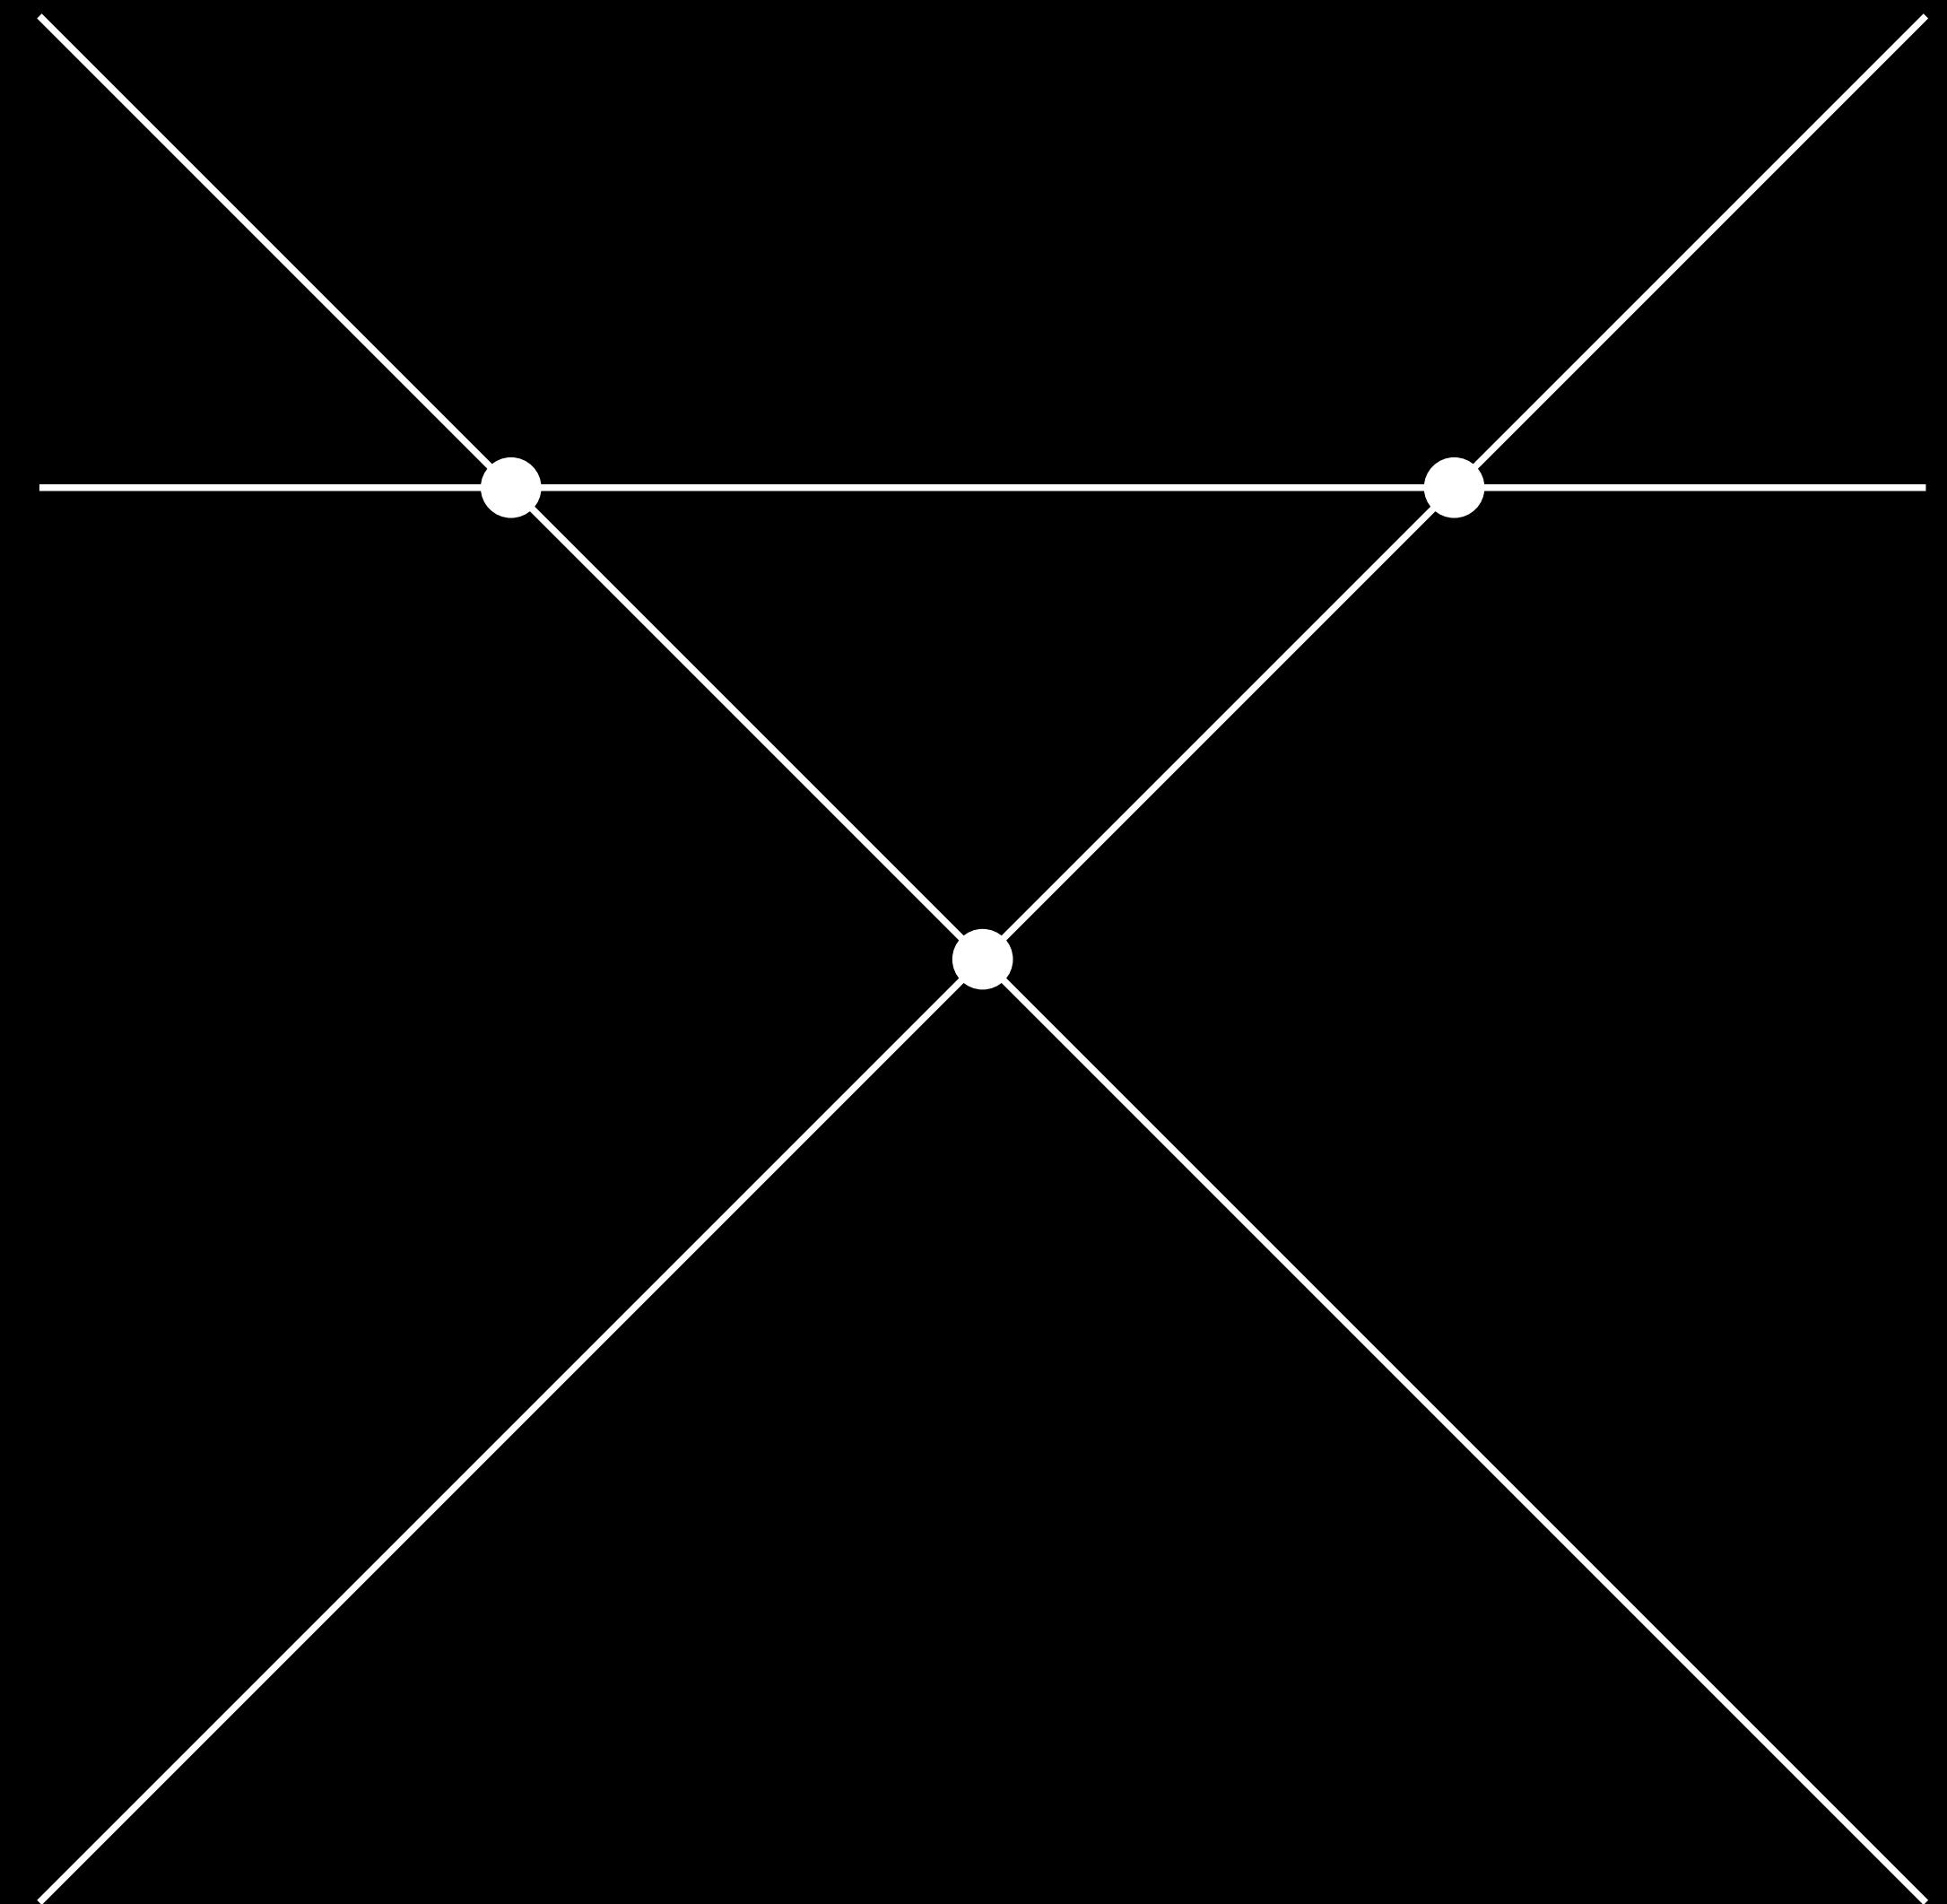
- Given a set of planes  $H$  and a set of points  $Q$ , compute the z-sorted sequence of planes  $\forall q \in Q$



# Algorithms (6)

## Batch Point Location

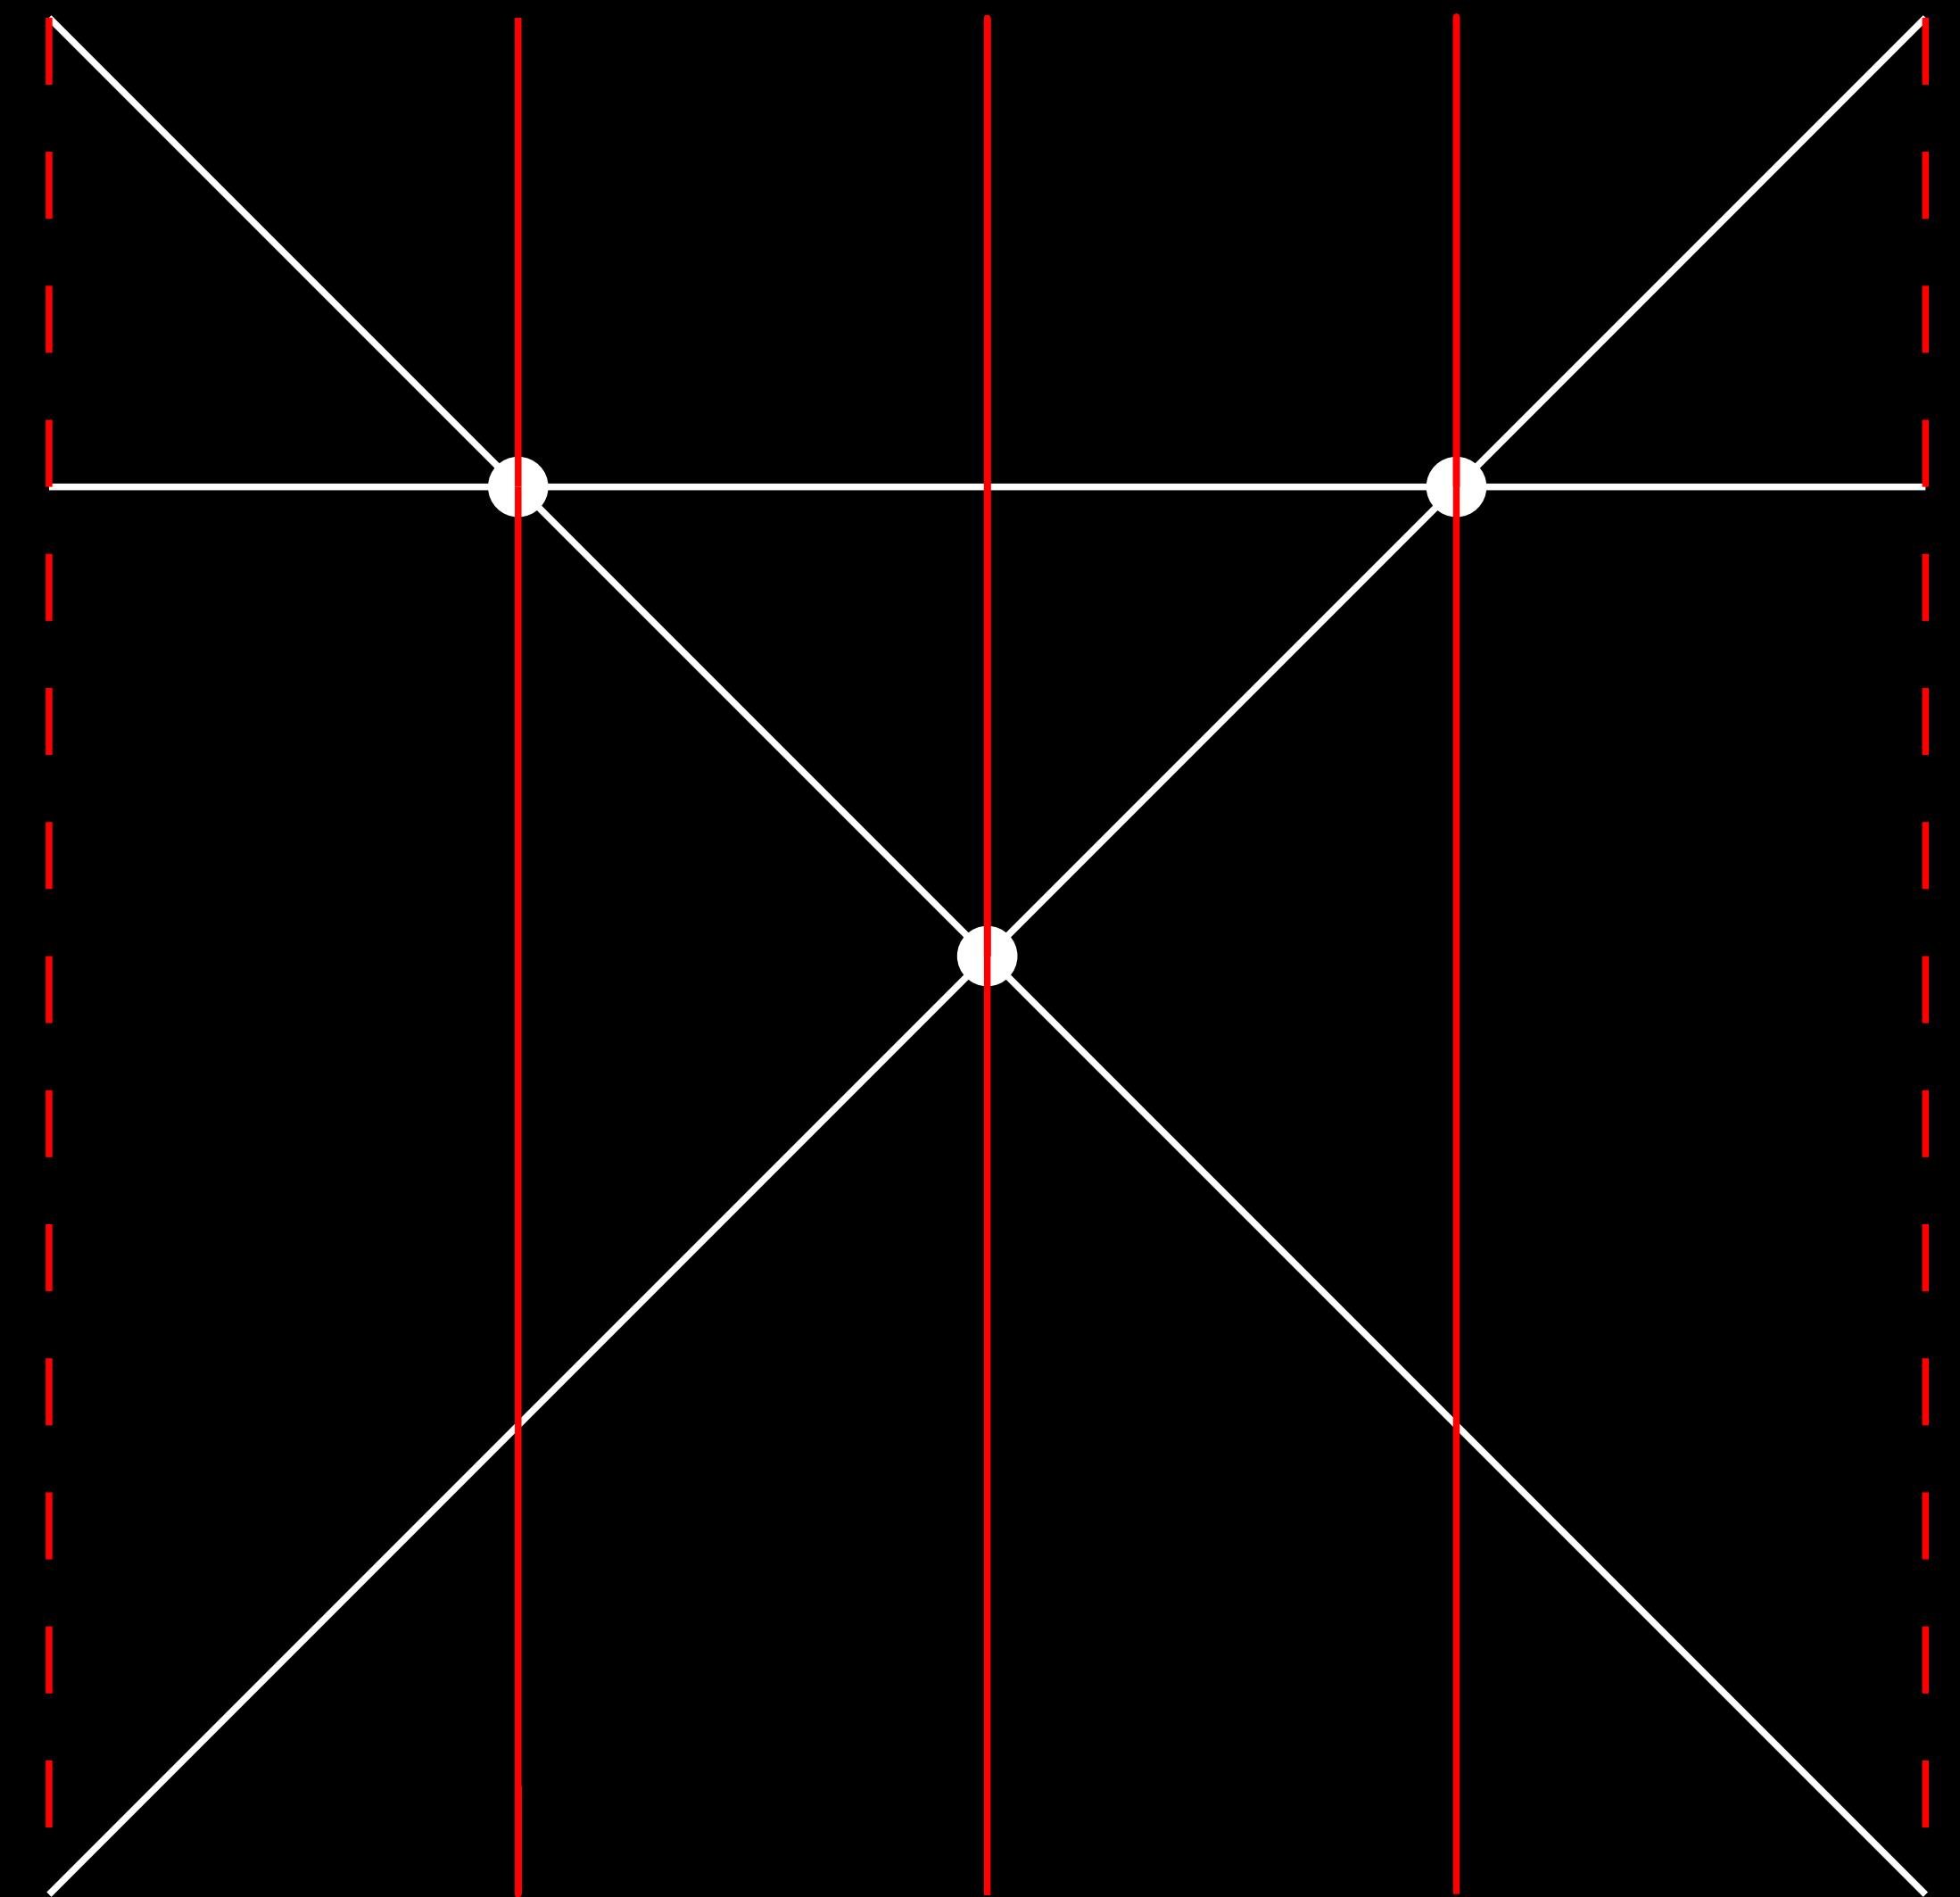
- Partition  $\mathbb{R}^2$  based on the projected intersection lines generated by  $H$



# Algorithms (6)

## Batch Point Location

- Partition  $\mathbb{R}^2$  based on the projected intersection lines generated by  $H$
- Further subdivide  $\mathbb{R}^2$  with vertical slabs

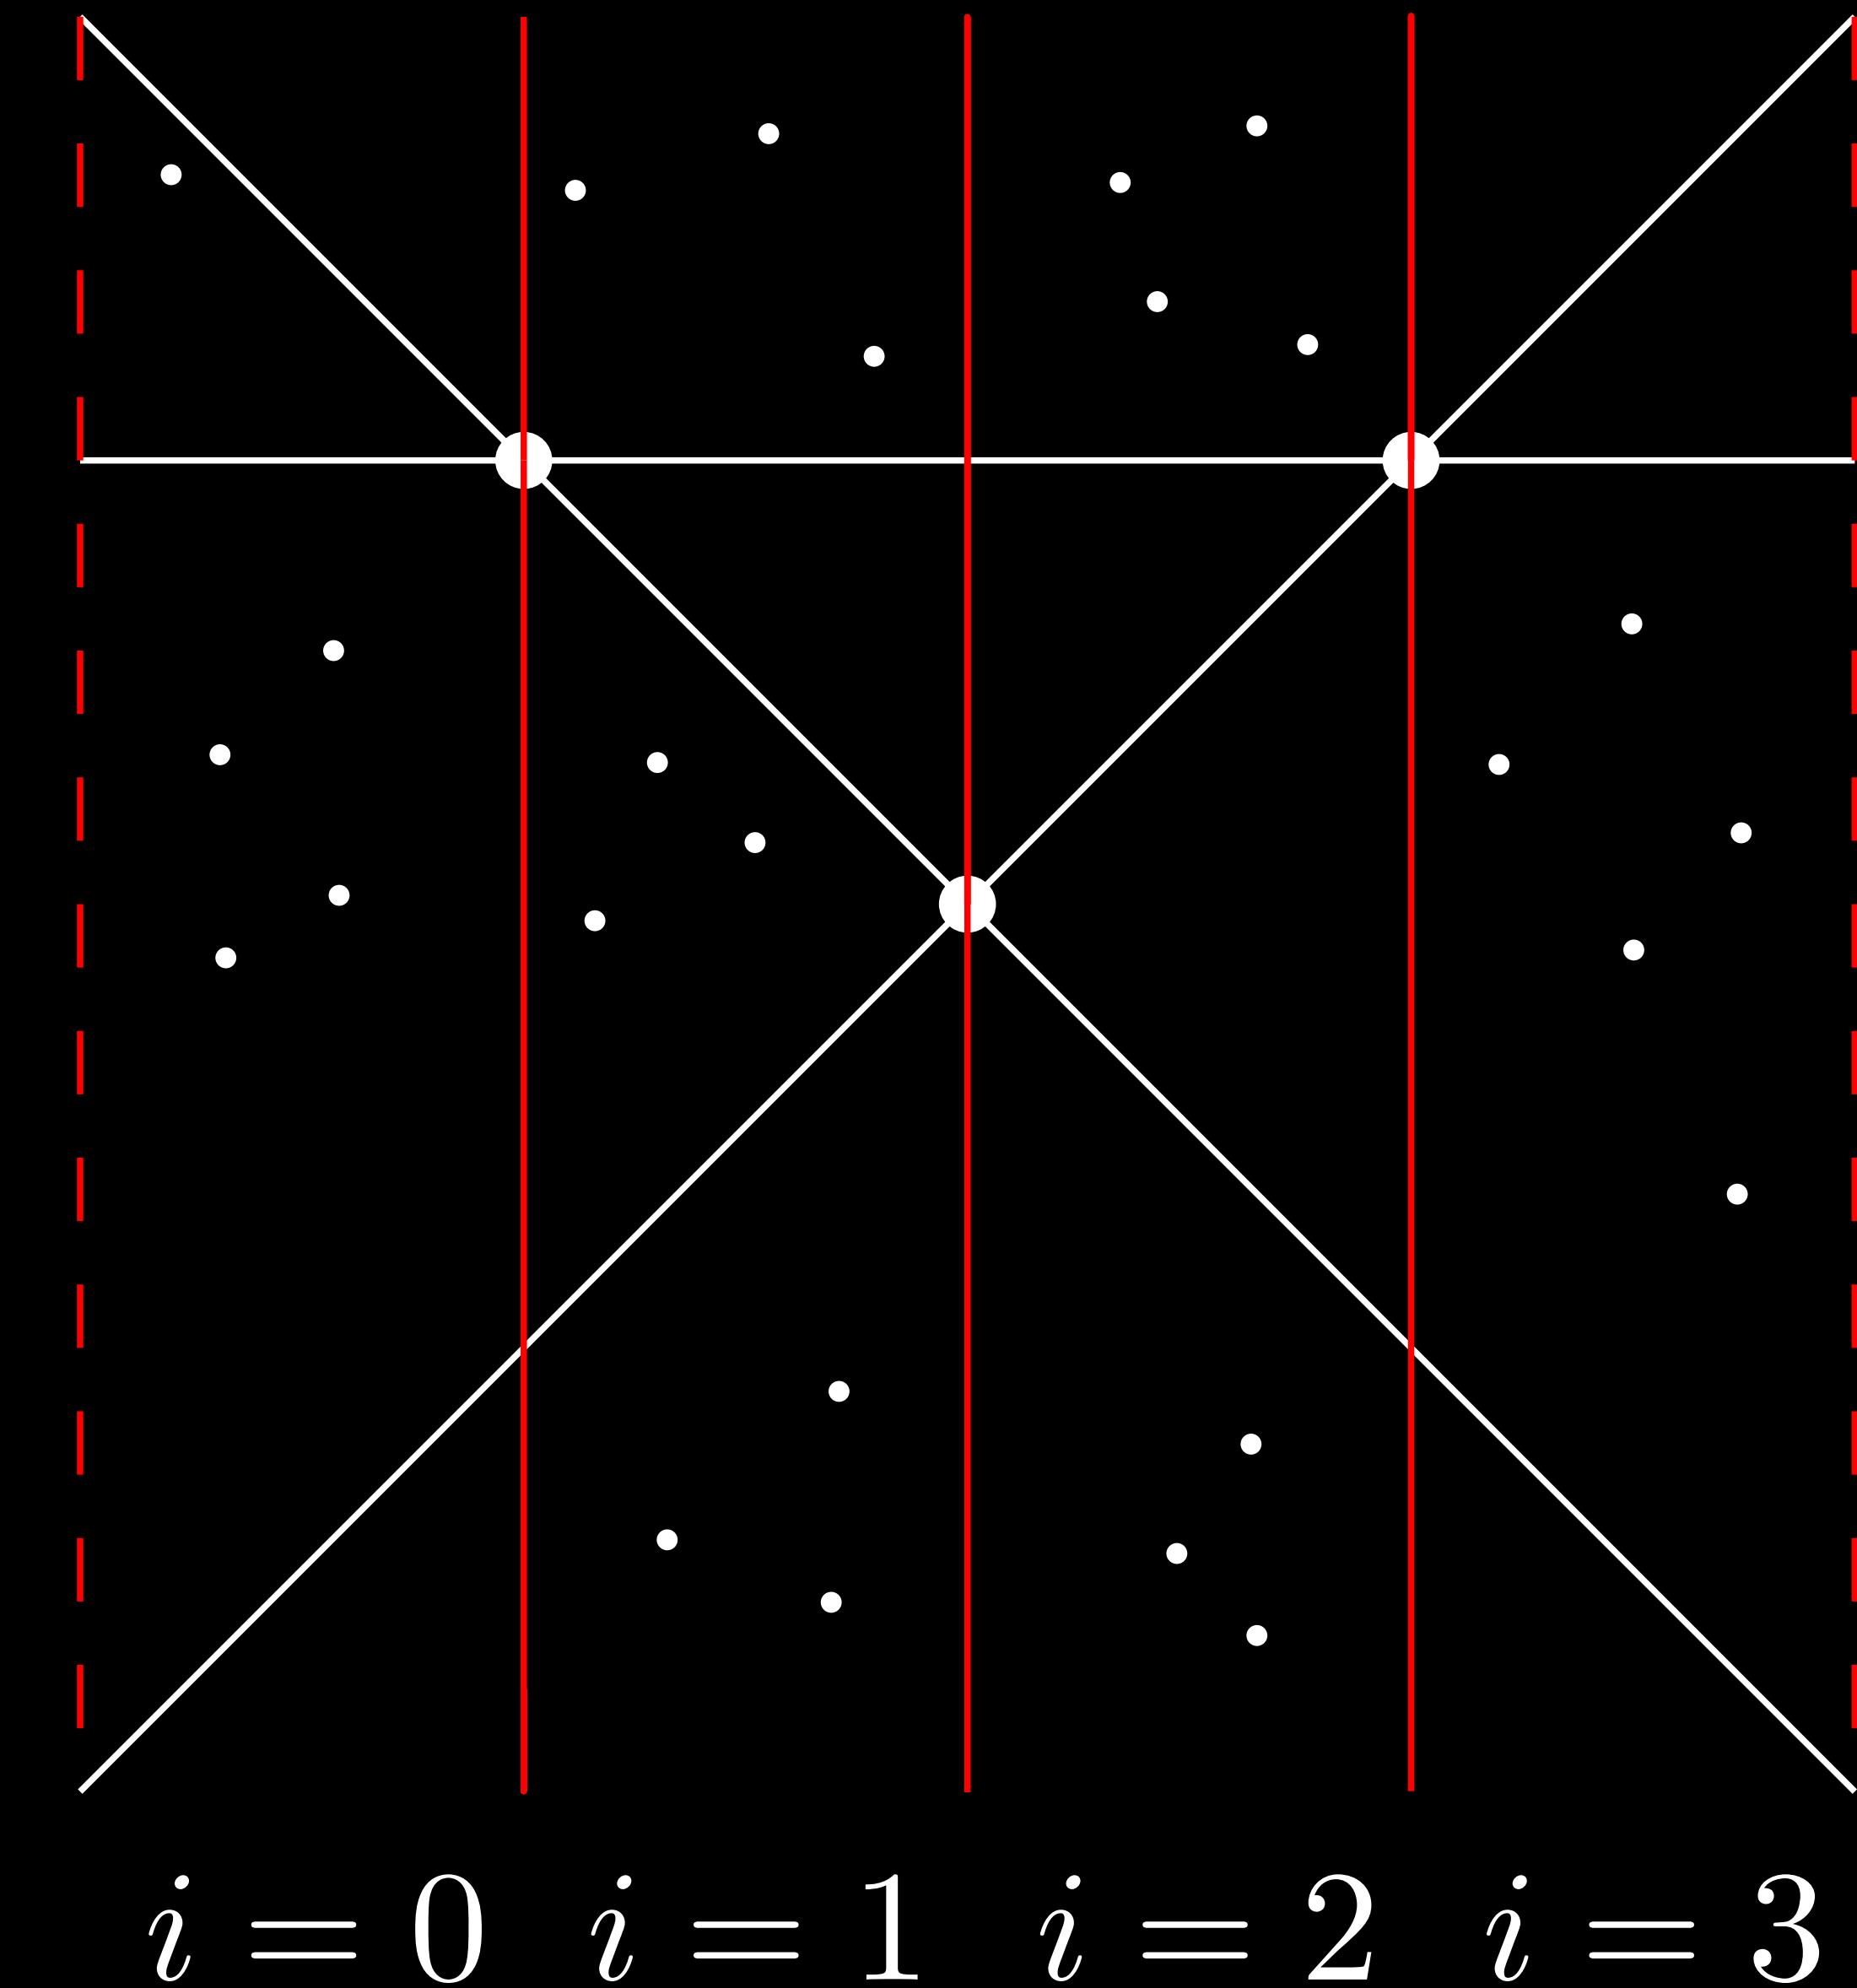


$i = 0 \quad i = 1 \quad i = 2 \quad i = 3$

# Algorithms (6)

## Batch Point Location

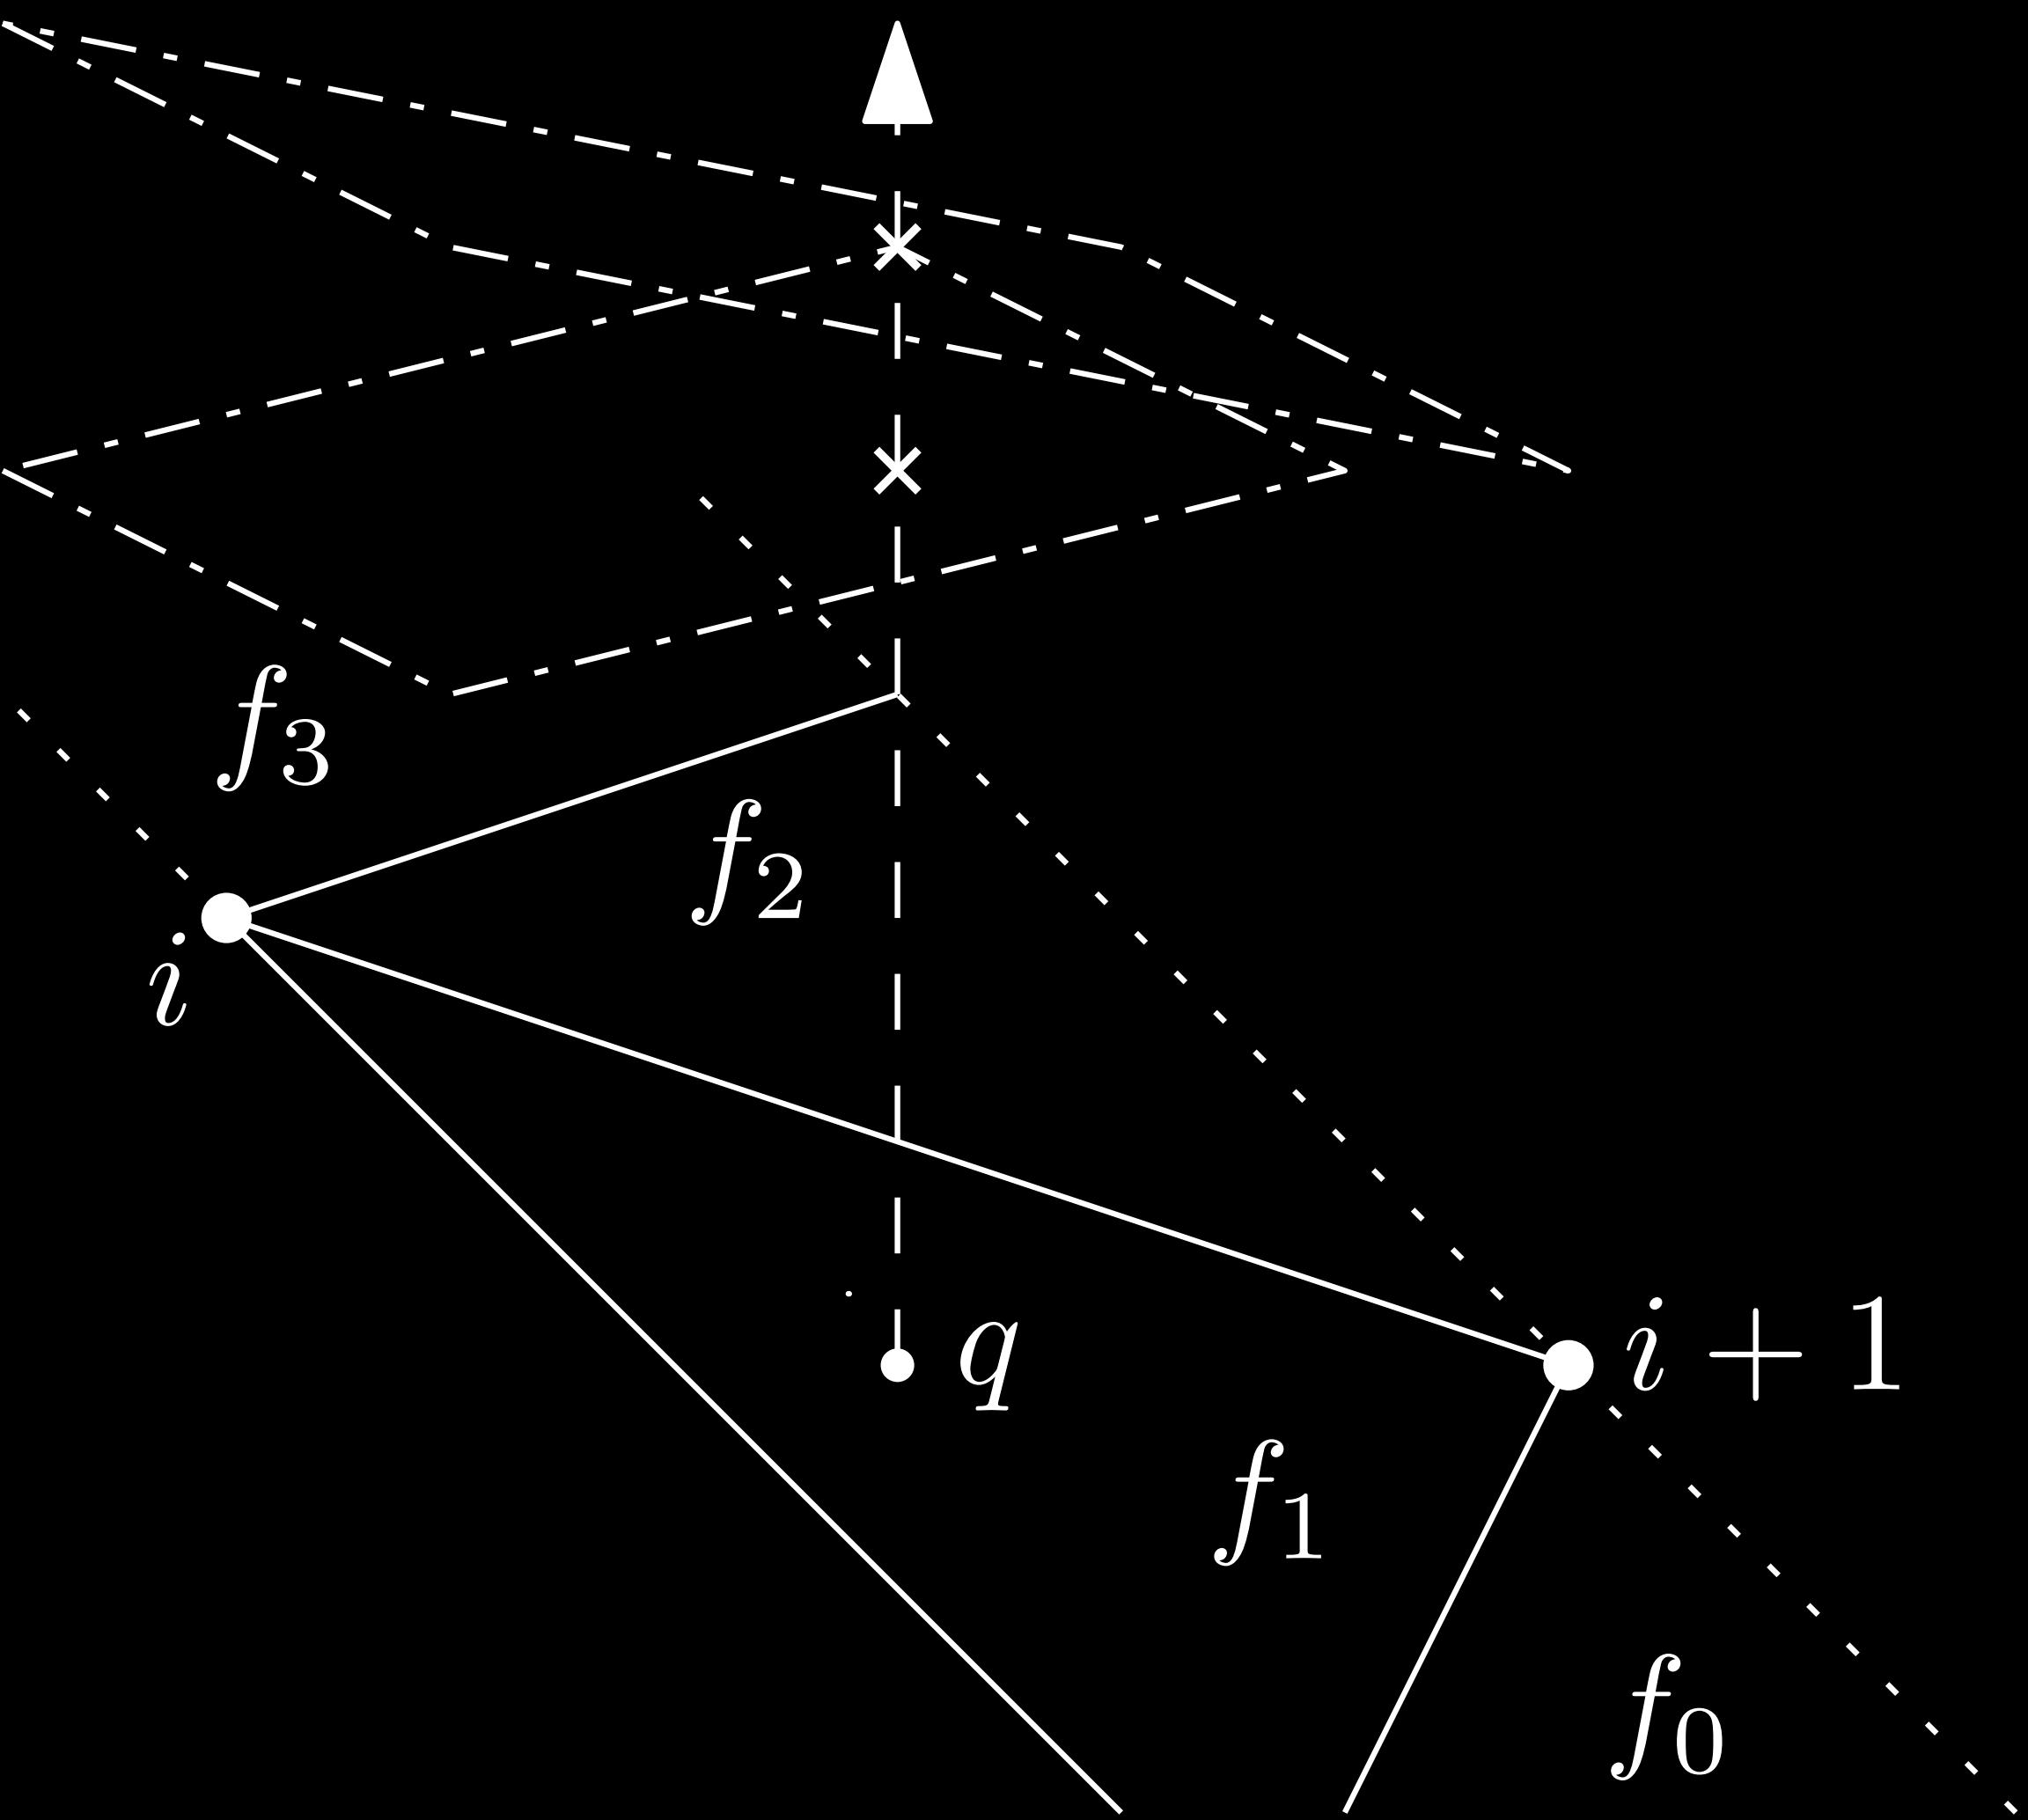
- Partition  $\mathbb{R}^2$  based on the projected intersection lines generated by  $H$
- Further subdivide  $\mathbb{R}^2$  with vertical slabs (Binary search)
- Multiple points may share the same face (Caching)
- Many faces are not used at all (Lazy Computation)



# Algorithms (7)

## Batch Point Location

- Divide computation into three independent phases:
  1. Slabs lookup
  2. Faces lookup
  3. Z-sorted planes computation



# Algorithms (8)

## Batch Point Location

- Given  $n$  input planes,  $k$  query points and  $p$  processors, the algorithm runs in

$$O\left(\frac{n^4 \log n + k \log n + on^2 \log n + fn \log n}{p}\right)$$

- With  $o = \#$  slabs used and  $f = \#$  faces used

# Algorithms (9)

## Batch Point Location

- Find the set of planes passing below each prism in  $L_0(S)$  with  $S \subset T$
- Small and bounded amount of prisms  $\Lambda_0(S)$
- Many planes to test in  $T - S$
- Achieve efficiency with duality:
  1.  $Q = (T - S)^\star$  is a set of query points (expected to be very large)
  2.  $H = \Lambda_0(S)^\star$  is a set of planes (expected to be very small)

# Algorithms (10)

## Missing algorithms

- Lower envelope triangulation, adjacency and merge algorithms

**Algorithm 9:** Assemble final envelope

---

**Data:**  $P, L = L_0(C_i)$   
**Result:**  $L_0(H)$

- 1  $E_0 \leftarrow (v_0^e, v_1^e, \Delta_{\text{idx}}) \forall \Delta \in L_0(S) \forall e \in \Delta;$
- 2  $L_0 \leftarrow L;$
- 3  $T_0 \leftarrow |P|$  booleans set to false;
- 4 **while**  $\exists a, b \in E_i \mid a_2 \neq b_2$  **do**
- 5     Lexicographically sort  $E_i$ ;
- 6     Remove edges defined by only one triangle;
- 7     **foreach** Pair of consecutive edges  $(a, b) \in E_{i-1}$  **do**
- 8         **if**  $T[a] \vee T[b]$  **then**
- 9             **continue**;
- 10         **end**
- 11          $T[a], T[b] \leftarrow \text{true, true};$
- 12          $m \leftarrow \text{Merge envelopes } (a, b) \text{ using } \text{Merge}_{\Delta};$
- 13         Store envelope in  $L_i$ ;
- 14          $E_i \leftarrow \text{Insert } (a_0, a_i, \min(a_{\text{idx}}, b_{\text{idx}}) \forall a \in E_{i-1};$
- 15         **end**
- 16         Set all entries in  $T$  to false;
- 17 **end**

---

**Algorithm 11:** Triangulates each face of  $L_0(H)$

---

**Data:**  $L_0(H)$   
**Result:** A triangulated  $L_0(H)$

- 1  $s, p, E \leftarrow 0, -1, <_{\Delta} L_0(H);$
- 2 **for**  $i \leftarrow 1$  **to**  $|E|$  **do**
- 3     **if**  $i = |E| \vee E[s]_{\text{plane}} \neq E[i]_{\text{plane}}$  **then**
- 4         **for**  $k \leftarrow s + 1$  **to**  $i$  **do**
- 5             **if**  $ps > -1 \wedge (k = p \vee E[k]_{\text{end}} = E[p]_{\text{start}})$  **then**
- 6                 **continue**;
- 7             **end**
- 8              $E \leftarrow \text{Insert edge } \langle E[s]_{\text{start}}, E[k]_{\text{end}} \rangle;$
- 9              $E \leftarrow \text{Insert edge } \langle E[k]_{\text{end}}, E[s]_{\text{start}} \rangle;$
- 10         **end**
- 11          $s, p \leftarrow i, -1;$
- 12     **end**
- 13     **else if**  $E[i]_{\text{end}} = E[s]_{\text{start}}$  **then**
- 14          $ps \leftarrow i;$
- 15     **end**
- 16 **end**
- 17 **return**  $< E;$

---

# Implementation Details

# Implementation Details (1)

## Development Environment and 3rd Party Libraries

- Using CMake and Git
- Using C++ 17 with Visual Studio 2022 (Windows)
- Using CGAL for exact arithmetic
- Using SFML for visualisation
- Using Catch2 for Test-Driven Development
- Using OneAPI Intel TBB for C++ code parallelisation

# Implementation Details (2)

## Pair-wise Intersections

- Reduced from  $n^2$  to  $\approx \frac{n^2}{2}$  intersection tests
- Work divided in tiles based on the amount of instructions
- Contiguous tiles scheduled to the same thread
- Results stored locally to each thread and spliced at the end

# Implementation Details (3)

## Pair-wise Intersections

```
class Observer : public tbb::task_scheduler_observer
{
    size_t& myFinalSize;
    std::vector<Point2>& myOutData;
public:
    Observer(std::vector<Point2>& someOutData, size_t& aFinalOutCount)
        : tbb::task_scheduler_observer(), myOutData(someOutData), myFinalSize(aFinalOutCount)
    {
        observe(true);
    }

    ~Observer()
    {
        observe(false);
        for (const auto& chunk : myData)
        {
            myFinalSize += chunk.size();
        }
    }

    void on_scheduler_entry(bool /*worker*/)
    {
        CGAL_precondition(myData.local().empty());
    }

    void on_scheduler_exit(bool /*worker*/)
    {
        const int start = myCounter.fetch_add(myData.local().size());
        std::copy(myData.local().begin(), myData.local().end(), myOutData.begin() + start);
    }
};

tbb::enumerable_thread_specific<std::vector<Point2>> myData;
std::atomic<size_t> myCounter{ 0 };
};
```

```
{
    Observer obs(anOutData.myVertices, finalSize);
    static tbb::affinity_partitioner ap;
    tbb::parallel_for(tbb::blocked_range2d<int, int>(0, anOutData.myLines.size(), 0, anOutData.myLines.size()),
        [&](tbb::blocked_range2d<int, int>& r) {
            if (r.rows().end() > r.cols().begin())
            {
                for (int i = r.rows().begin(); i < r.rows().end(); ++i)
                {
                    for (int k = r.cols().begin(); k < r.cols().end(); ++k)
                    {
                        if (i <= k) continue;
                        const auto intersection = CGAL::intersection(anOutData.myLines[k], anOutData.myLines[i]);
                        if (intersection)
                        {
                            const Point2* point = boost::get<Point2>(&*intersection);
                            CGAL_precondition(point != nullptr);
                            obs.myData.local().emplace_back(*point);
                        }
                    }
                }
            }
        });
}

const int start = obs.myCounter.fetch_add(obs.myData.local().size());
std::copy(obs.myData.local().begin(), obs.myData.local().end(), anOutData.myVertices.begin() + start);
};
```

# Implementation Details (4)

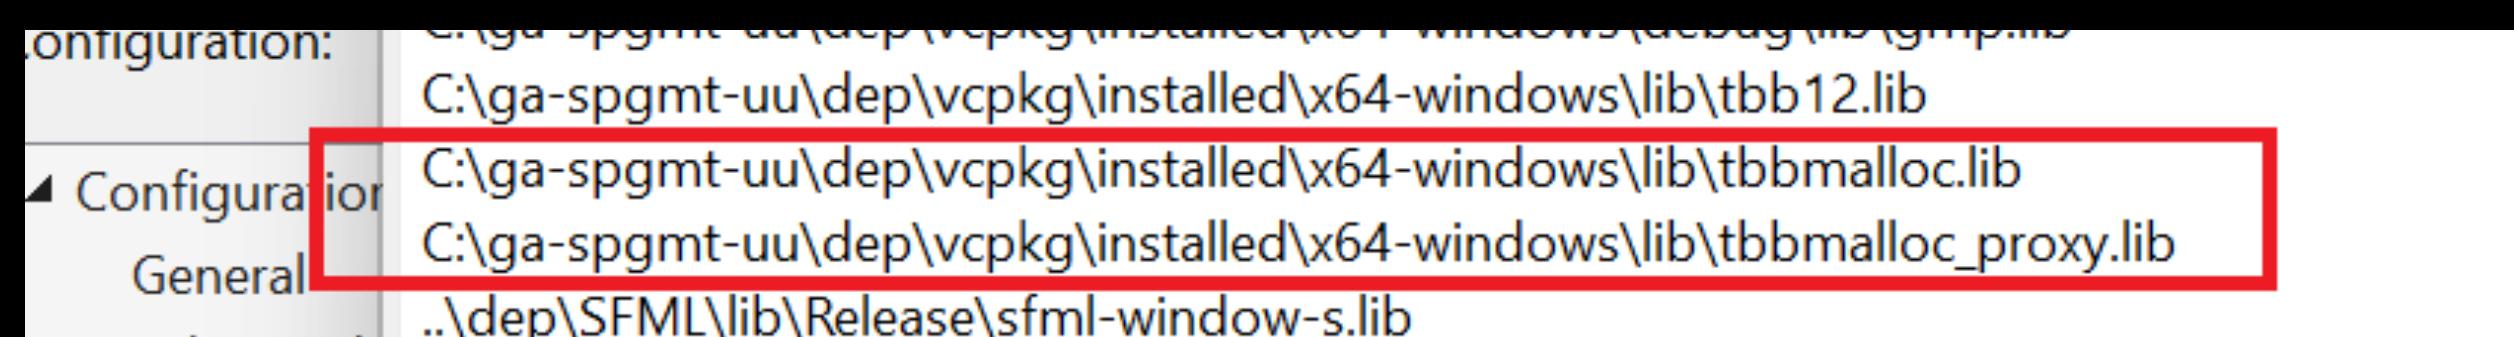
## Allocation and Deallocation Overhead

- Around 35% of the execution time spent for allocation and deallocation
- Nested data structures are expensive to deallocate
- Solutions:
  1. Use TBB scalable and cache aligned allocator
  2. Use plain arrays instead of arrays of nested structures

# Implementation Details (5)

## Allocation & Deallocation Overhead

```
// OLD  
std::vector<std::tuple<FT, std::vector<FT>>> mySlabs;  
  
// NEW  
std::vector<Point2> myVertices;
```



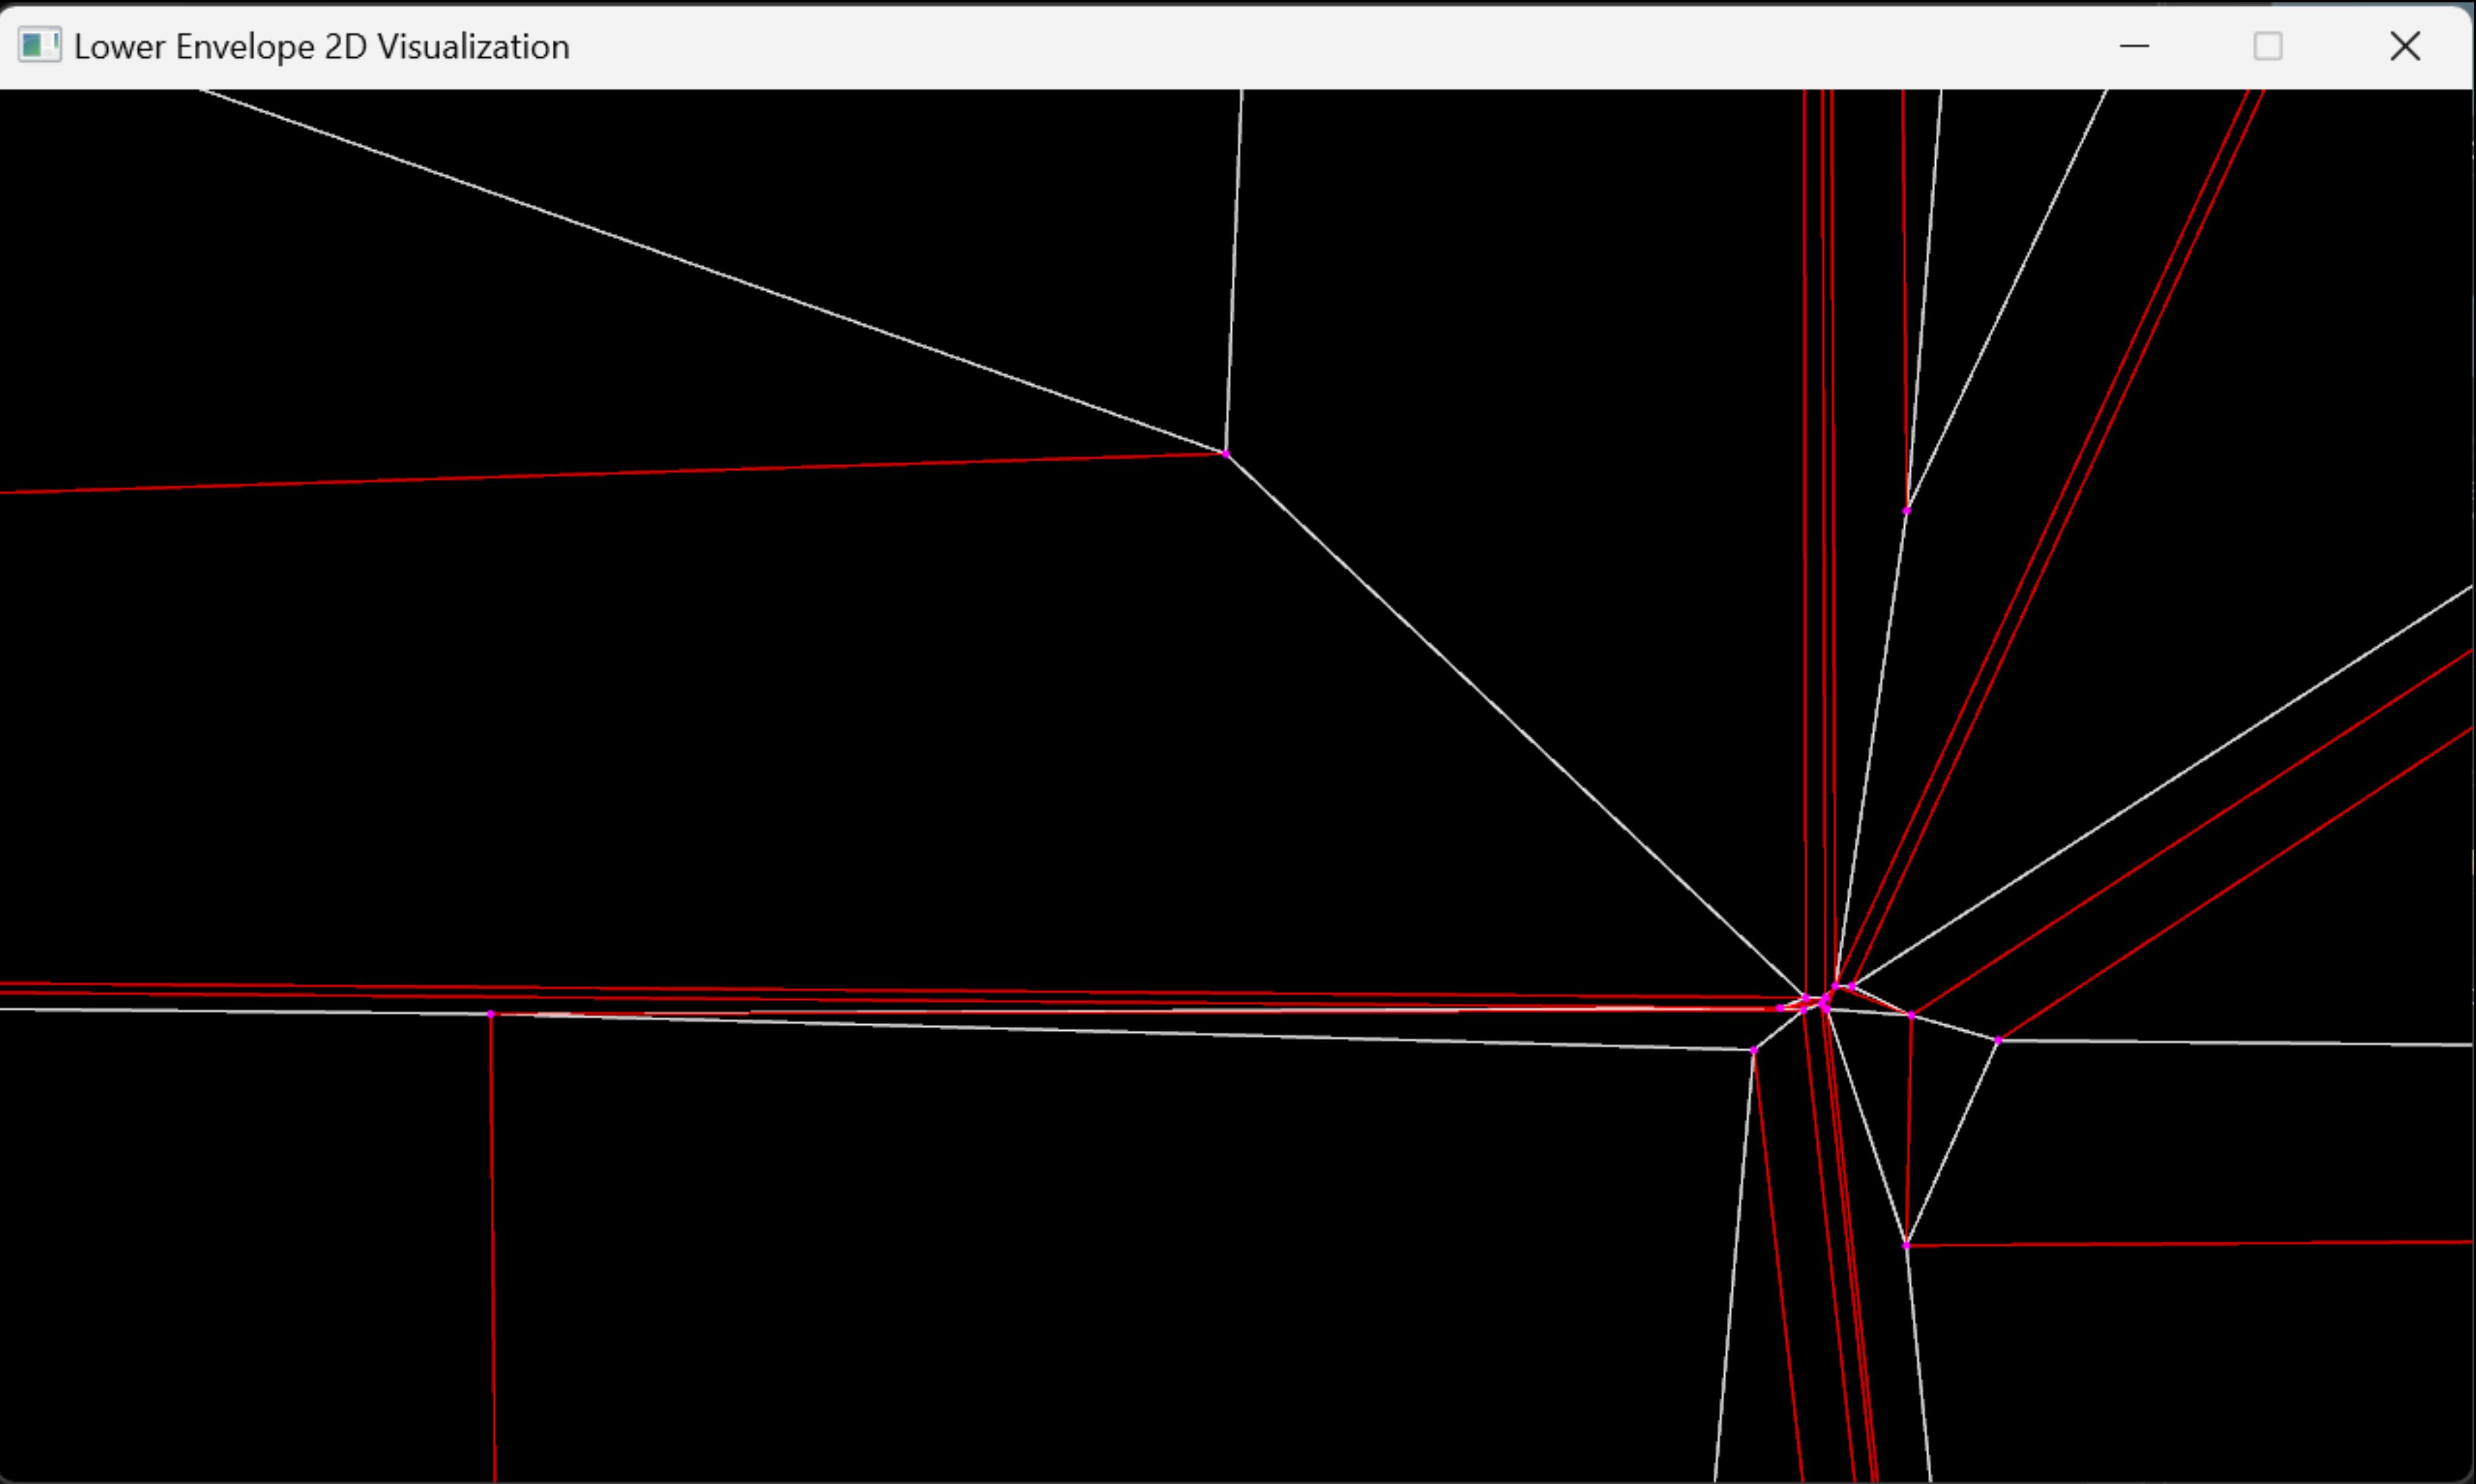
```
24  
25  [] #include <oneapi/tbb/tbbmalloc_proxy.h>  
26
```

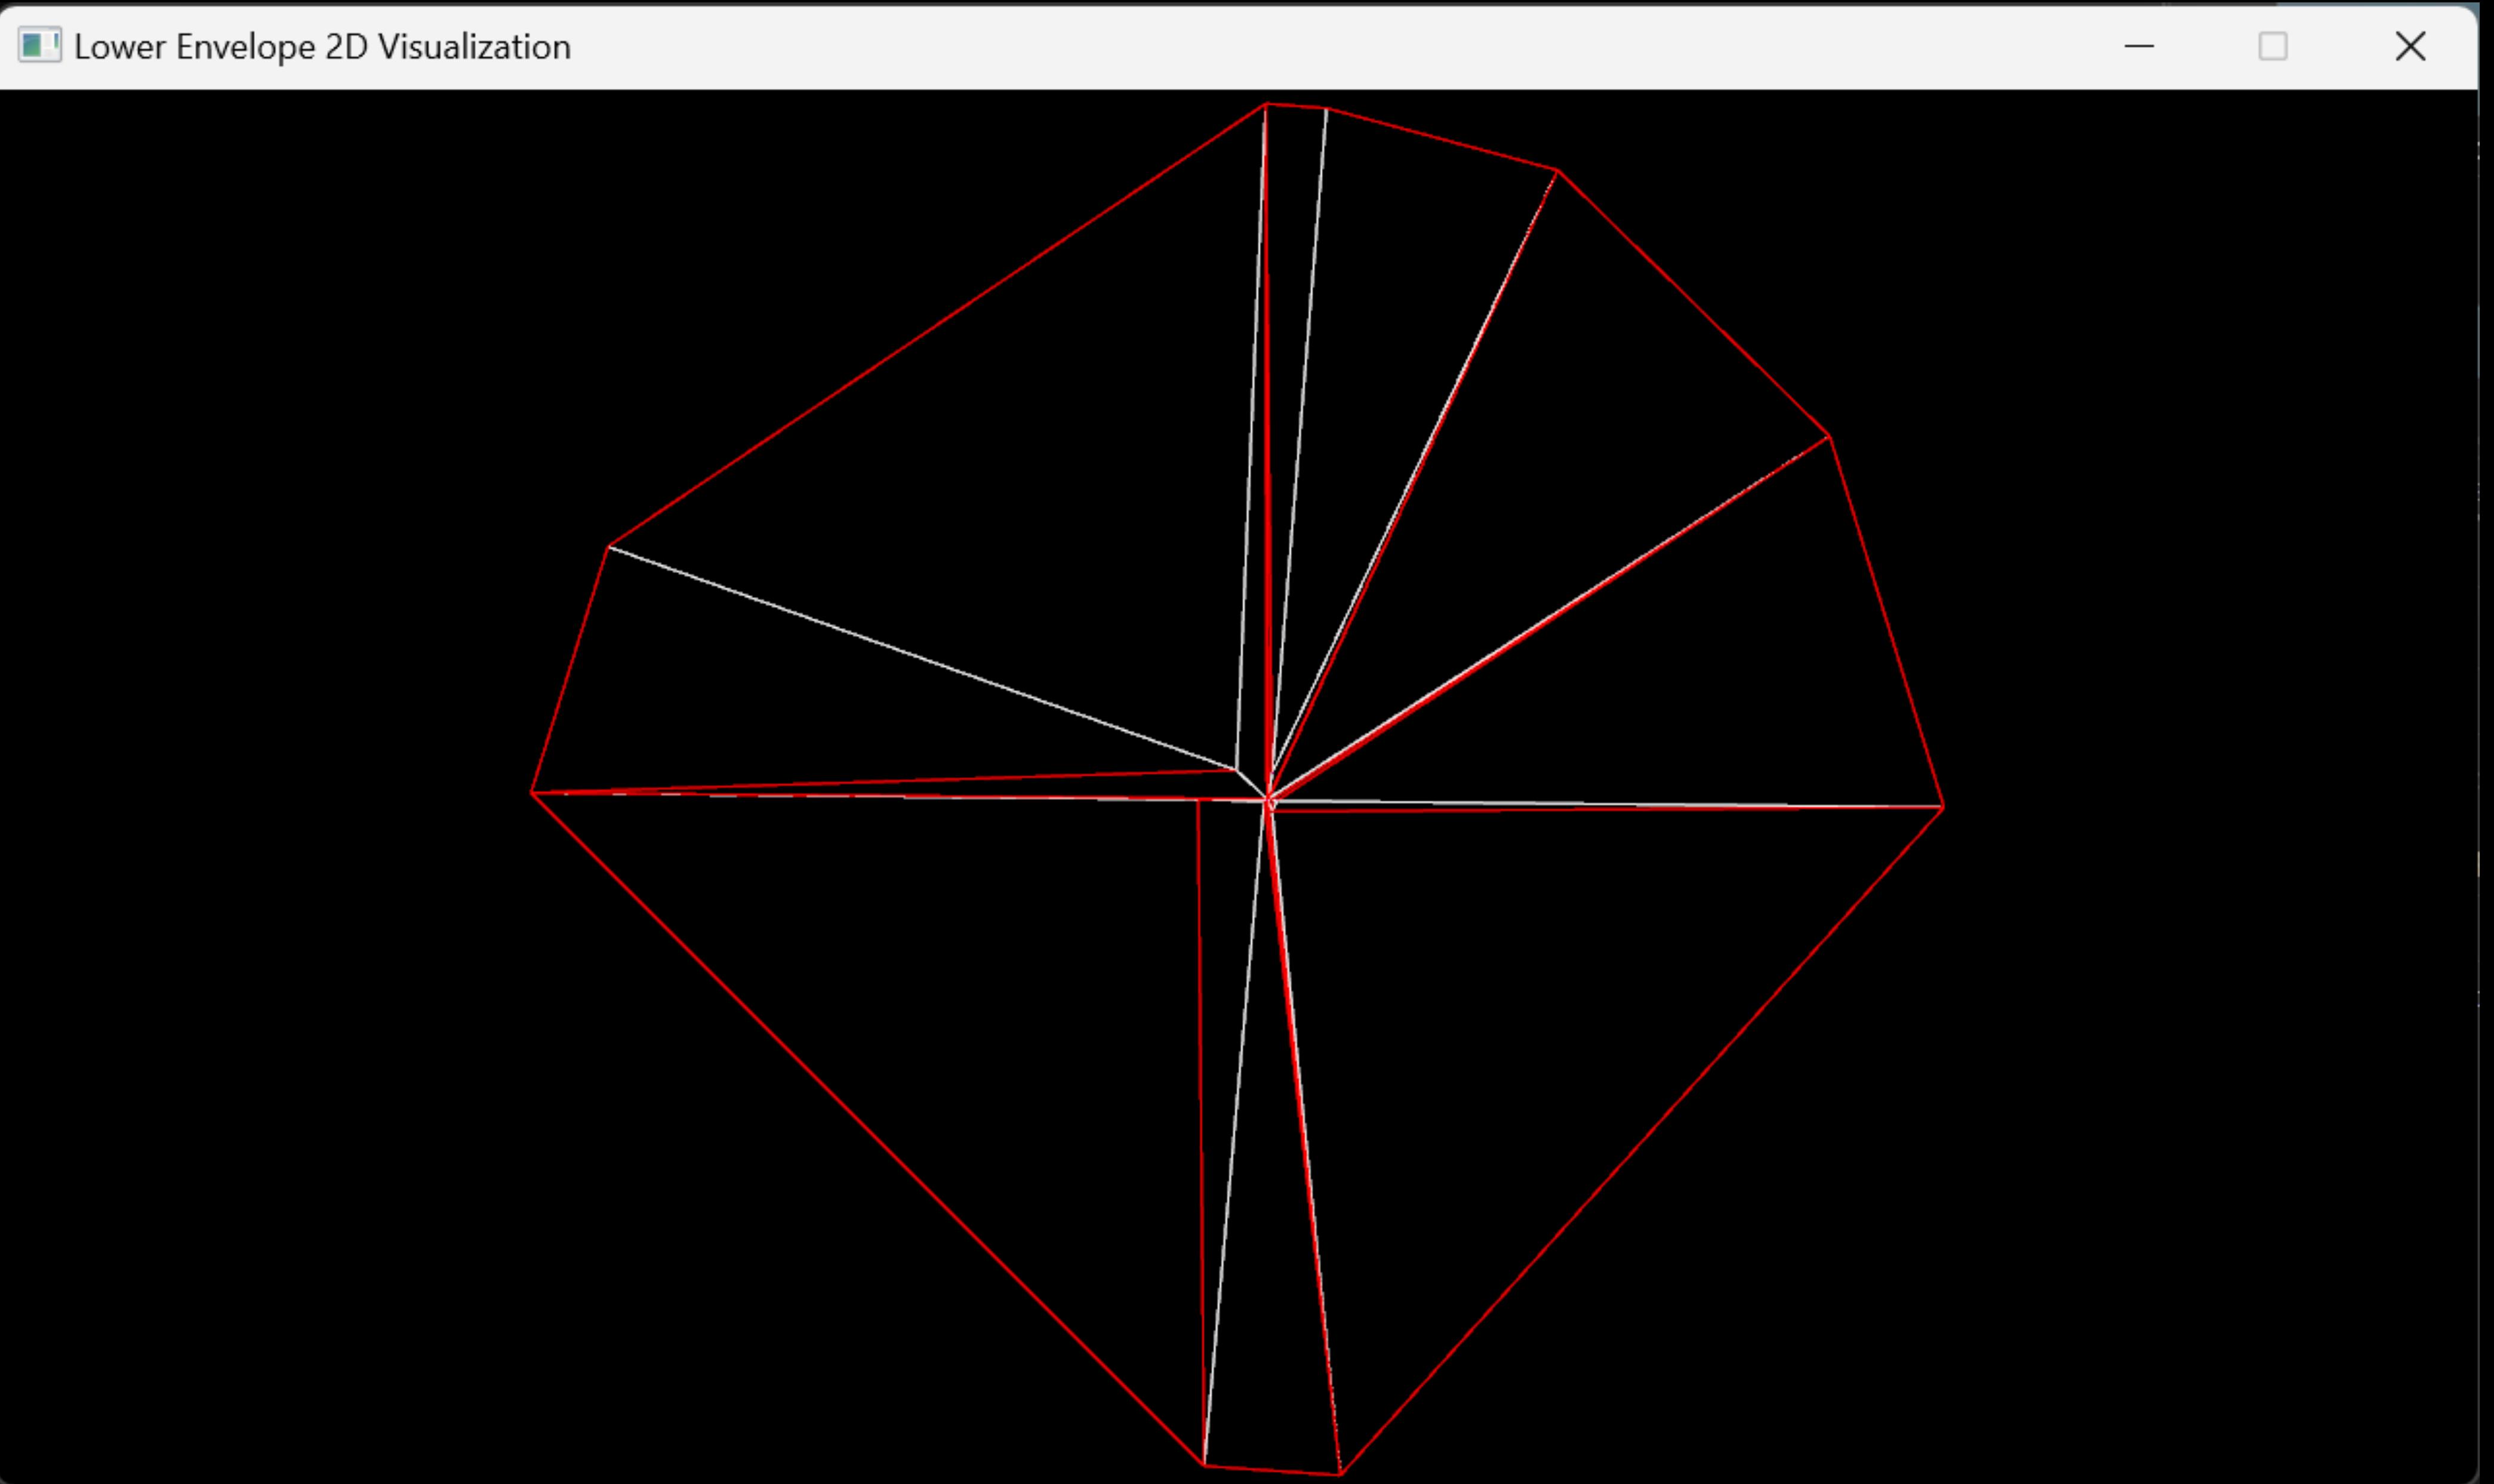
# Implementation Details (6)

## Exact Arithmetic

- Using CGAL::Lazy\_exact\_nt
- Around 25% penalty per arithmetic operation
- Dynamic allocations
- Stored values impossible to inspect while debugging
- Doubles or floats would significantly improve performances, but inexact computations require special treatment

# Visualisation





# Benchmarks and Results

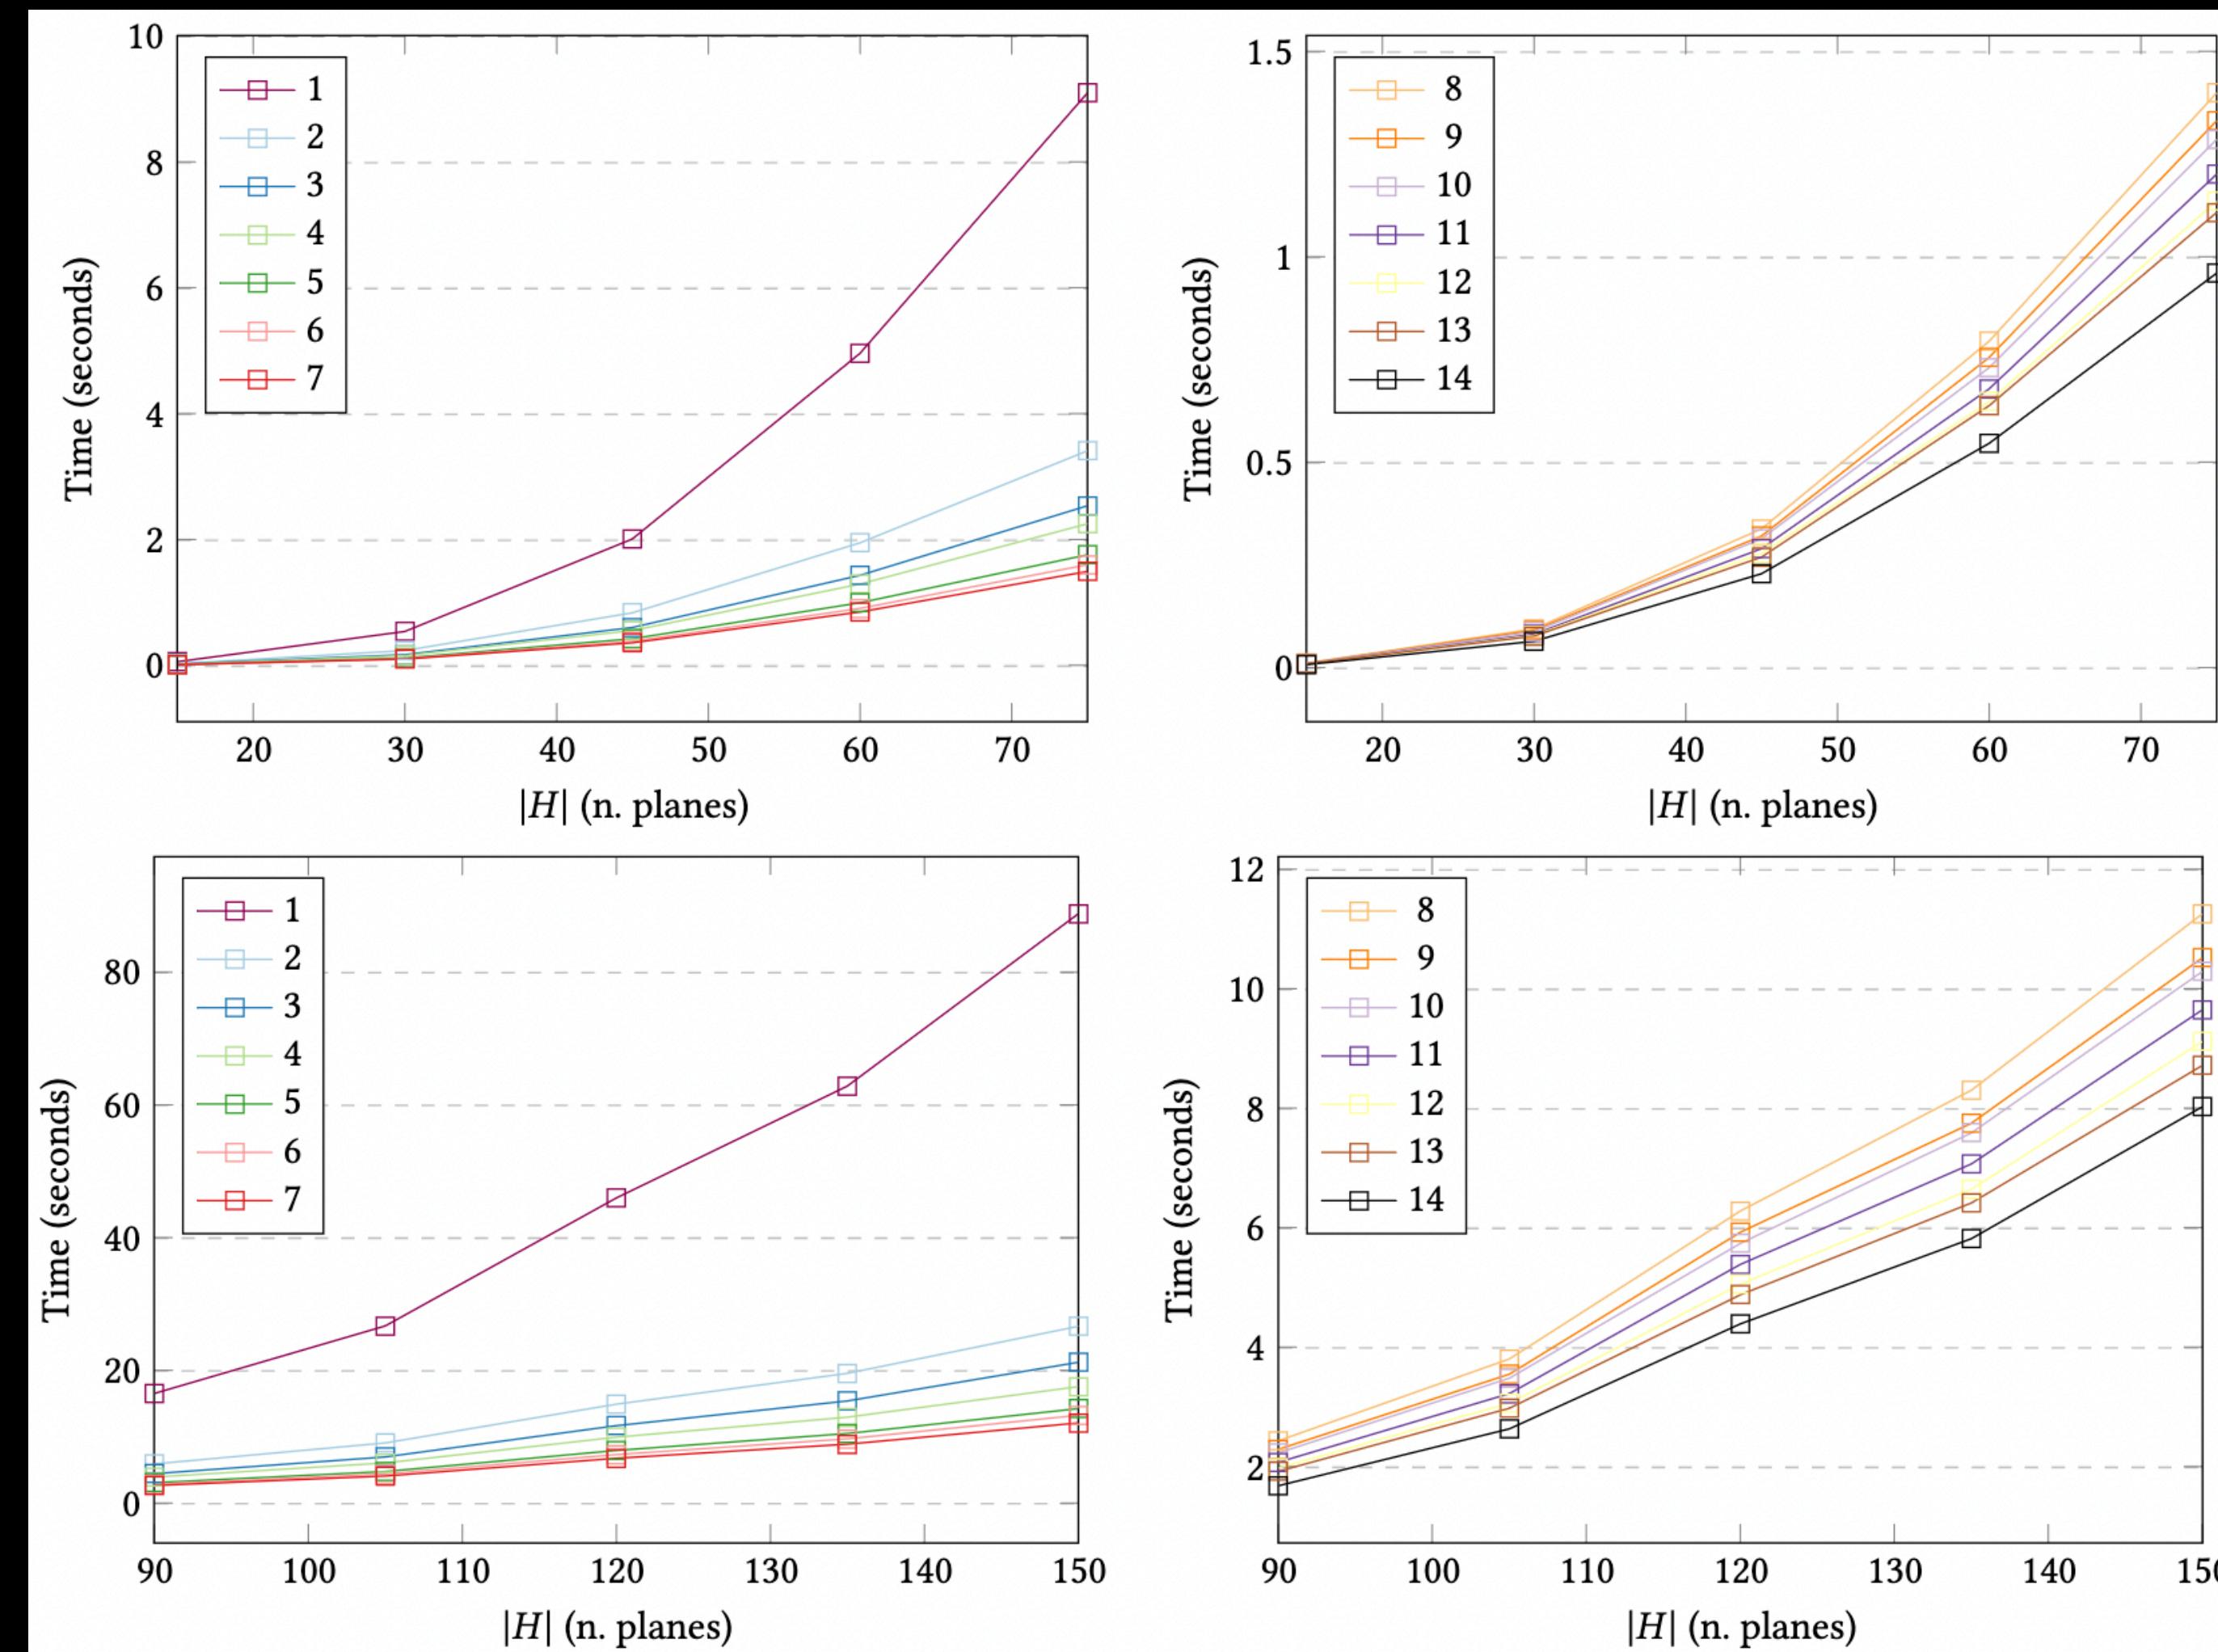
# Benchmarks and Results (1)

## Introduction

- Experiments run for 1 to 14 core versions
- CPU frequency scaling disabled
- 10 runs per input instance with SD  $\approx 1\%$  of total runtime

# Benchmarks and Results (2)

## Brute Force Algorithm



# Benchmarks and Results (3)

## Brute Force Algorithm

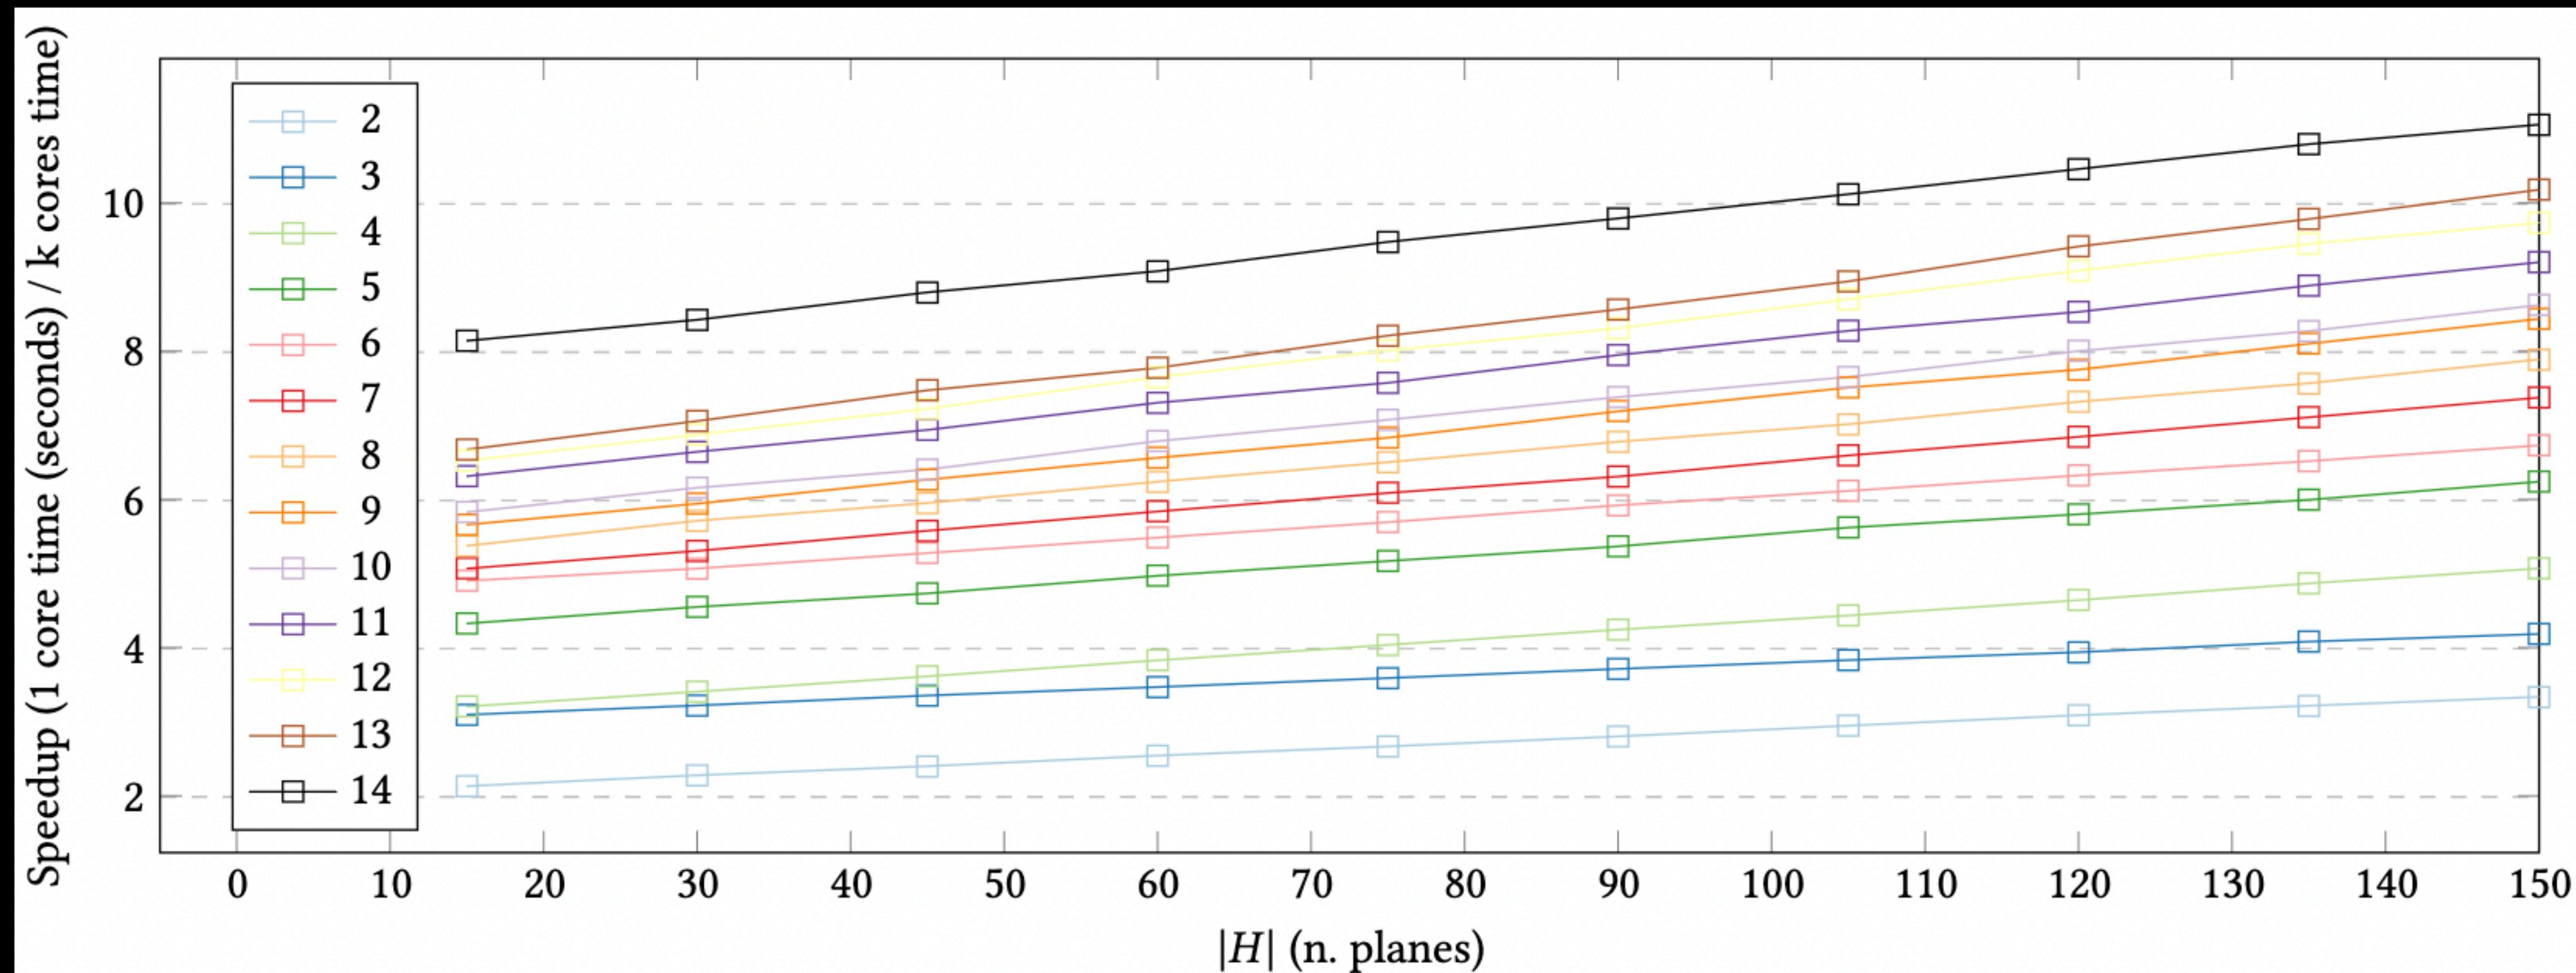
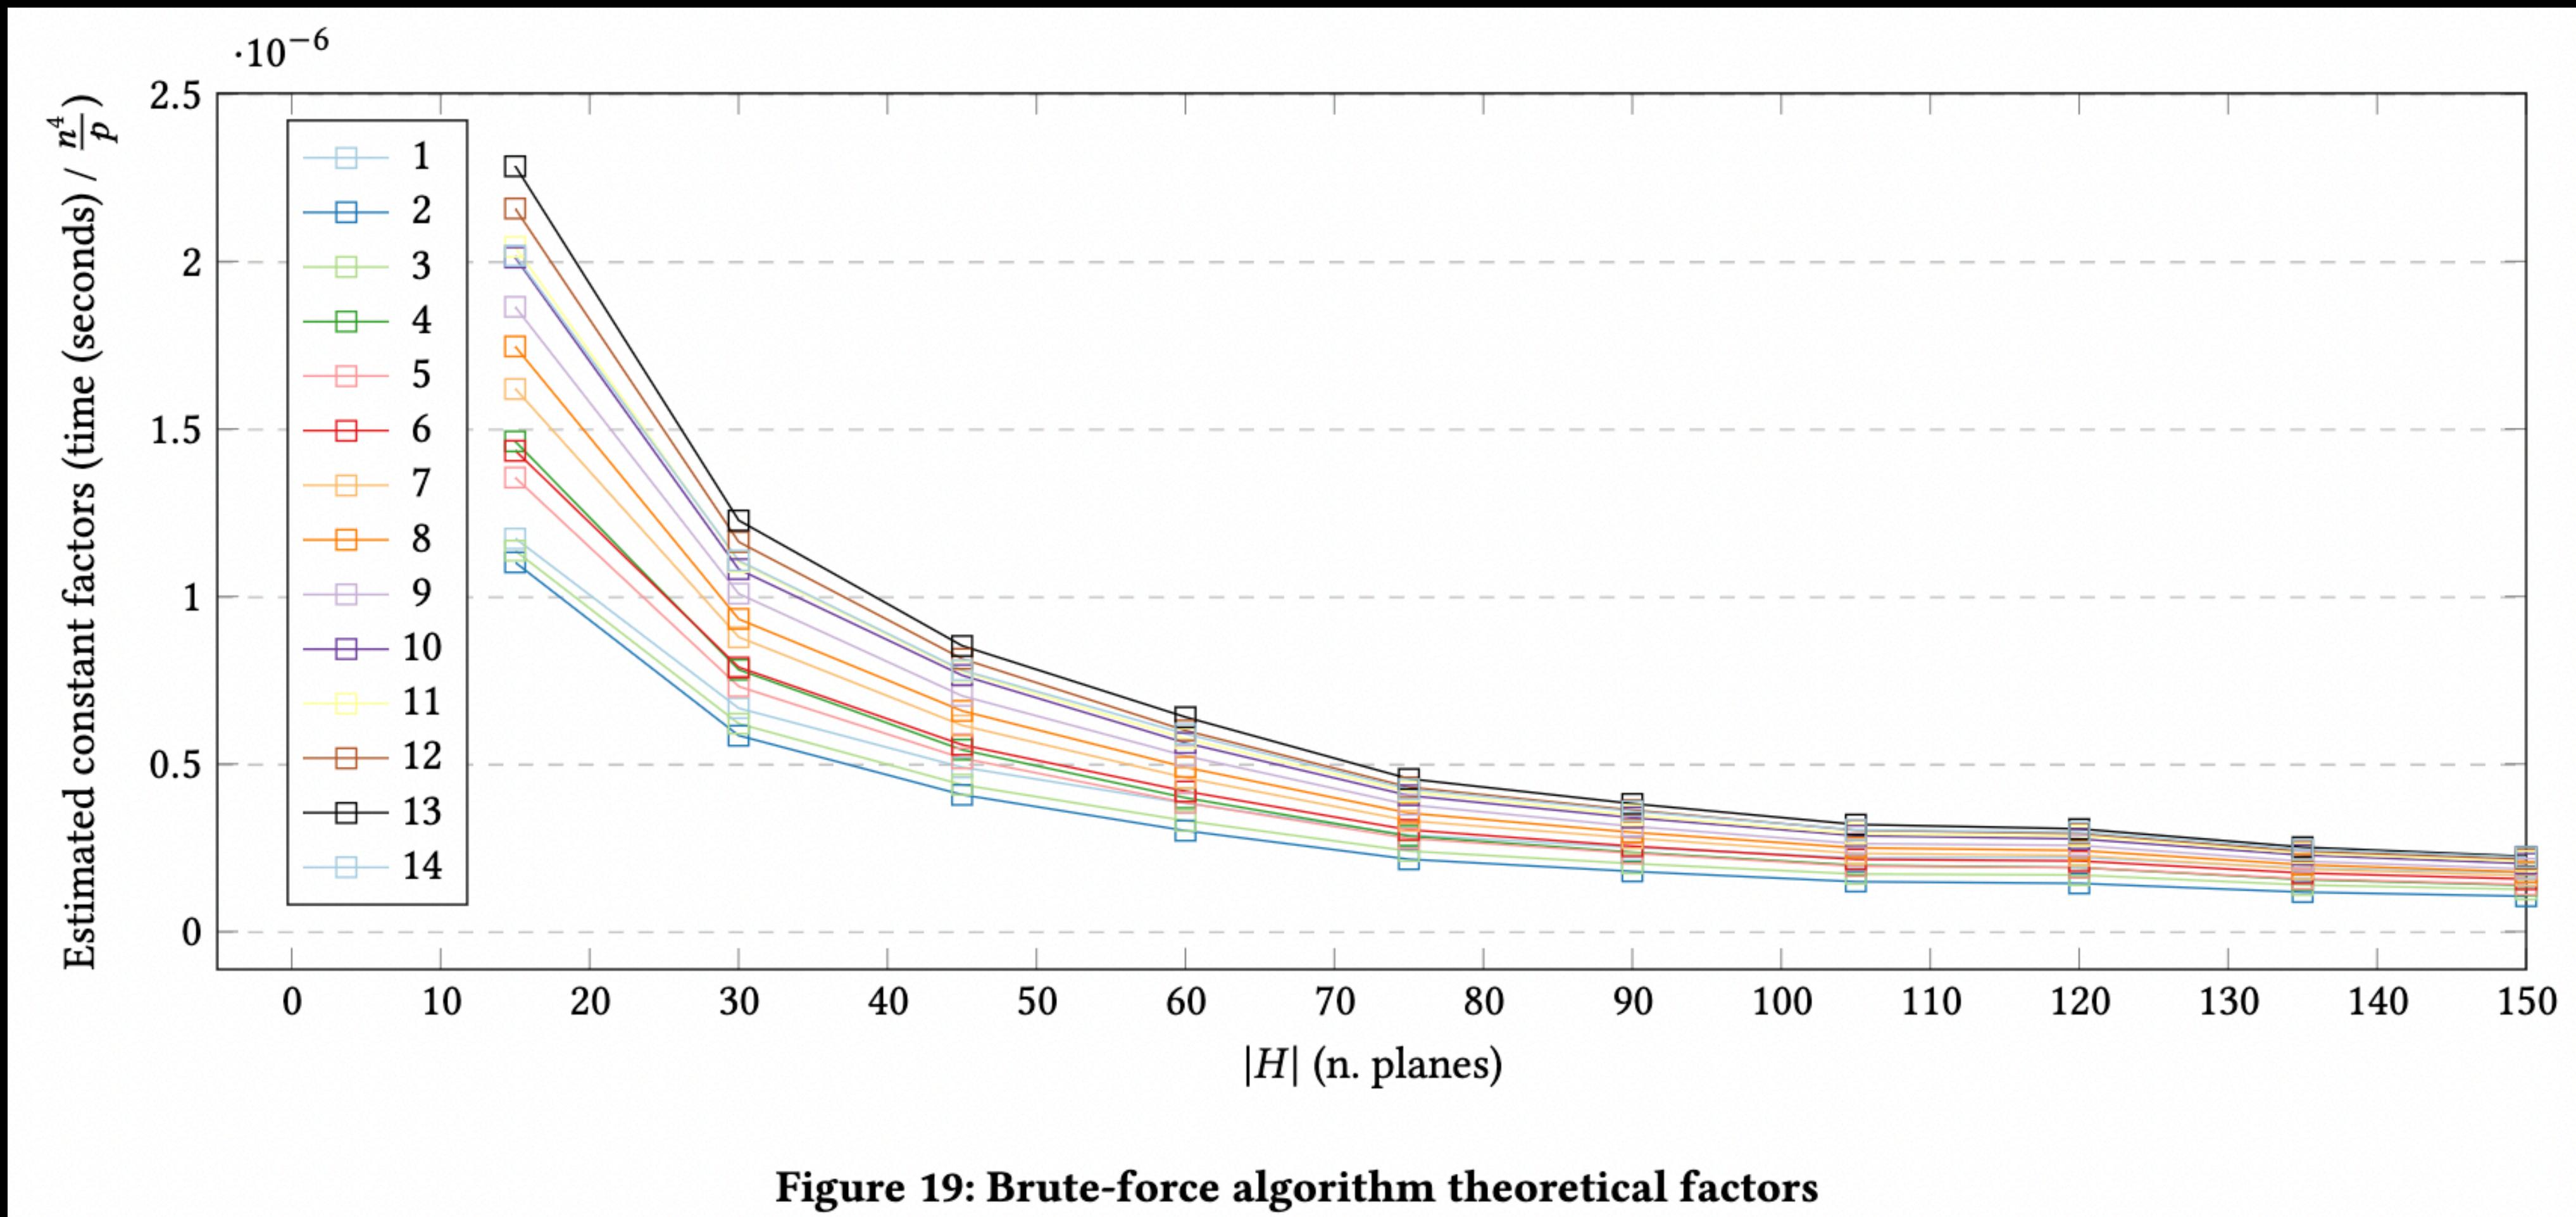


Figure 15: Brute-force algorithm parallel speed-ups

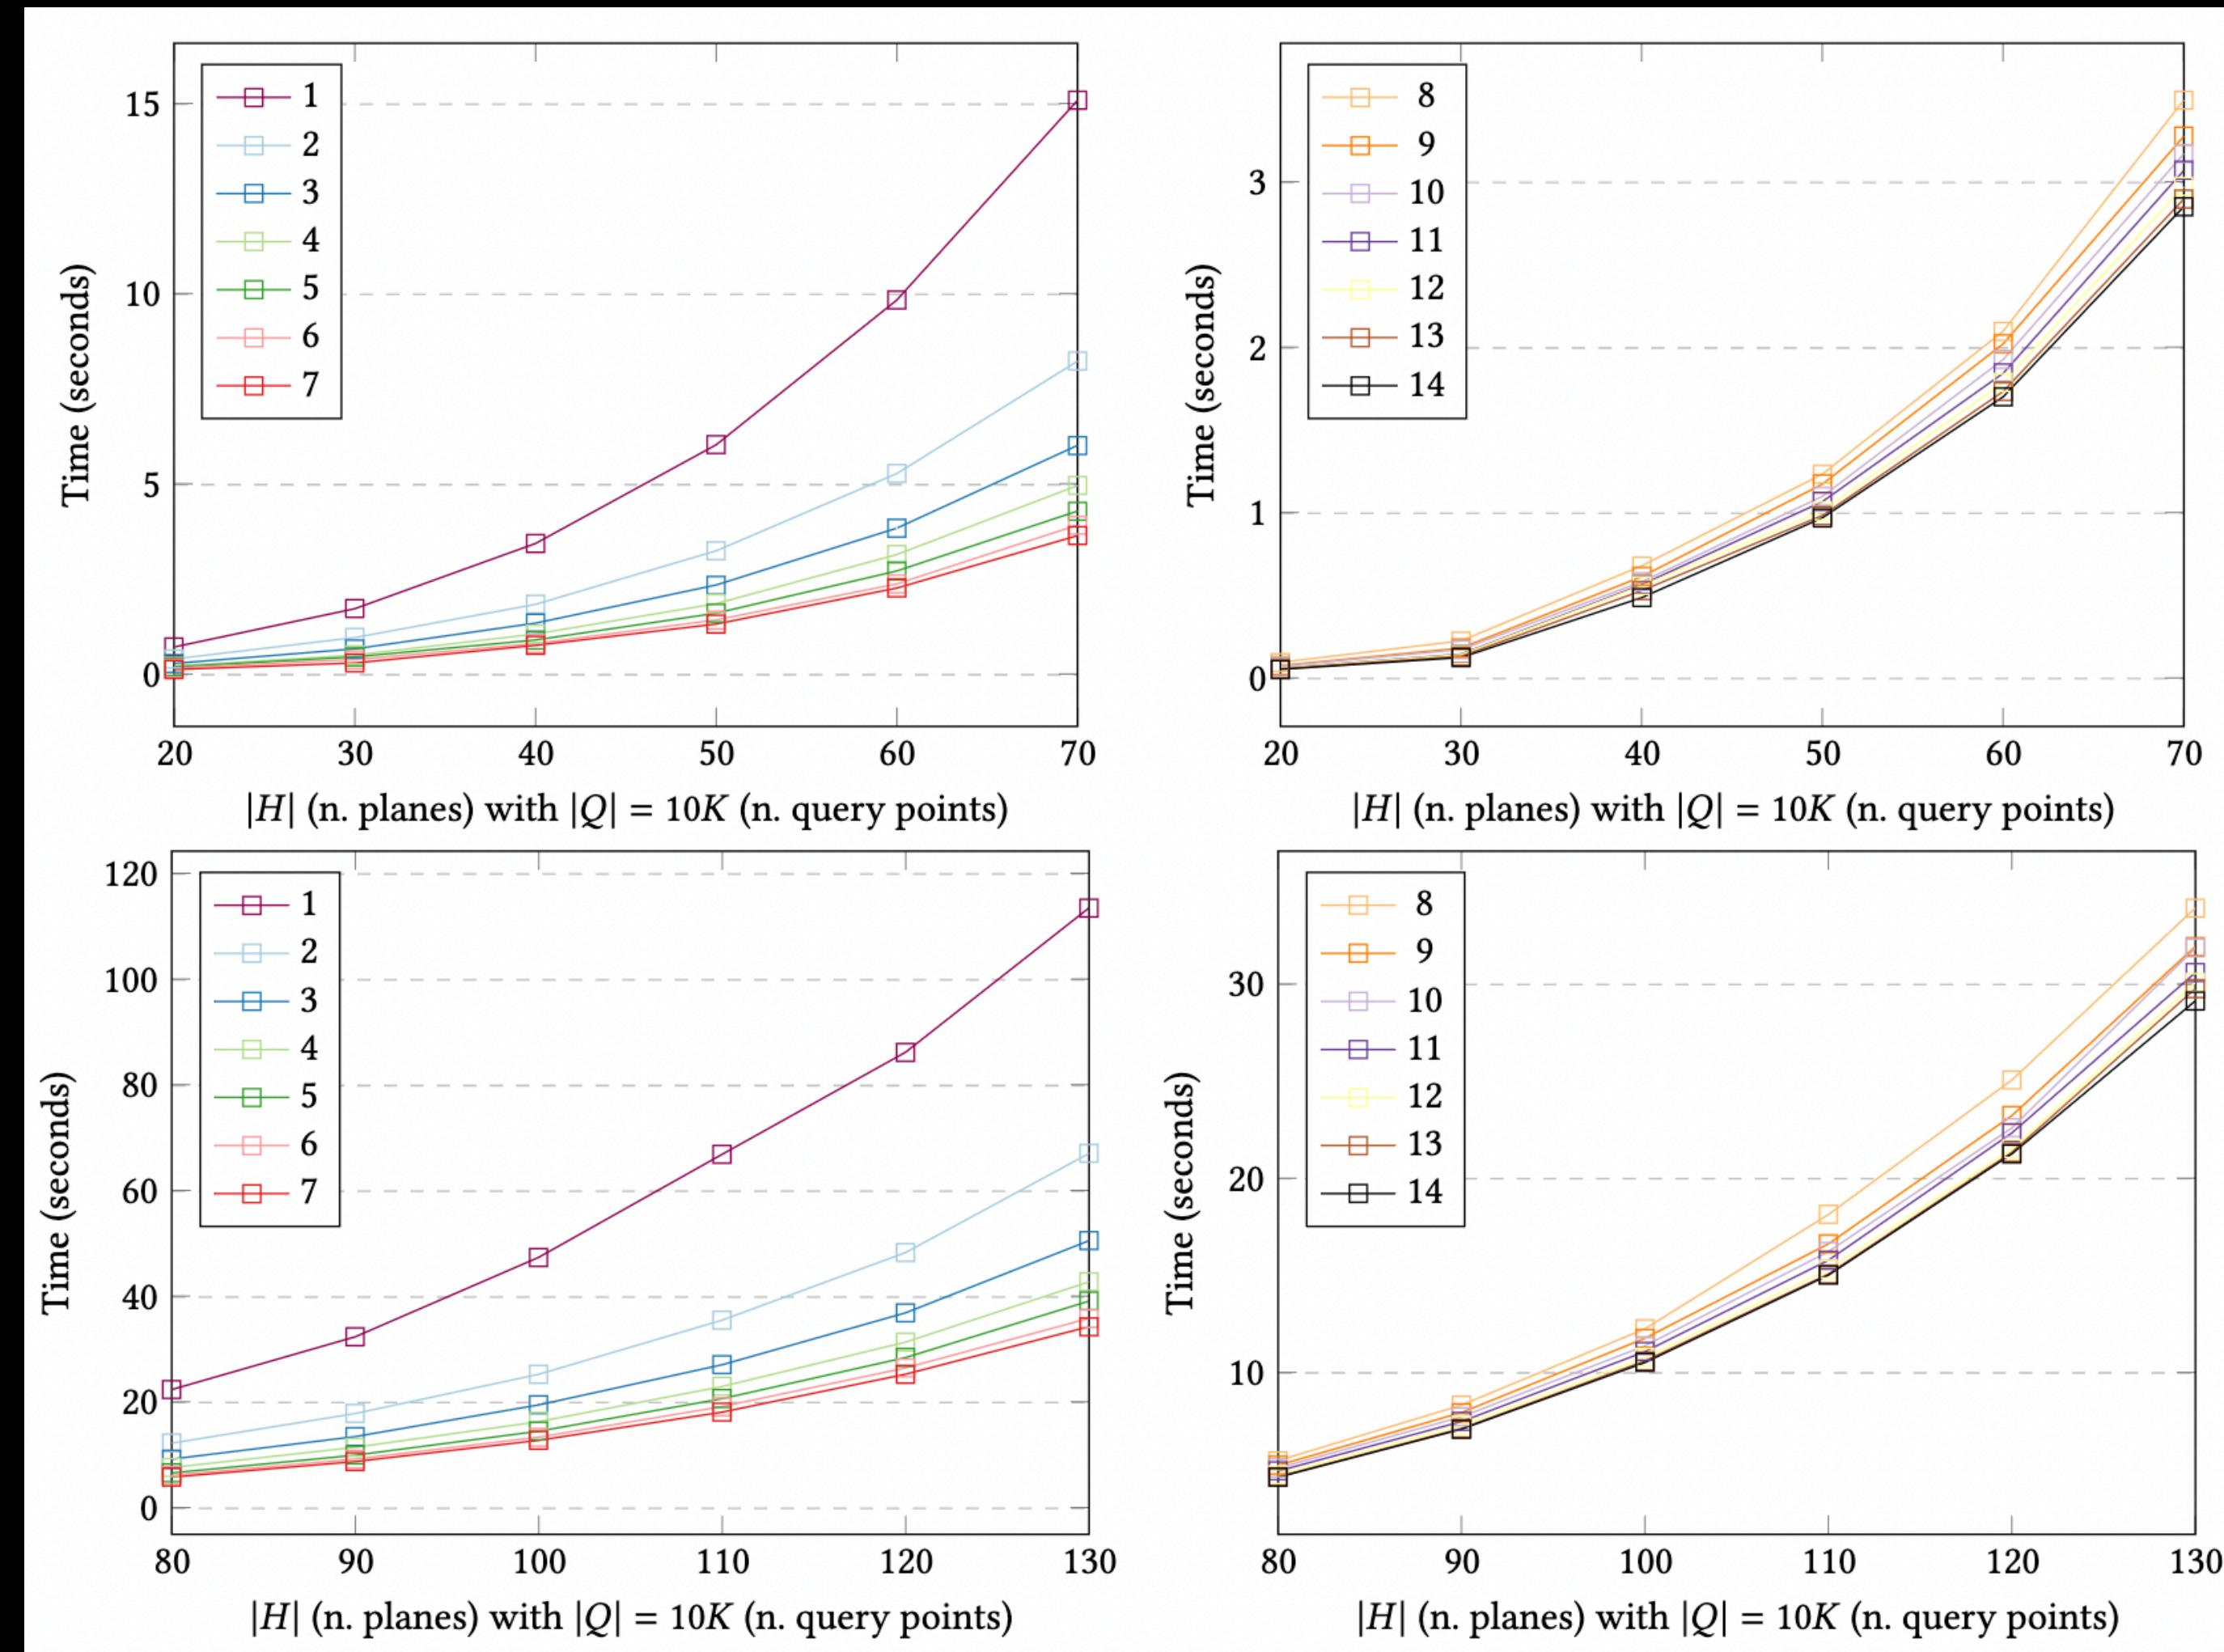
# Benchmarks and Results (4)

## Brute Force Algorithm



# Benchmarks and Results (5)

## Batch Point Location Algorithm



# Benchmarks and Results (6)

## Batch Point Location Algorithm

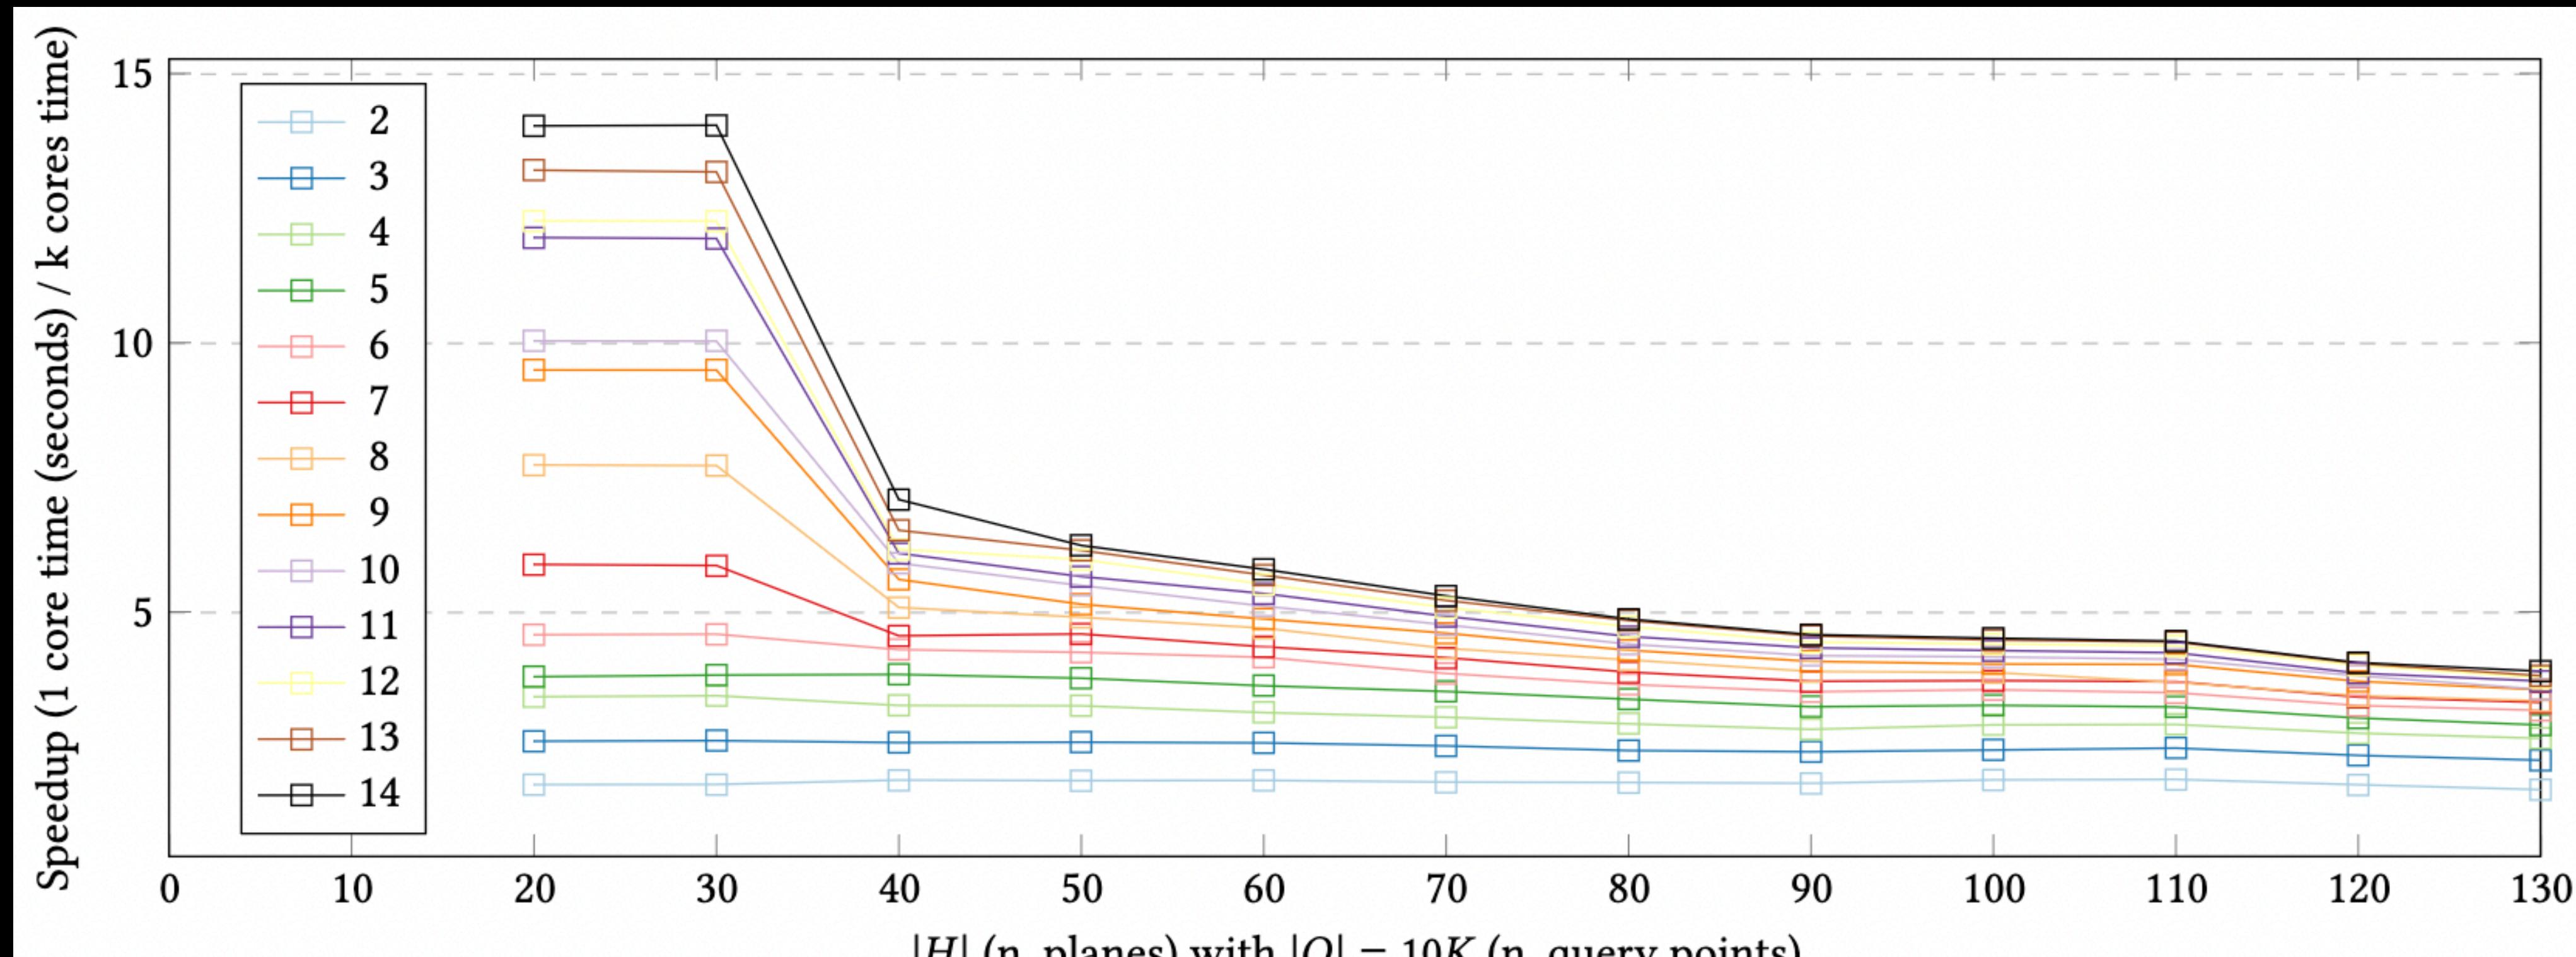


Figure 17: Batch-Point-Location algorithm parallel speed-ups

# Benchmarks and Results (7)

## Batch Point Location Algorithm

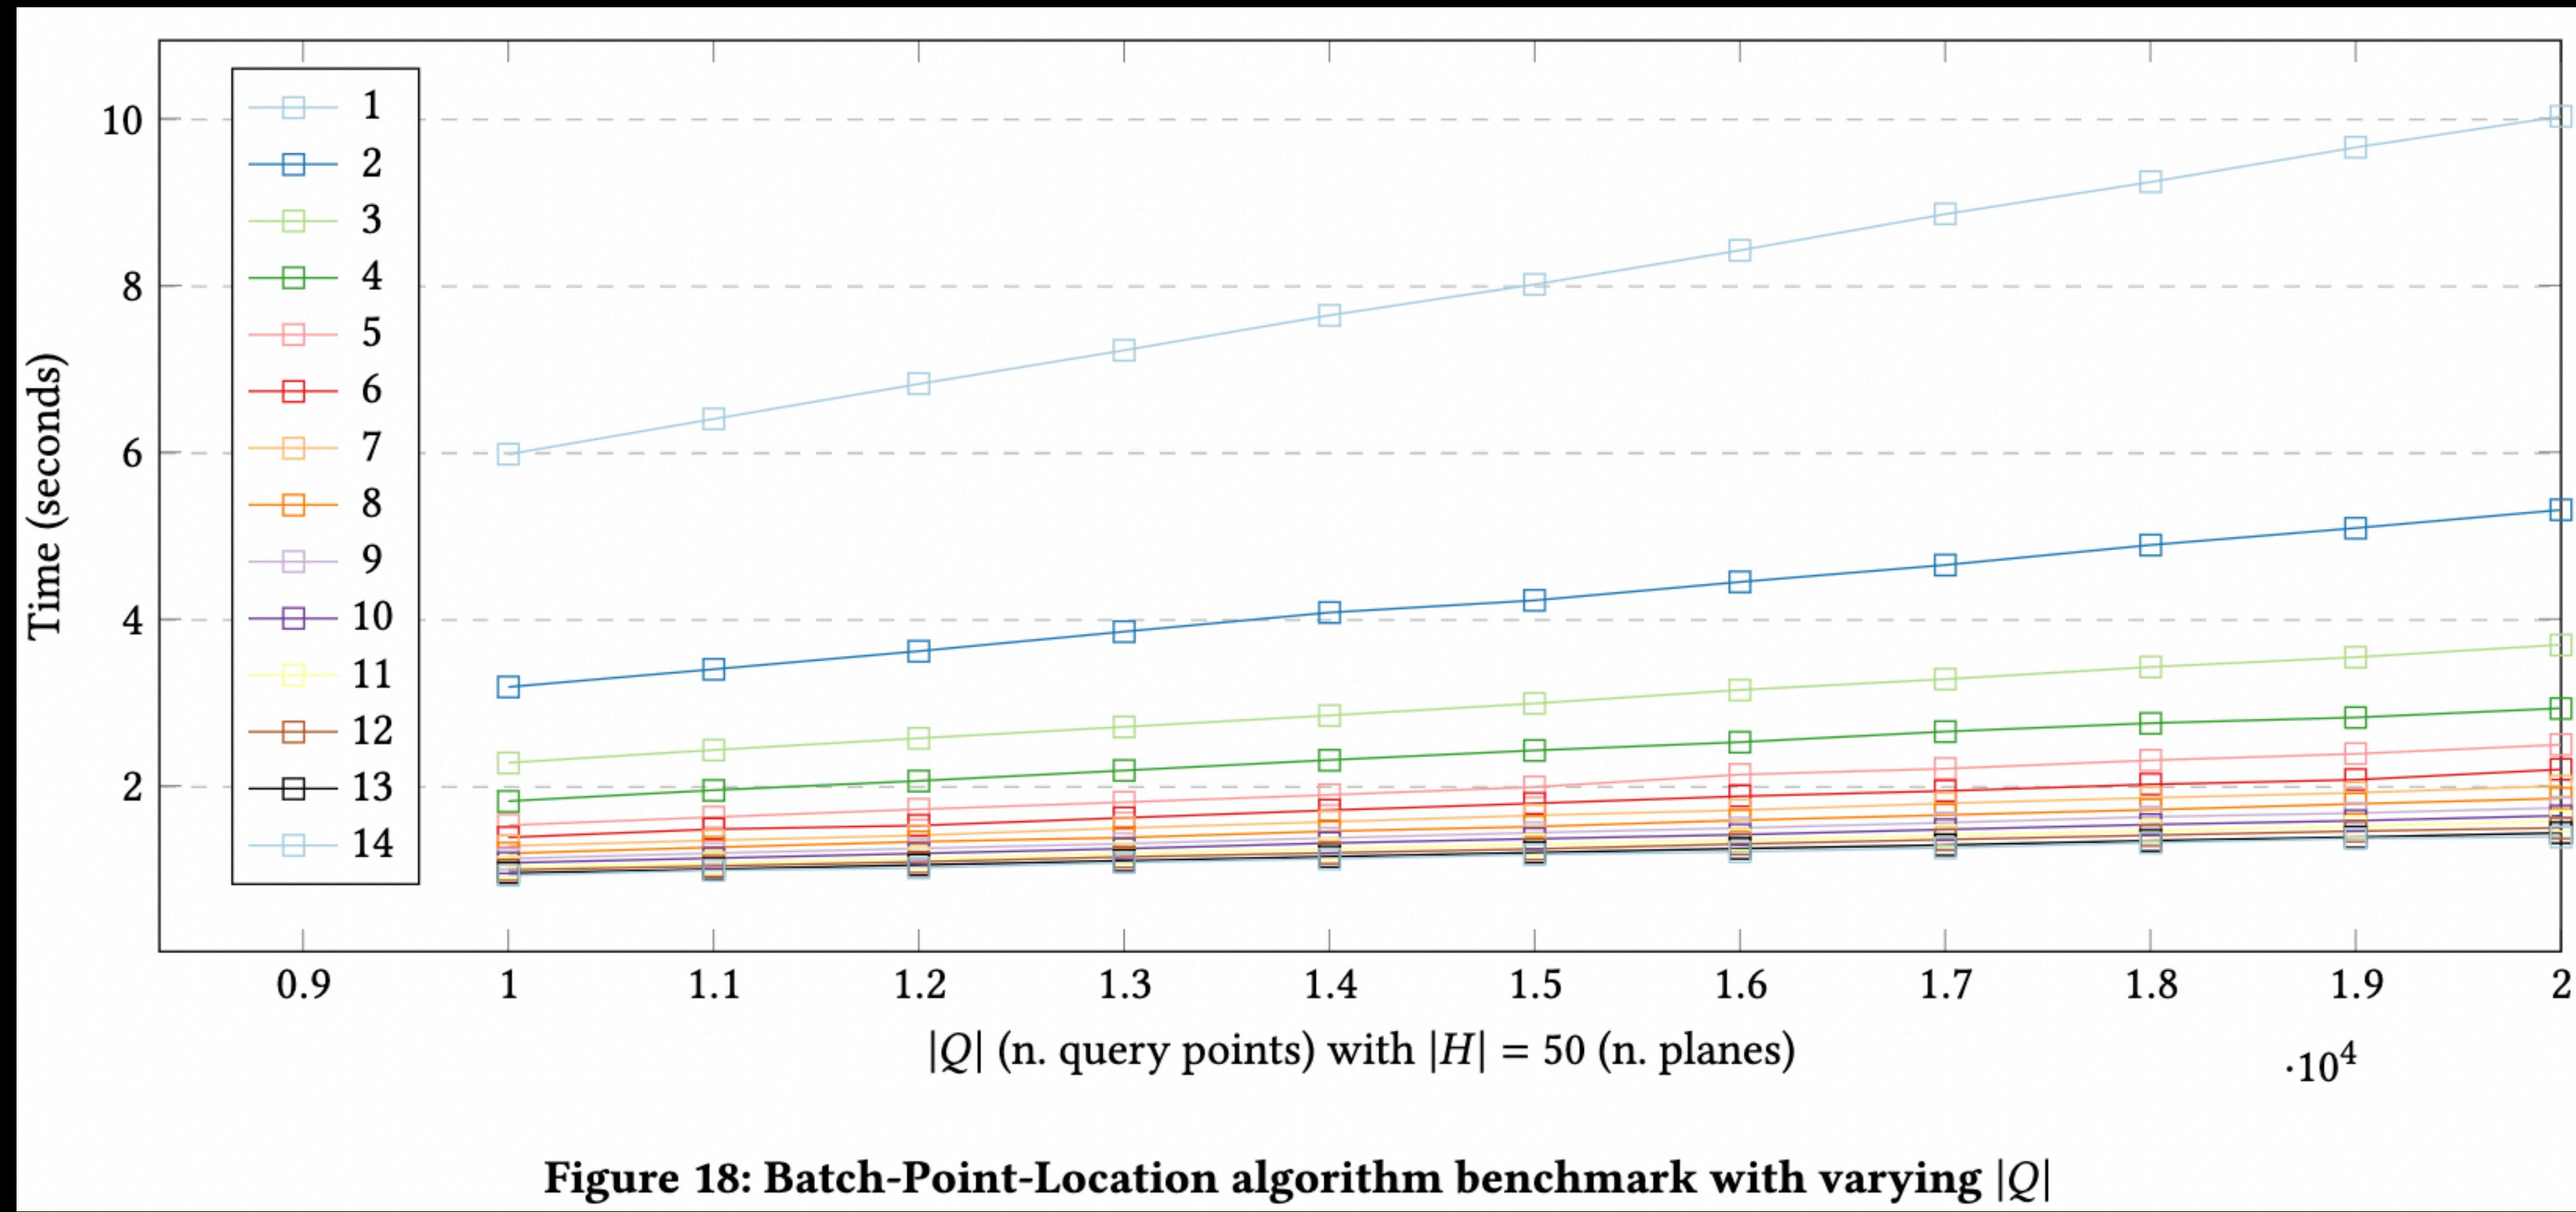
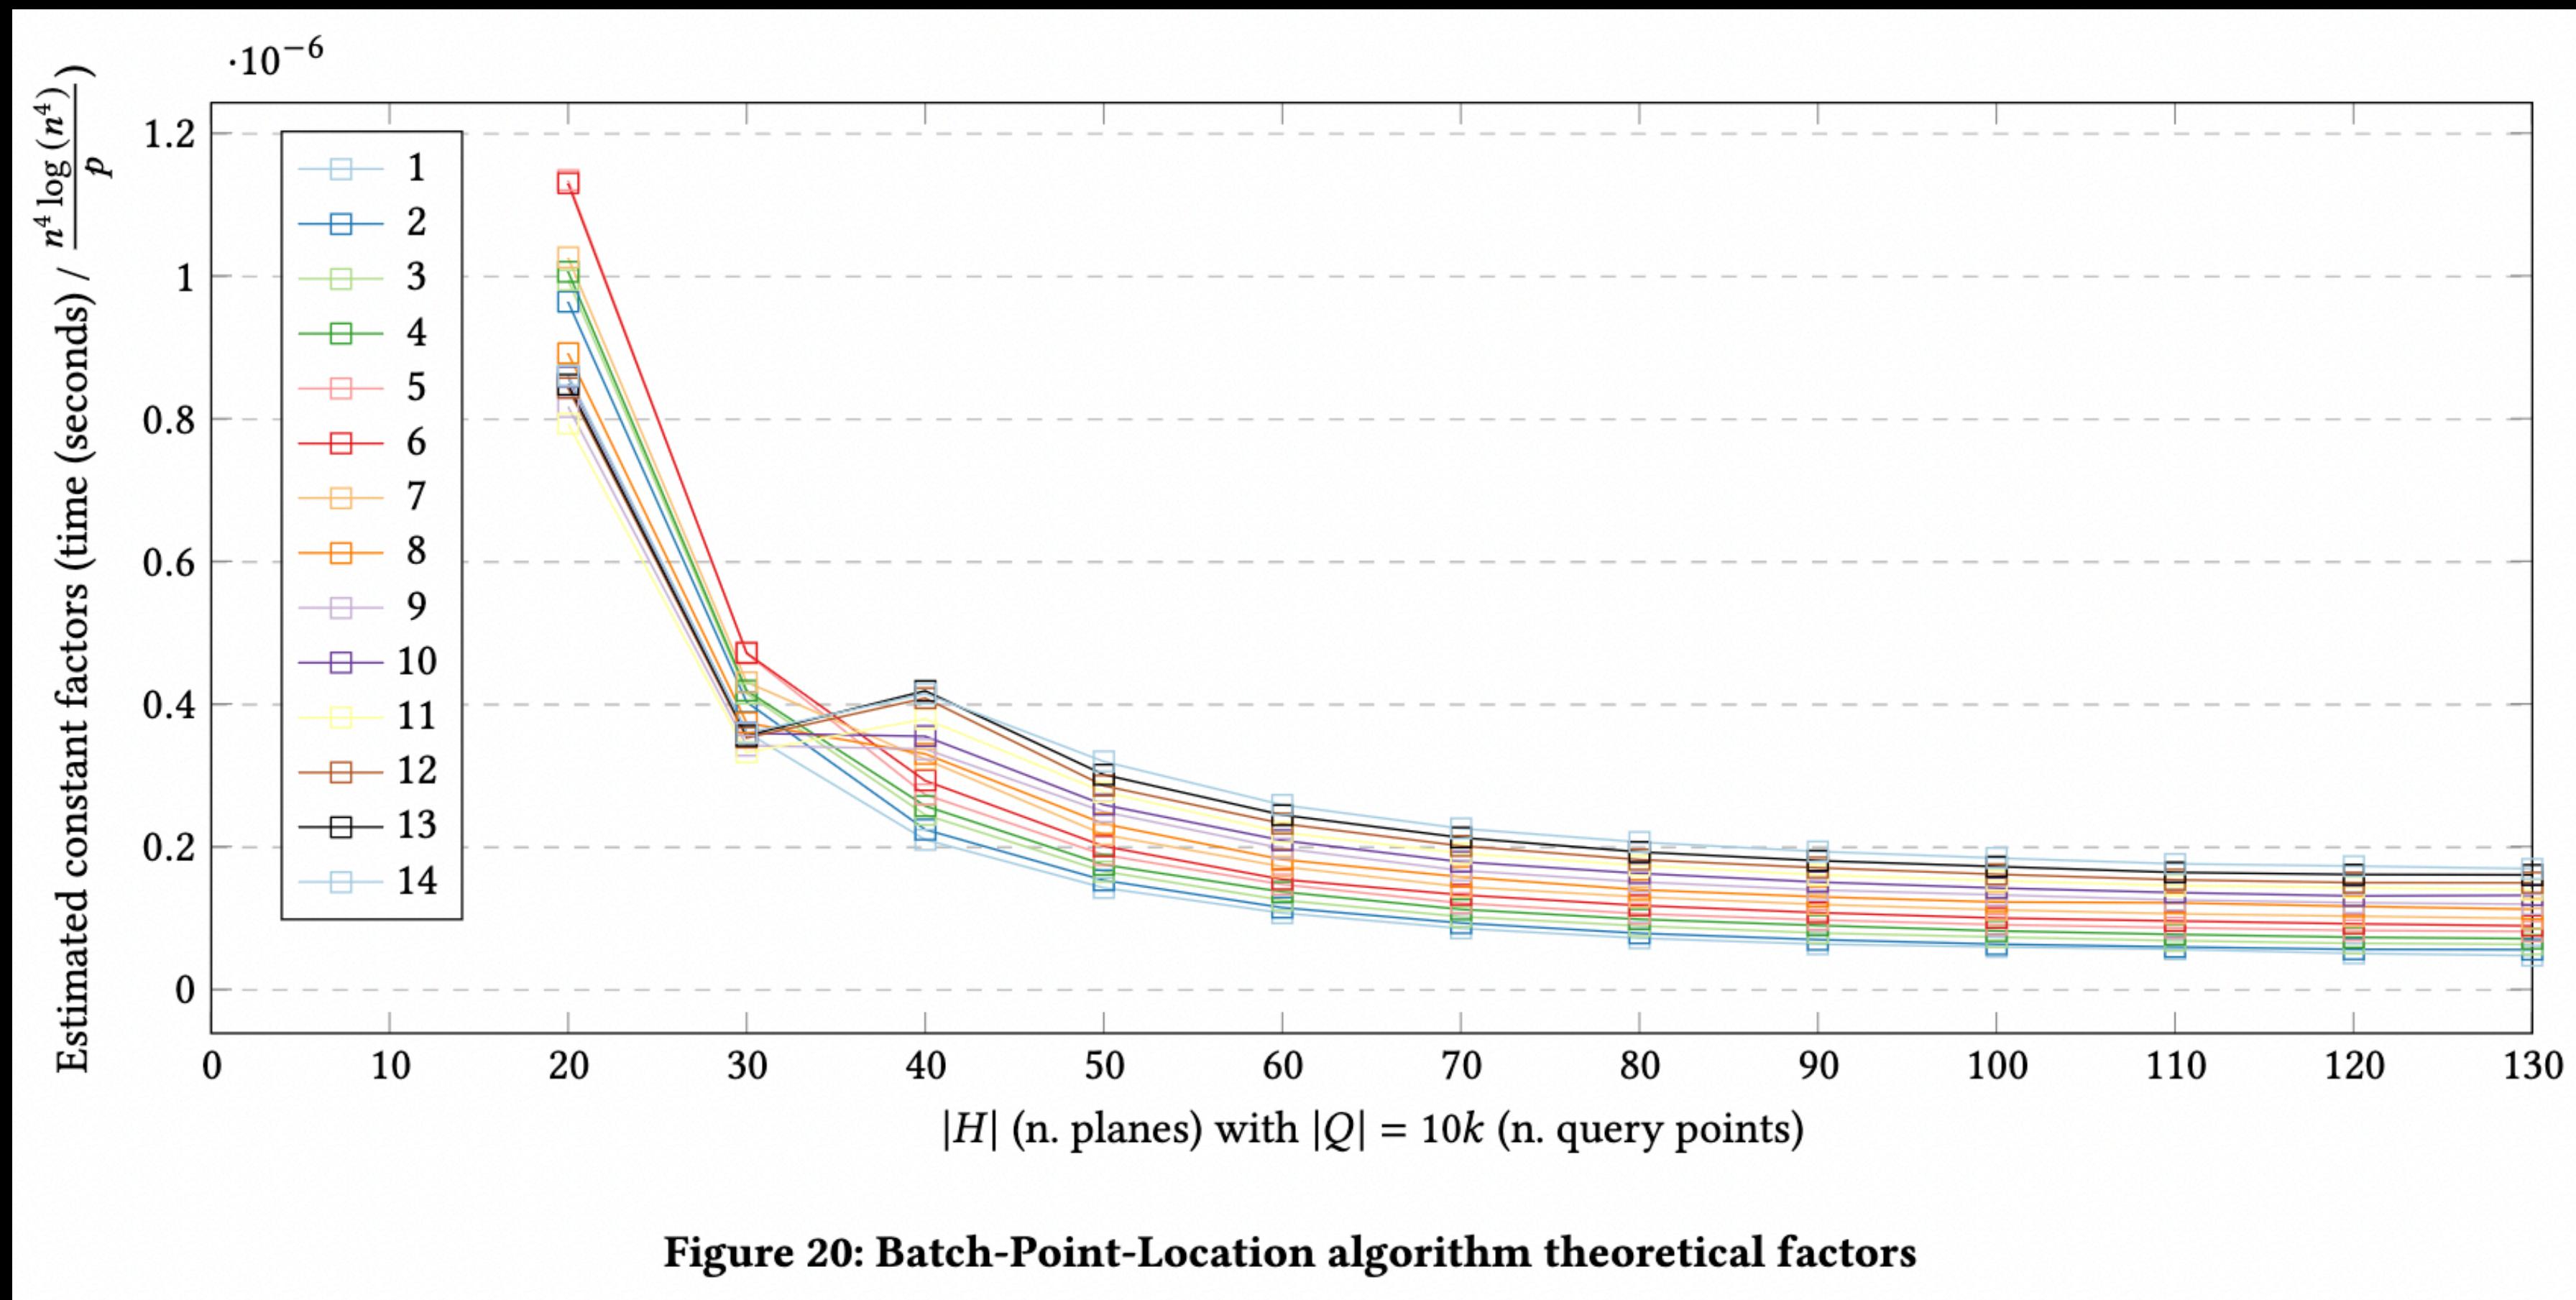


Figure 18: Batch-Point-Location algorithm benchmark with varying |Q|

# Benchmarks and Results (8)

## Batch Point Location Algorithm



# Benchmarks and Results (9)

## Future Work

- The main algorithm requires additional theoretical work concerning the sampling scheme used
- The batch point location algorithm requires I/O and cache-efficient slab and zone search routines
- It would be interesting to adapt the current implementations to use inexact arithmetic and measure the performance improvement

# References

1. Timothy Chan and Eric Chen. 2010. Optimal in-place and cache-oblivious algorithms for 3-d convex hulls and 2-d segment intersection. *Computational Geometry* 43 (10 2010), 636–646. <https://doi.org/10.1016/j.comgeo.2010.04.005>.
2. Frank Staals. 2023. Parallel 3D Lower Envelope/Convex Hull. Technical Report (2023).
3. Gianmarco Picarella. 2023. 3D Lower Envelope Algorithms. <https://github.com/gianmarcopicarella/gaspmt-uu>.

**Thank you! Questions?**